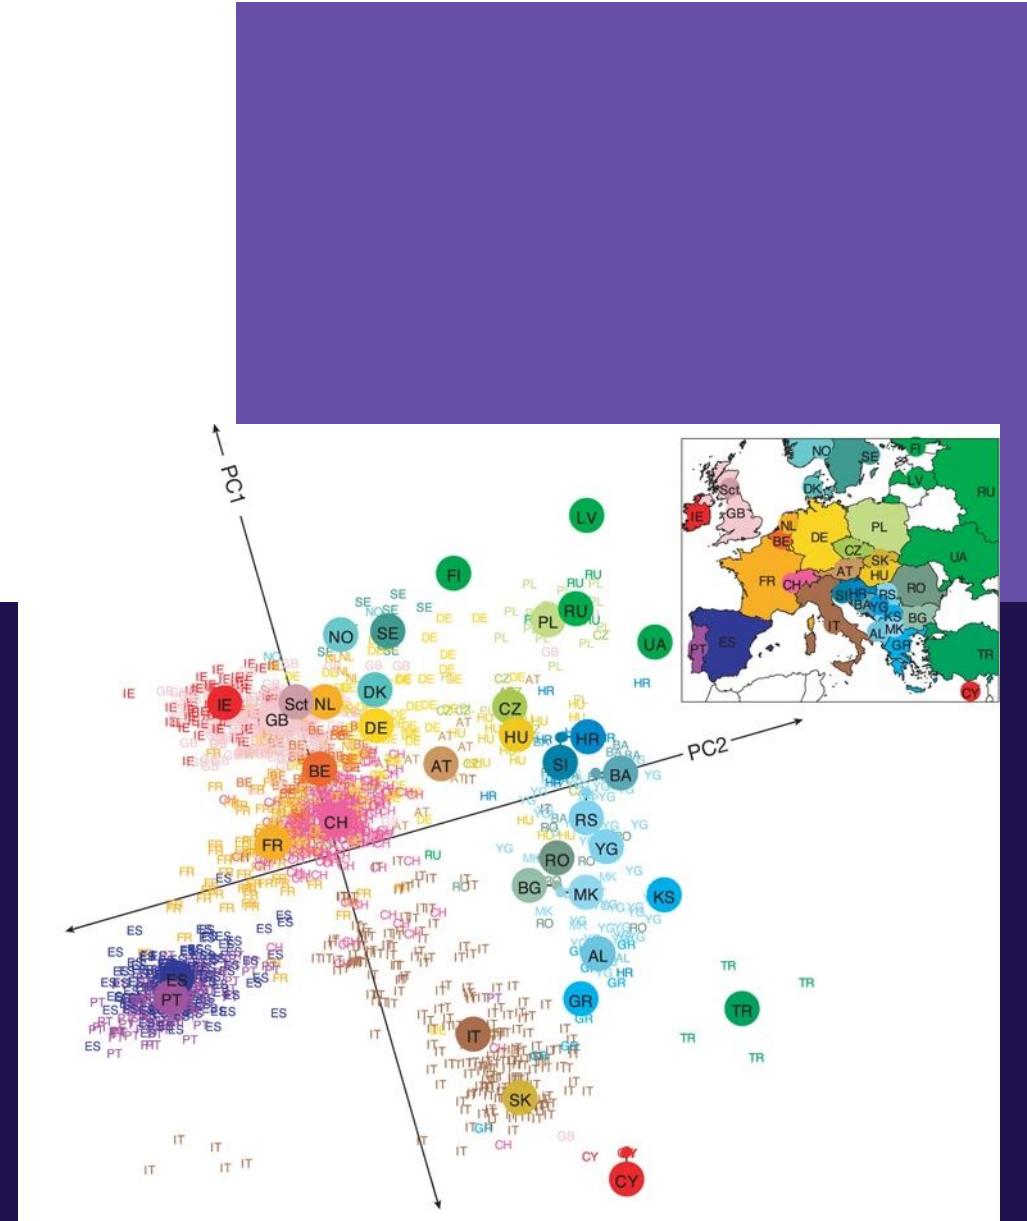


Genome scans

28/11/2024

Physalia course

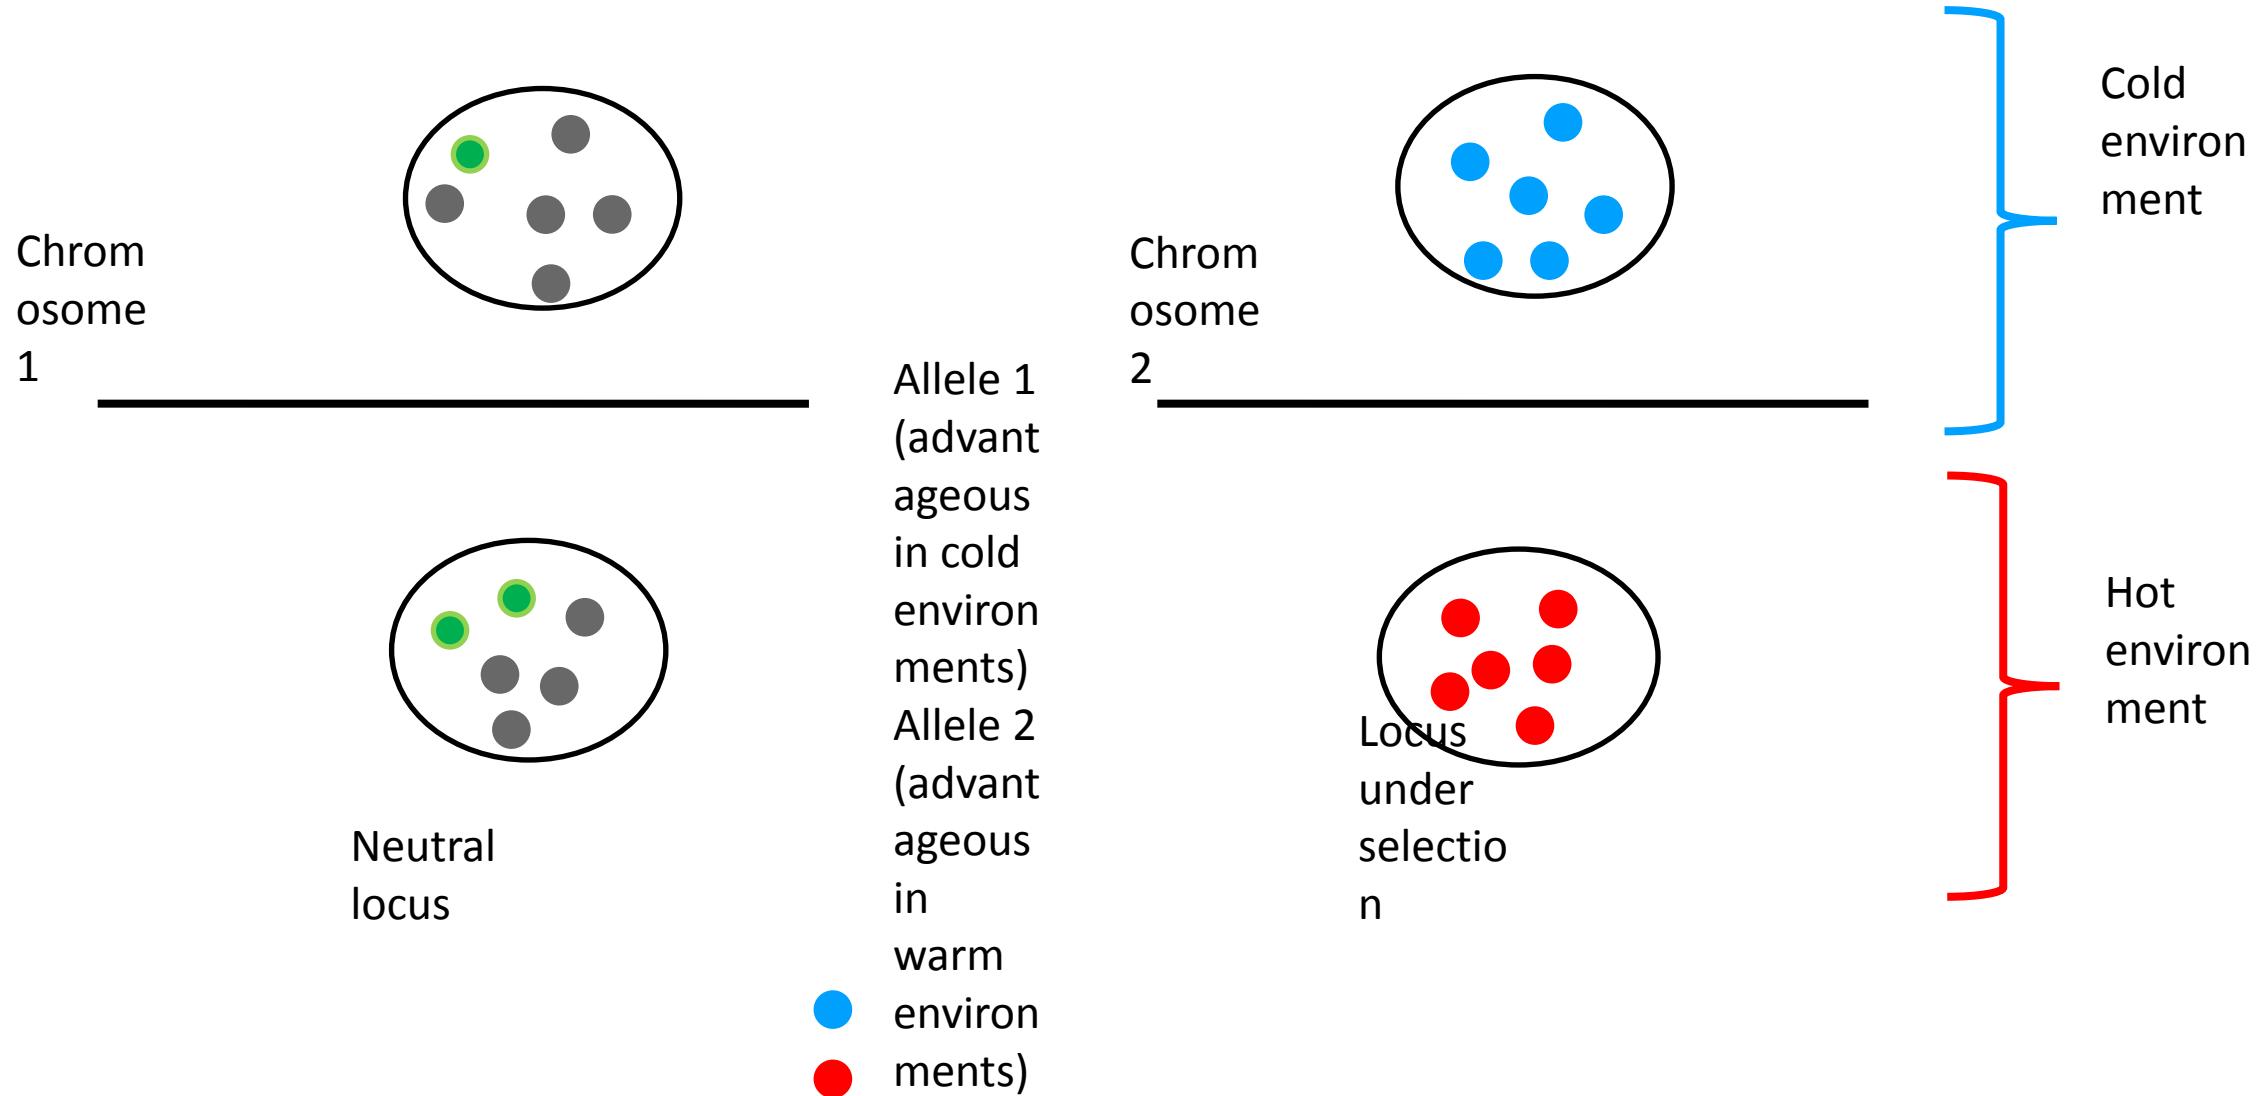
Yann Bourgeois, Thibault Leroy



Goals for today's lecture

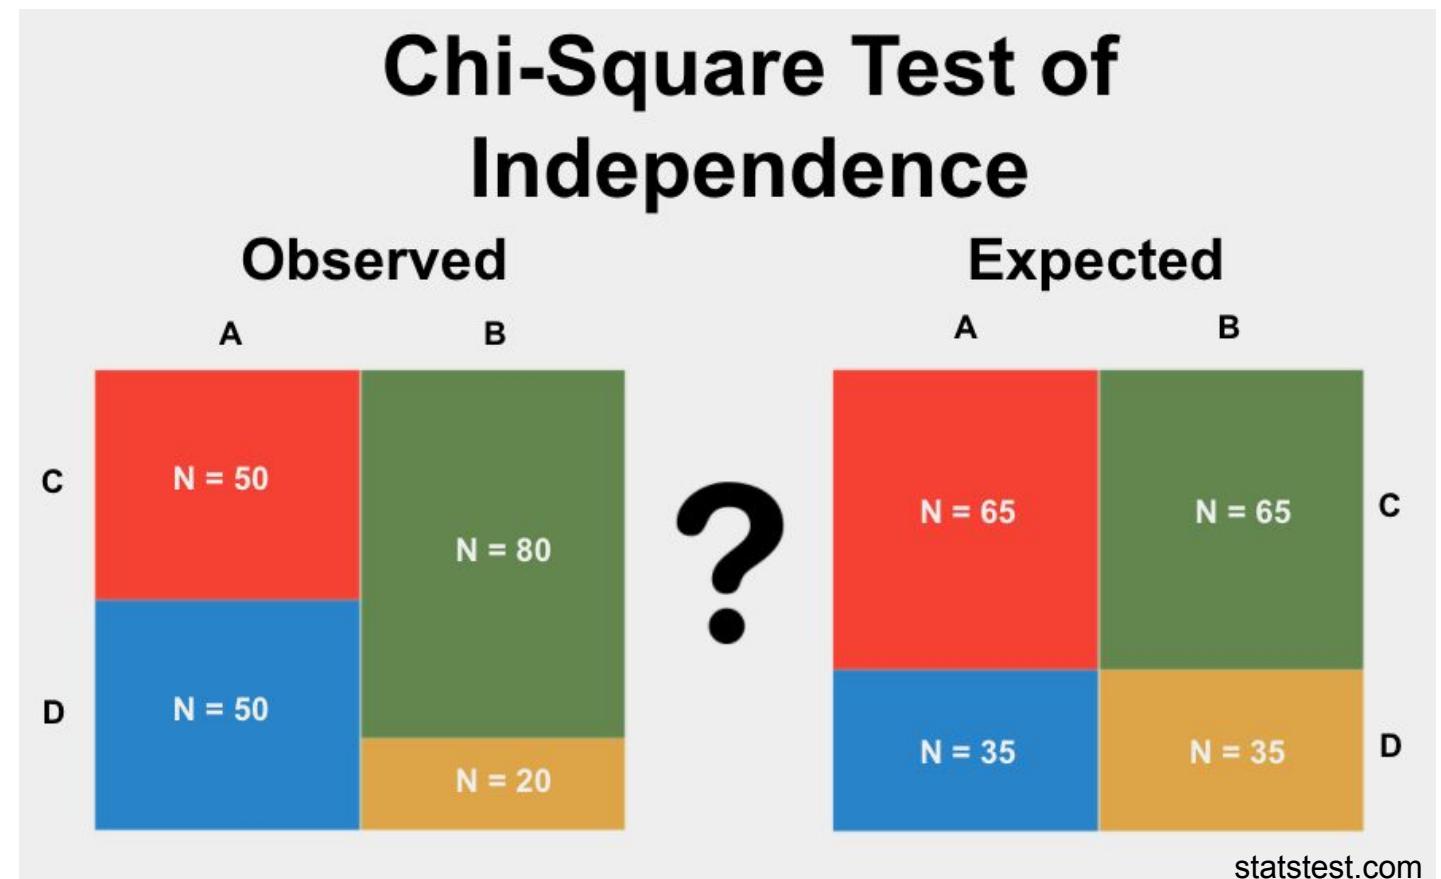
- What is a genome-wide association study (GWAS)?
- How can we detect selection in a single population?
- How do we detect divergent selection?
- A few methods to detect balancing selection.
- The perils of interference and linkage.
- As usual: a few limitations to keep in mind.

Genome scans for association/adaptation



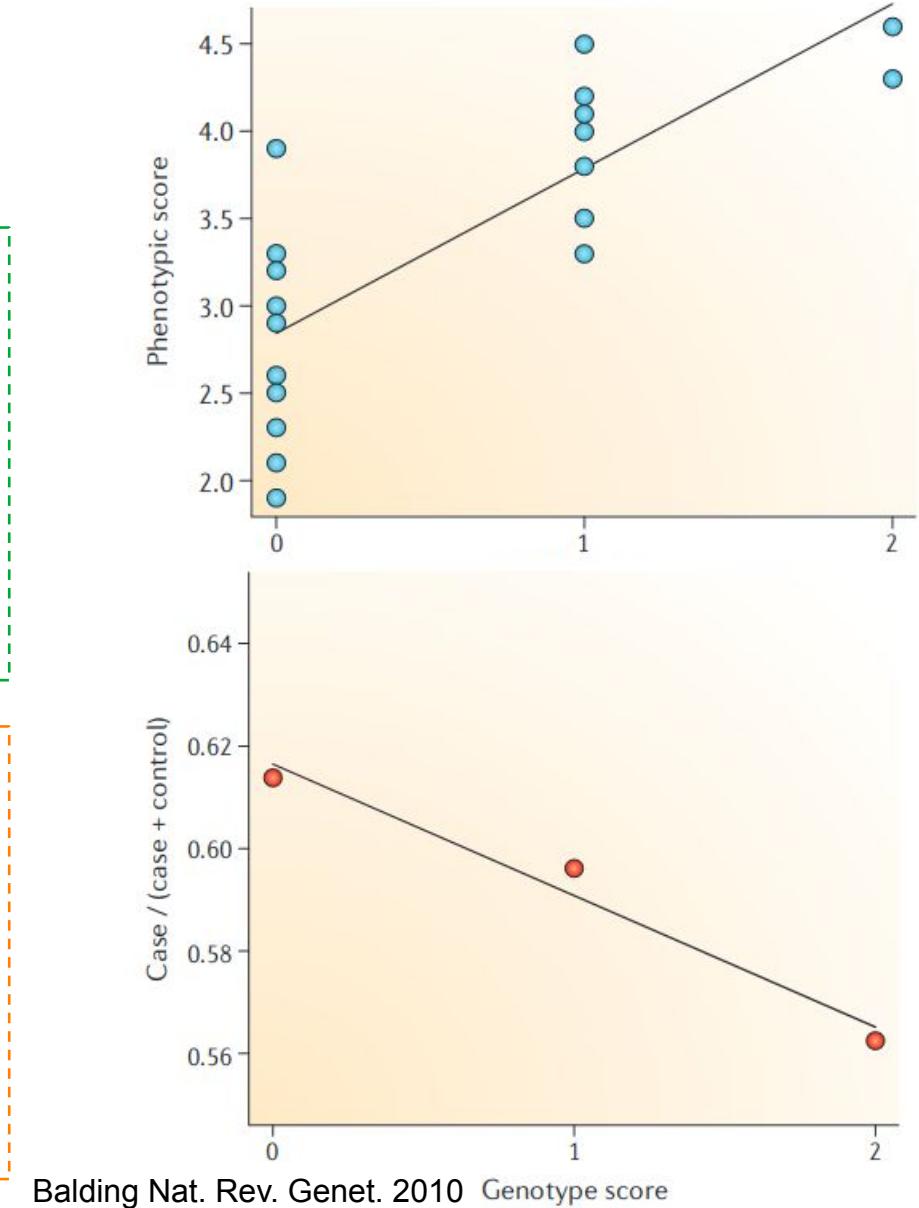
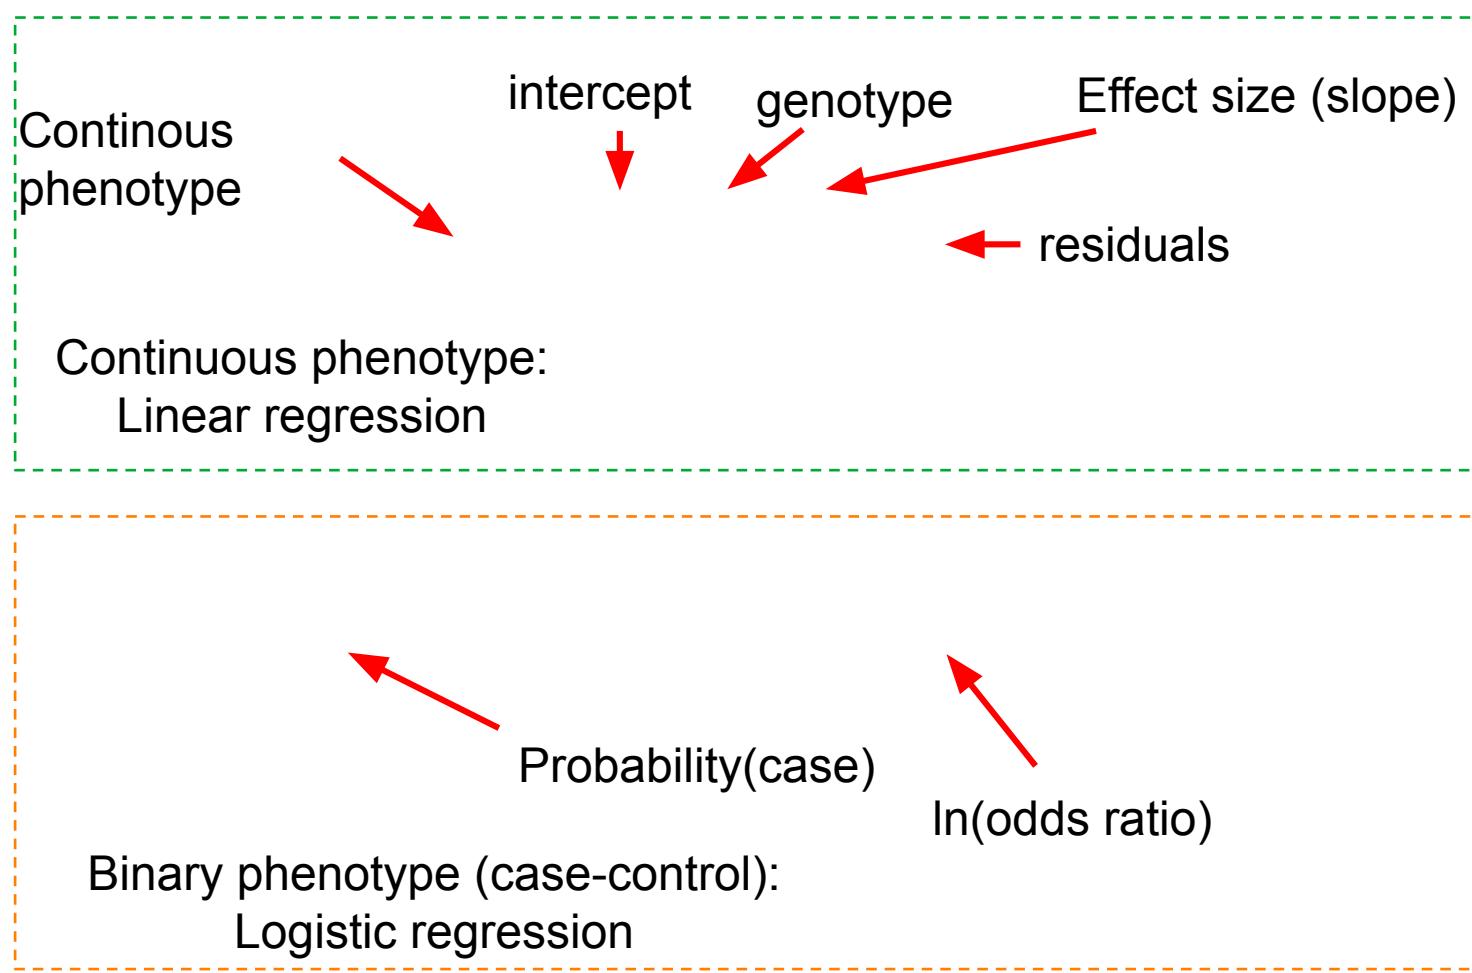
Genome Wide Scans of Association (GWAS)

- The simplest: χ^2 test (--assoc in PLINK1.9, cf workshop)
- Works only for discrete categories
- Categories are not ordered
- Not possible to directly account for structure or covariates.



Genome Wide Scans of Association (GWAS)

- Basic principle: a linear model of association.



Genome Wide Scans of Association (GWAS)

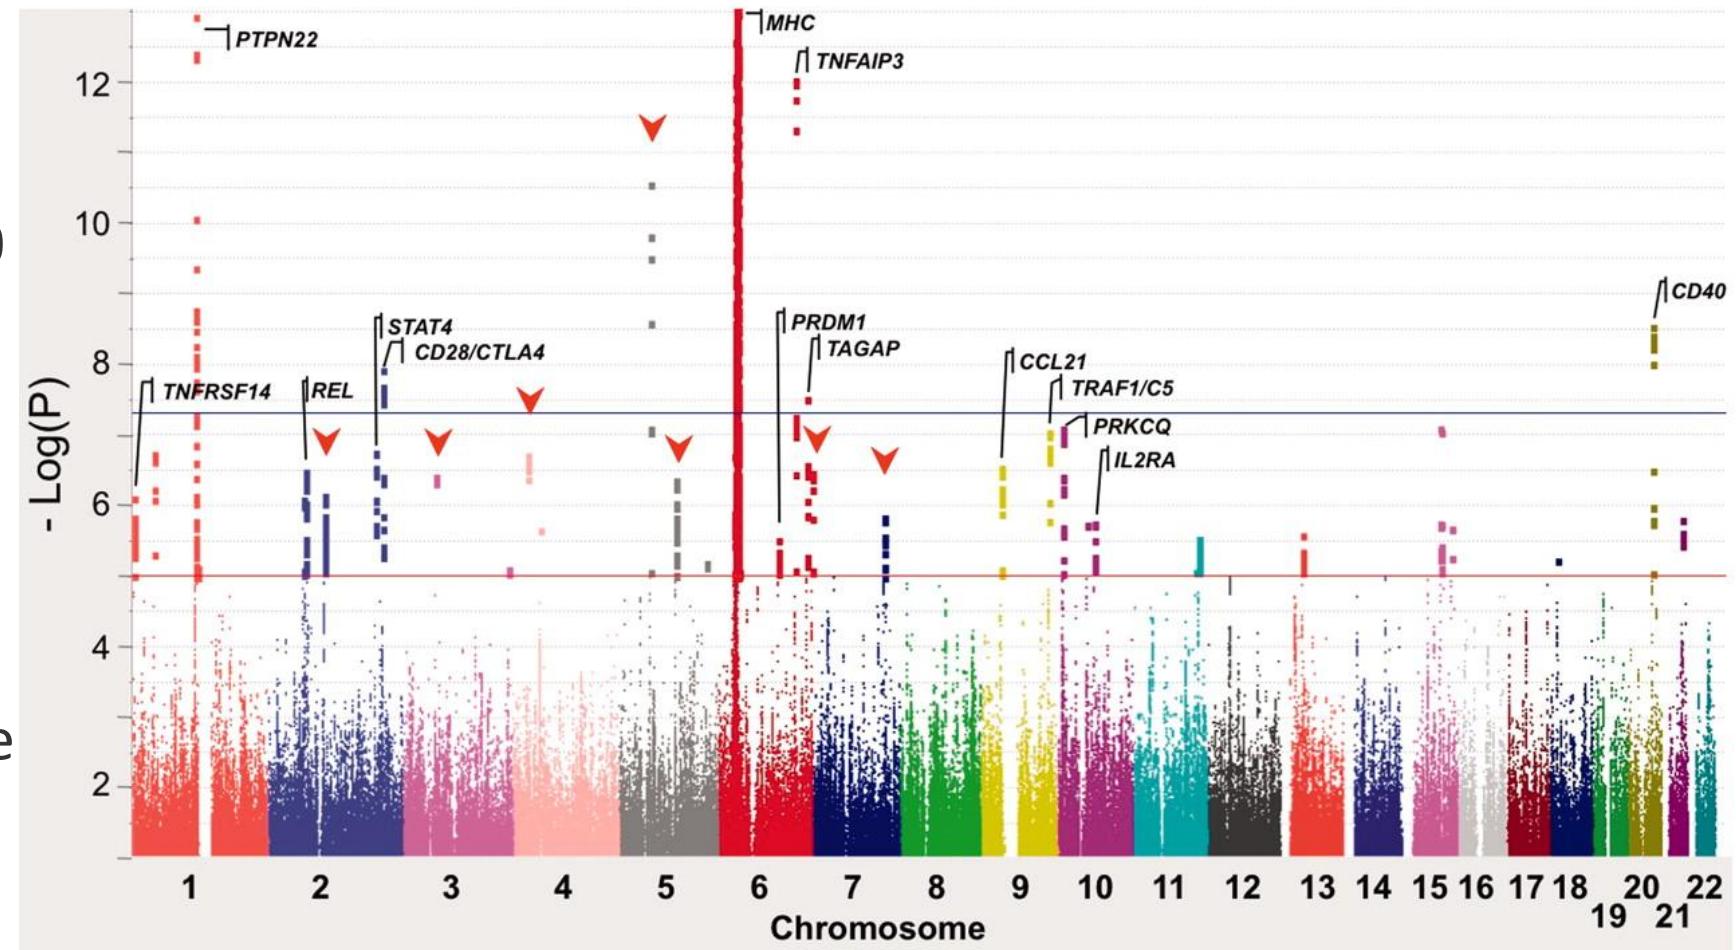
- We can add confounding variables and complexify the model if desired.
- Many statistical methods to increase power
- Unless you are working with very large panels, do not overinterpret differences between methods.



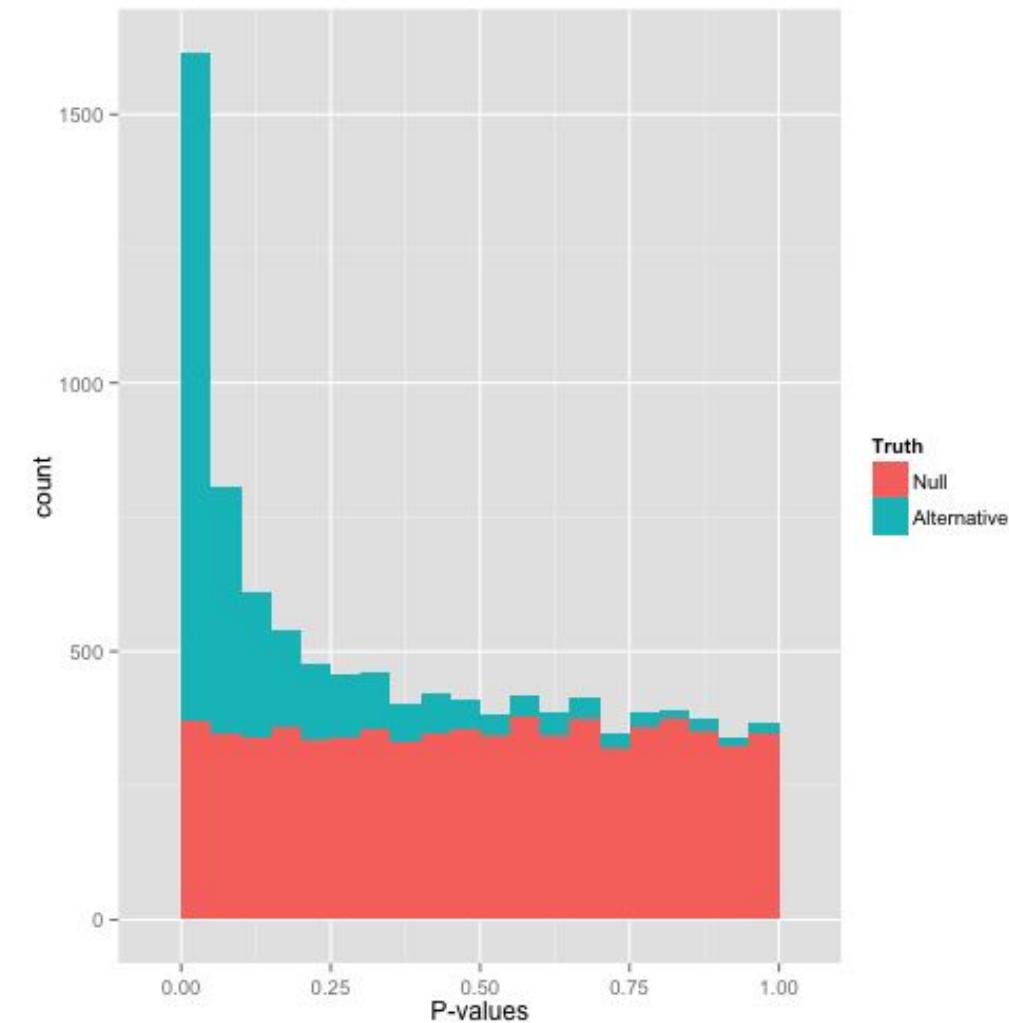
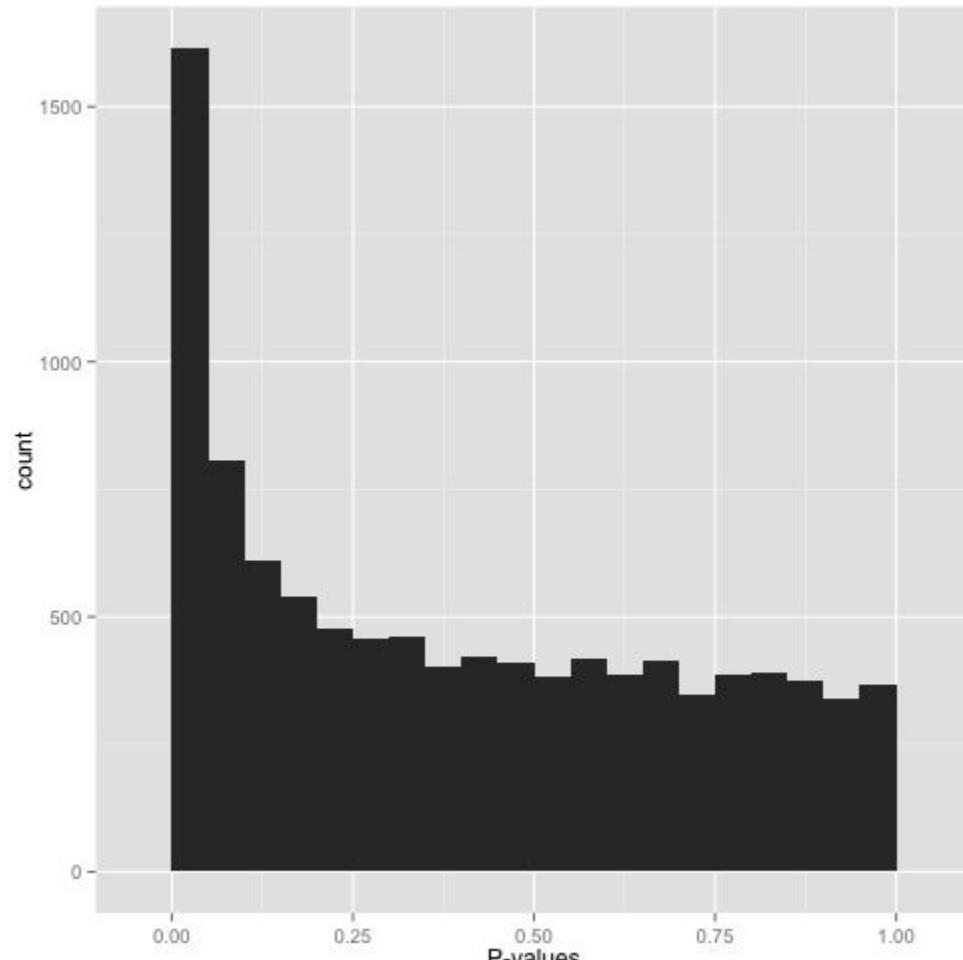
Correction for population
structure/relatedness

Genome Wide Scans of Association (GWAS)

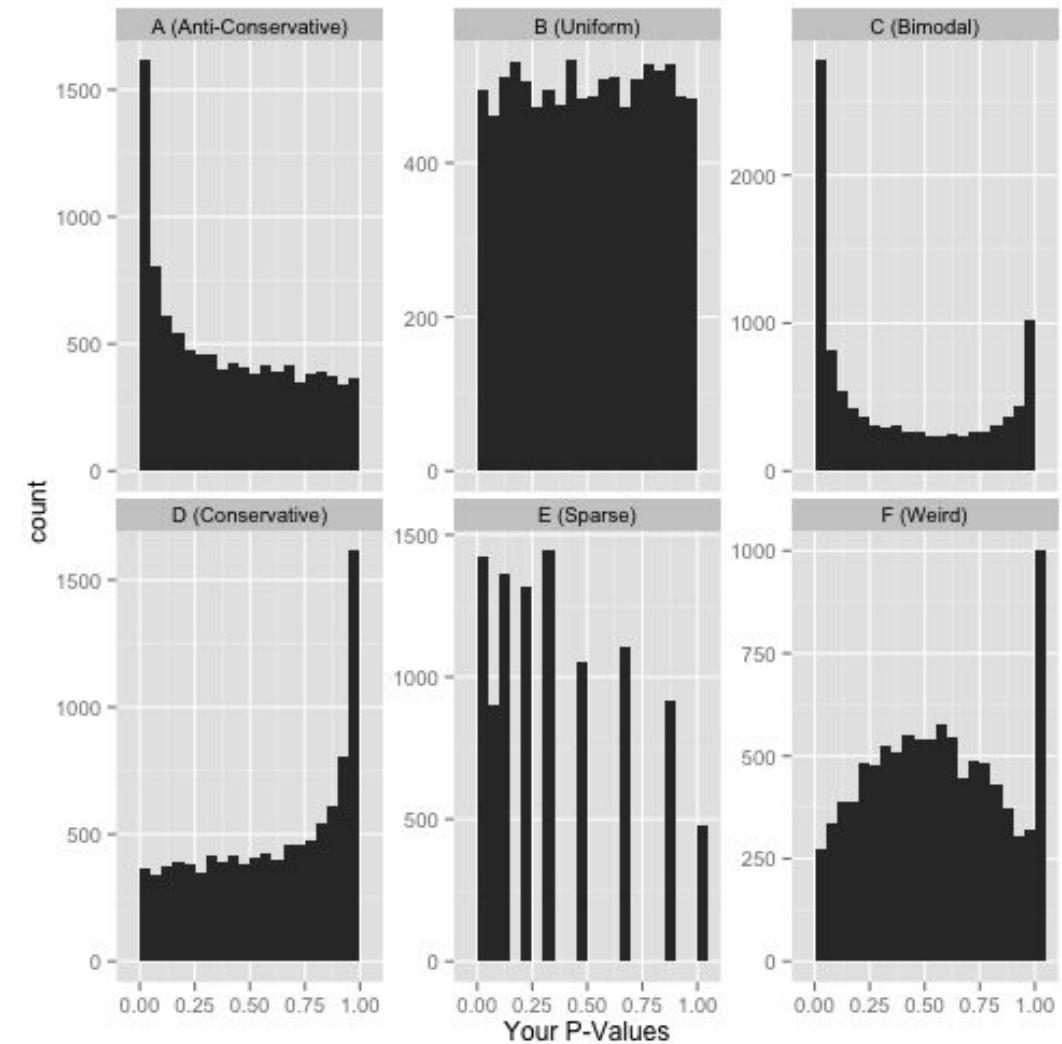
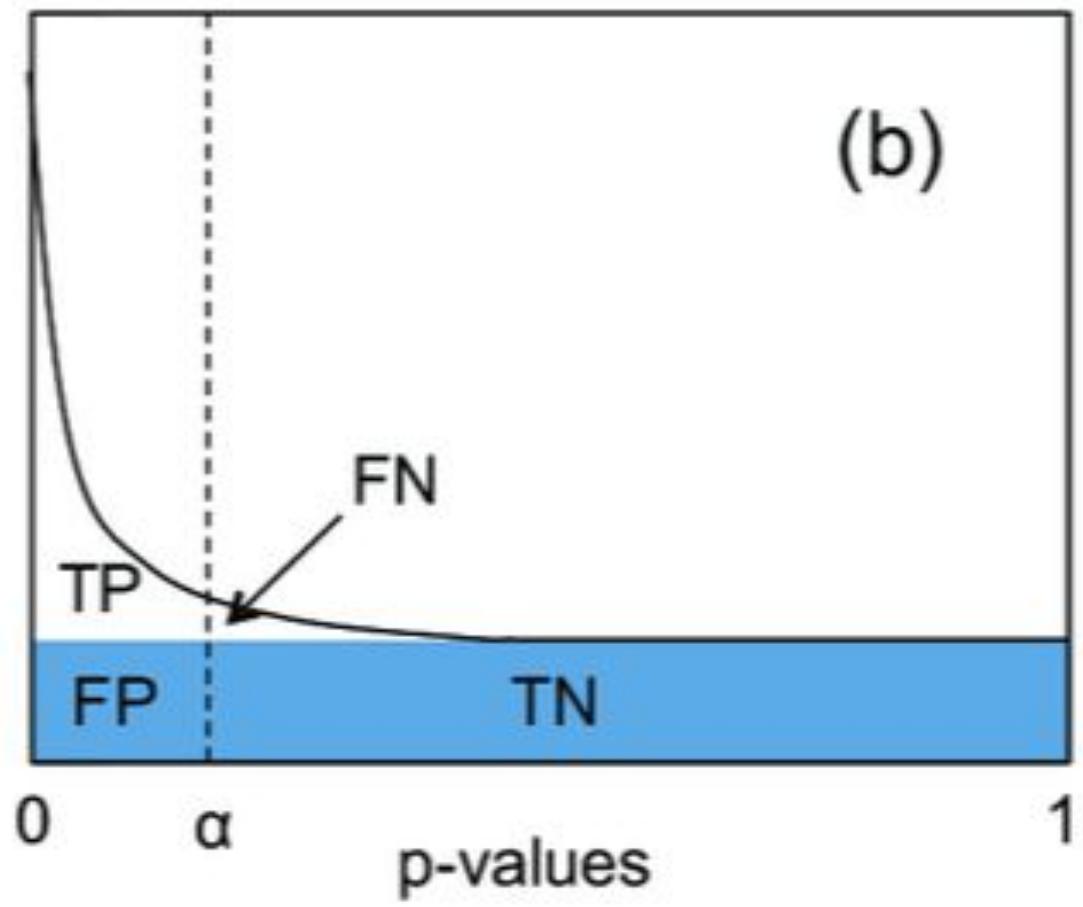
- Here a scan for association with rheumatoid arthritis
- Cohorts with 100,000 or even million individuals
- Useful to a point
- For polygenic traits: need high sample size



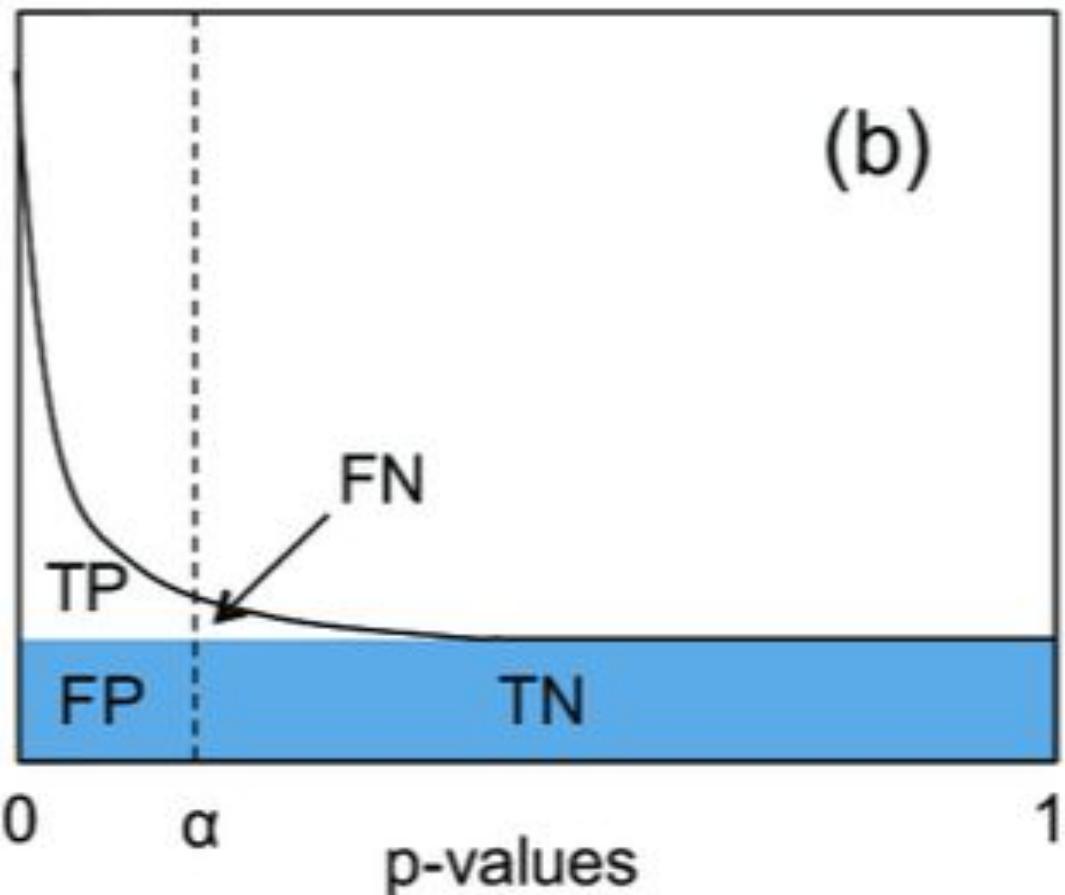
Always check your p-value distribution (not just for GWAS)



Always check your p-value distribution (not just for GWAS)

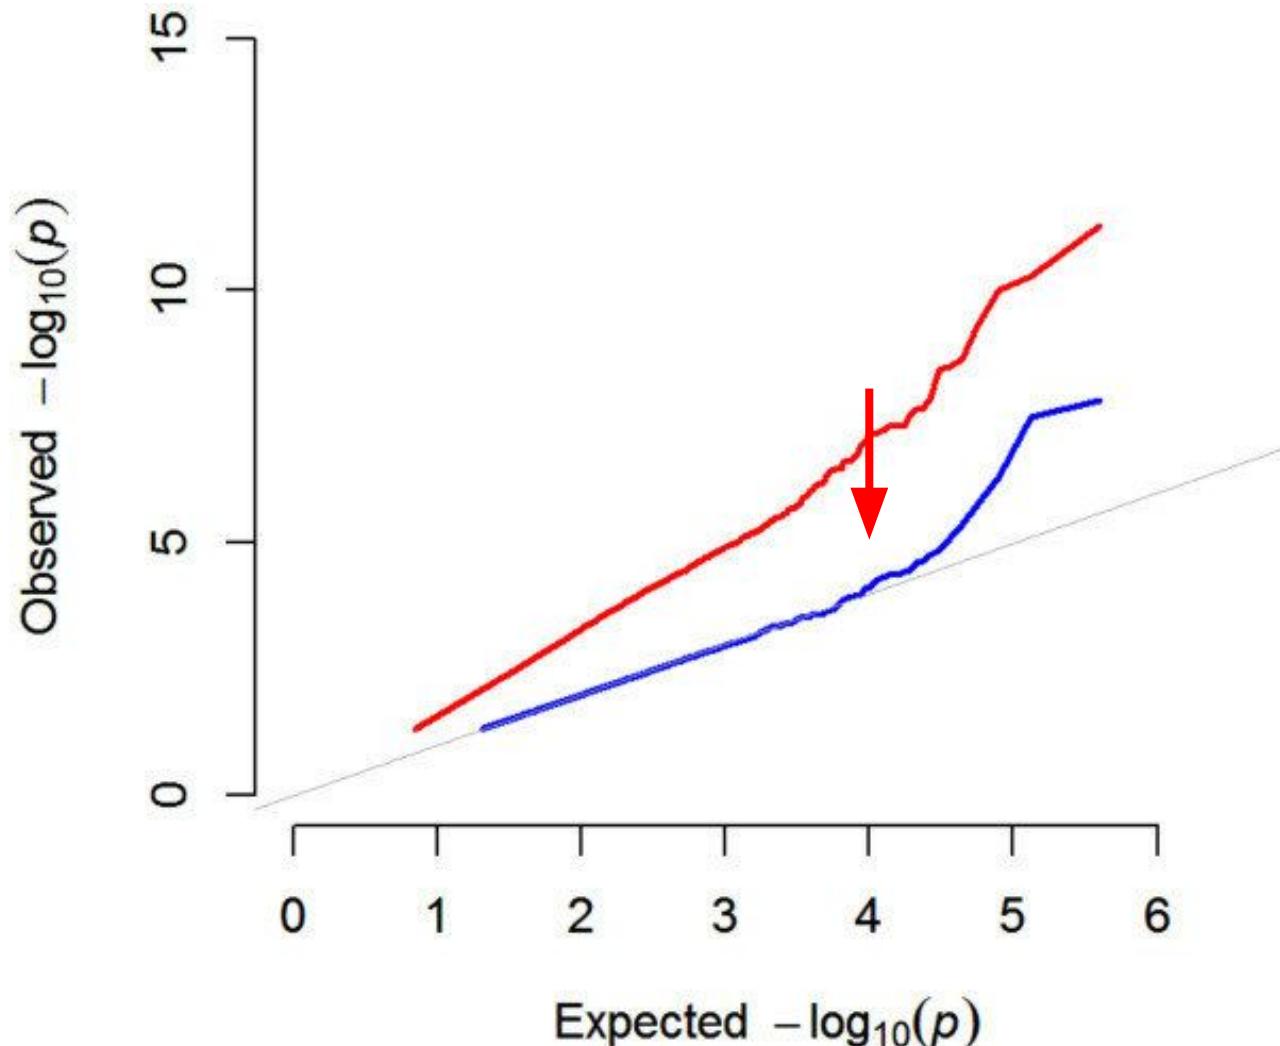


Always check your p-value distribution (not just for GWAS)



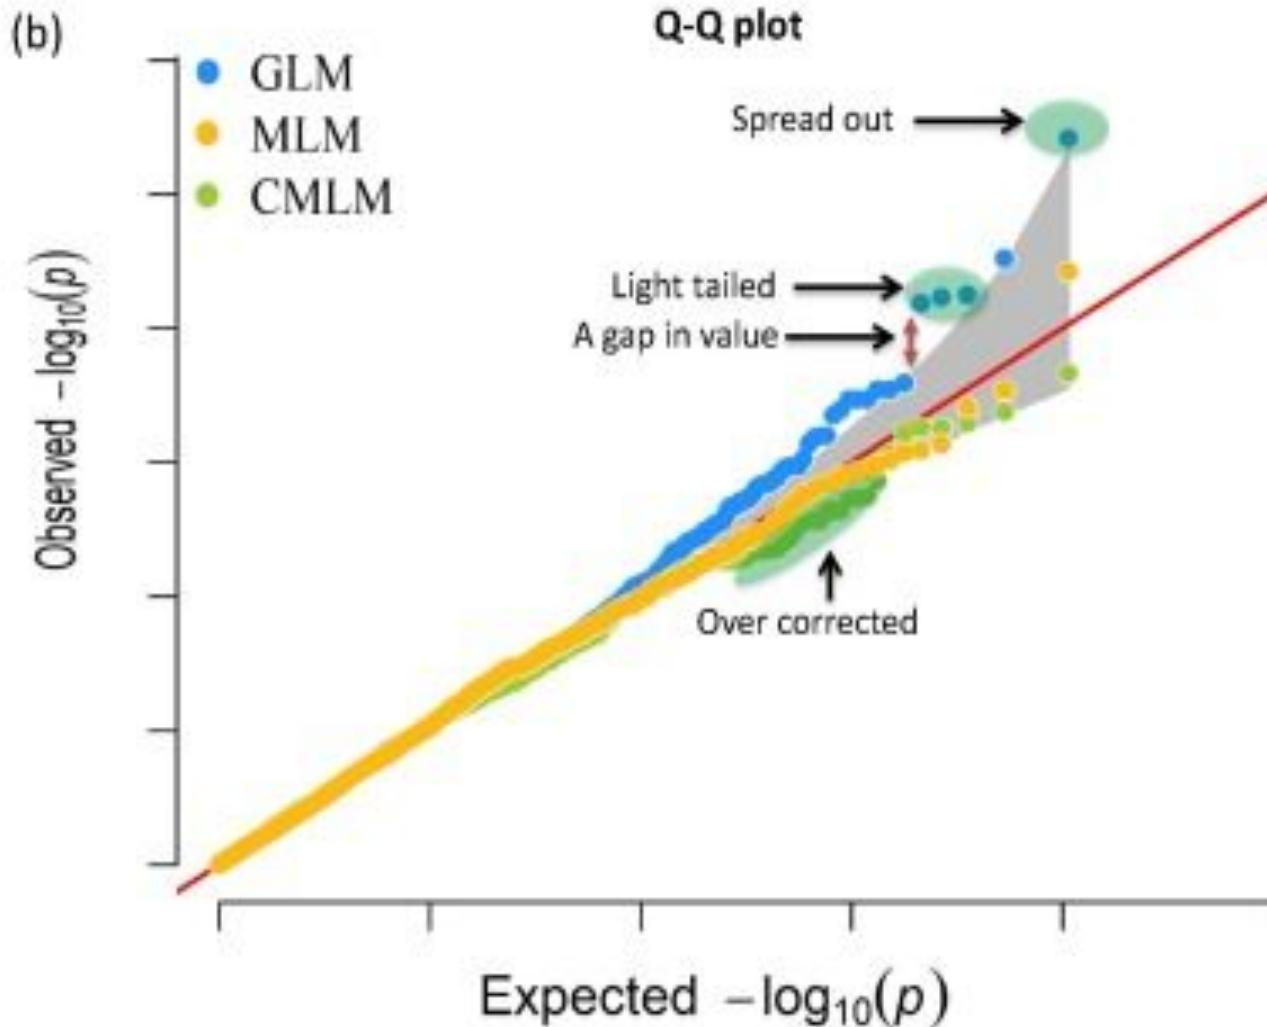
- You can apply a ‘brutal’ correction for multiple testing
- Bonferroni correction: multiply your p -value by the number of tests
- More subtle methods: q -value, the proportion of false positives for that particular p -value
- Local False Discovery Rate (FDR): the probability that H_0 is actually true given this particular p -value.
- In R package qvalue
- Requires the distribution of p -values to be uniform.

Correcting for p-value inflation

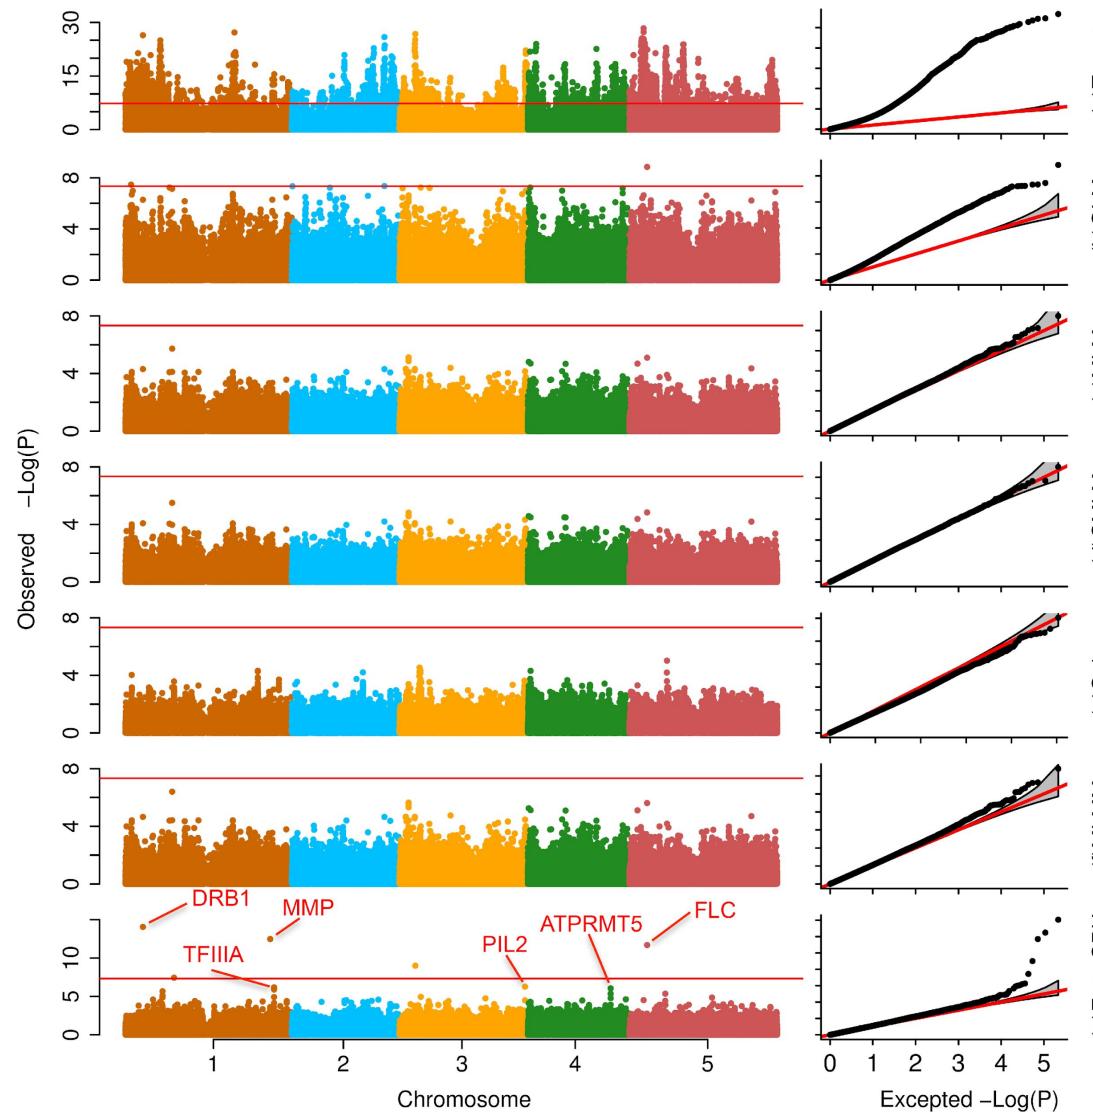


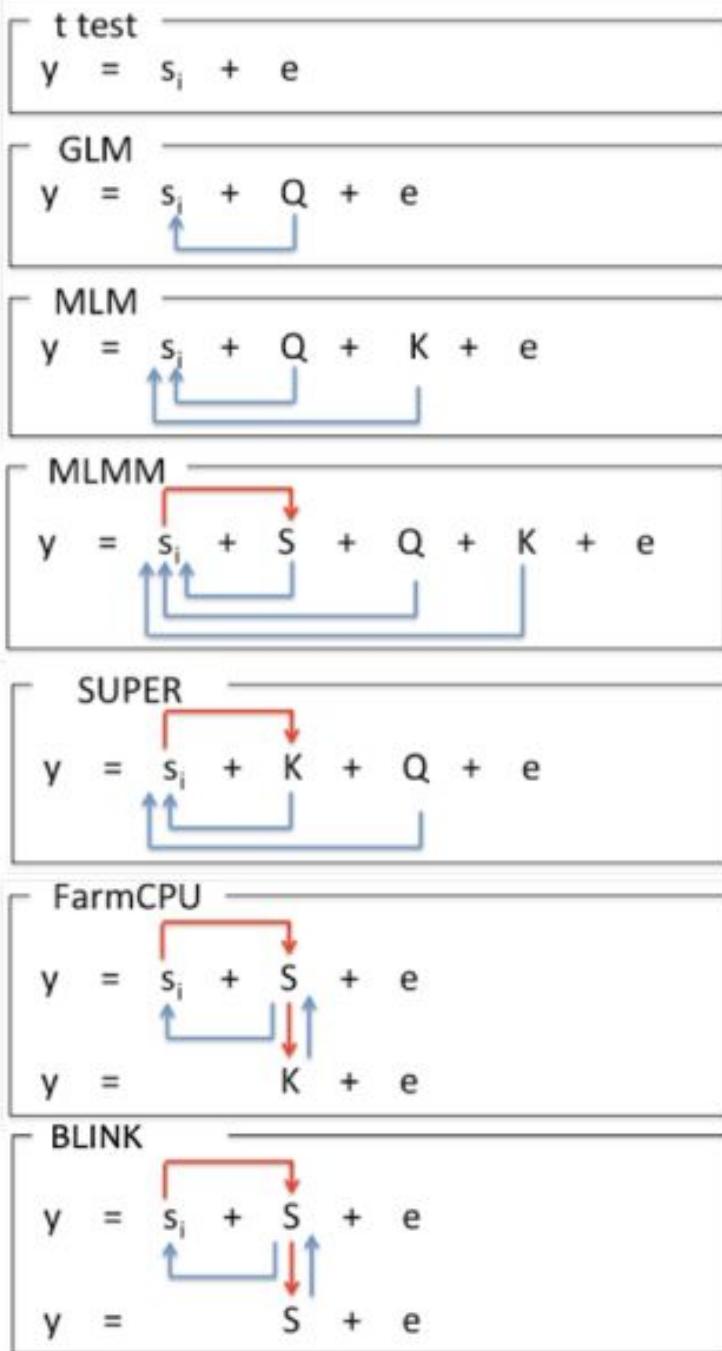
- The expectation is one drawn assuming a χ^2 distribution
- The ratio between observed and expected χ^2 values gives you a correction factor, called λ .
- Genomic control

Always check your p-value distribution (not just for GWAS)



A few GWAS models





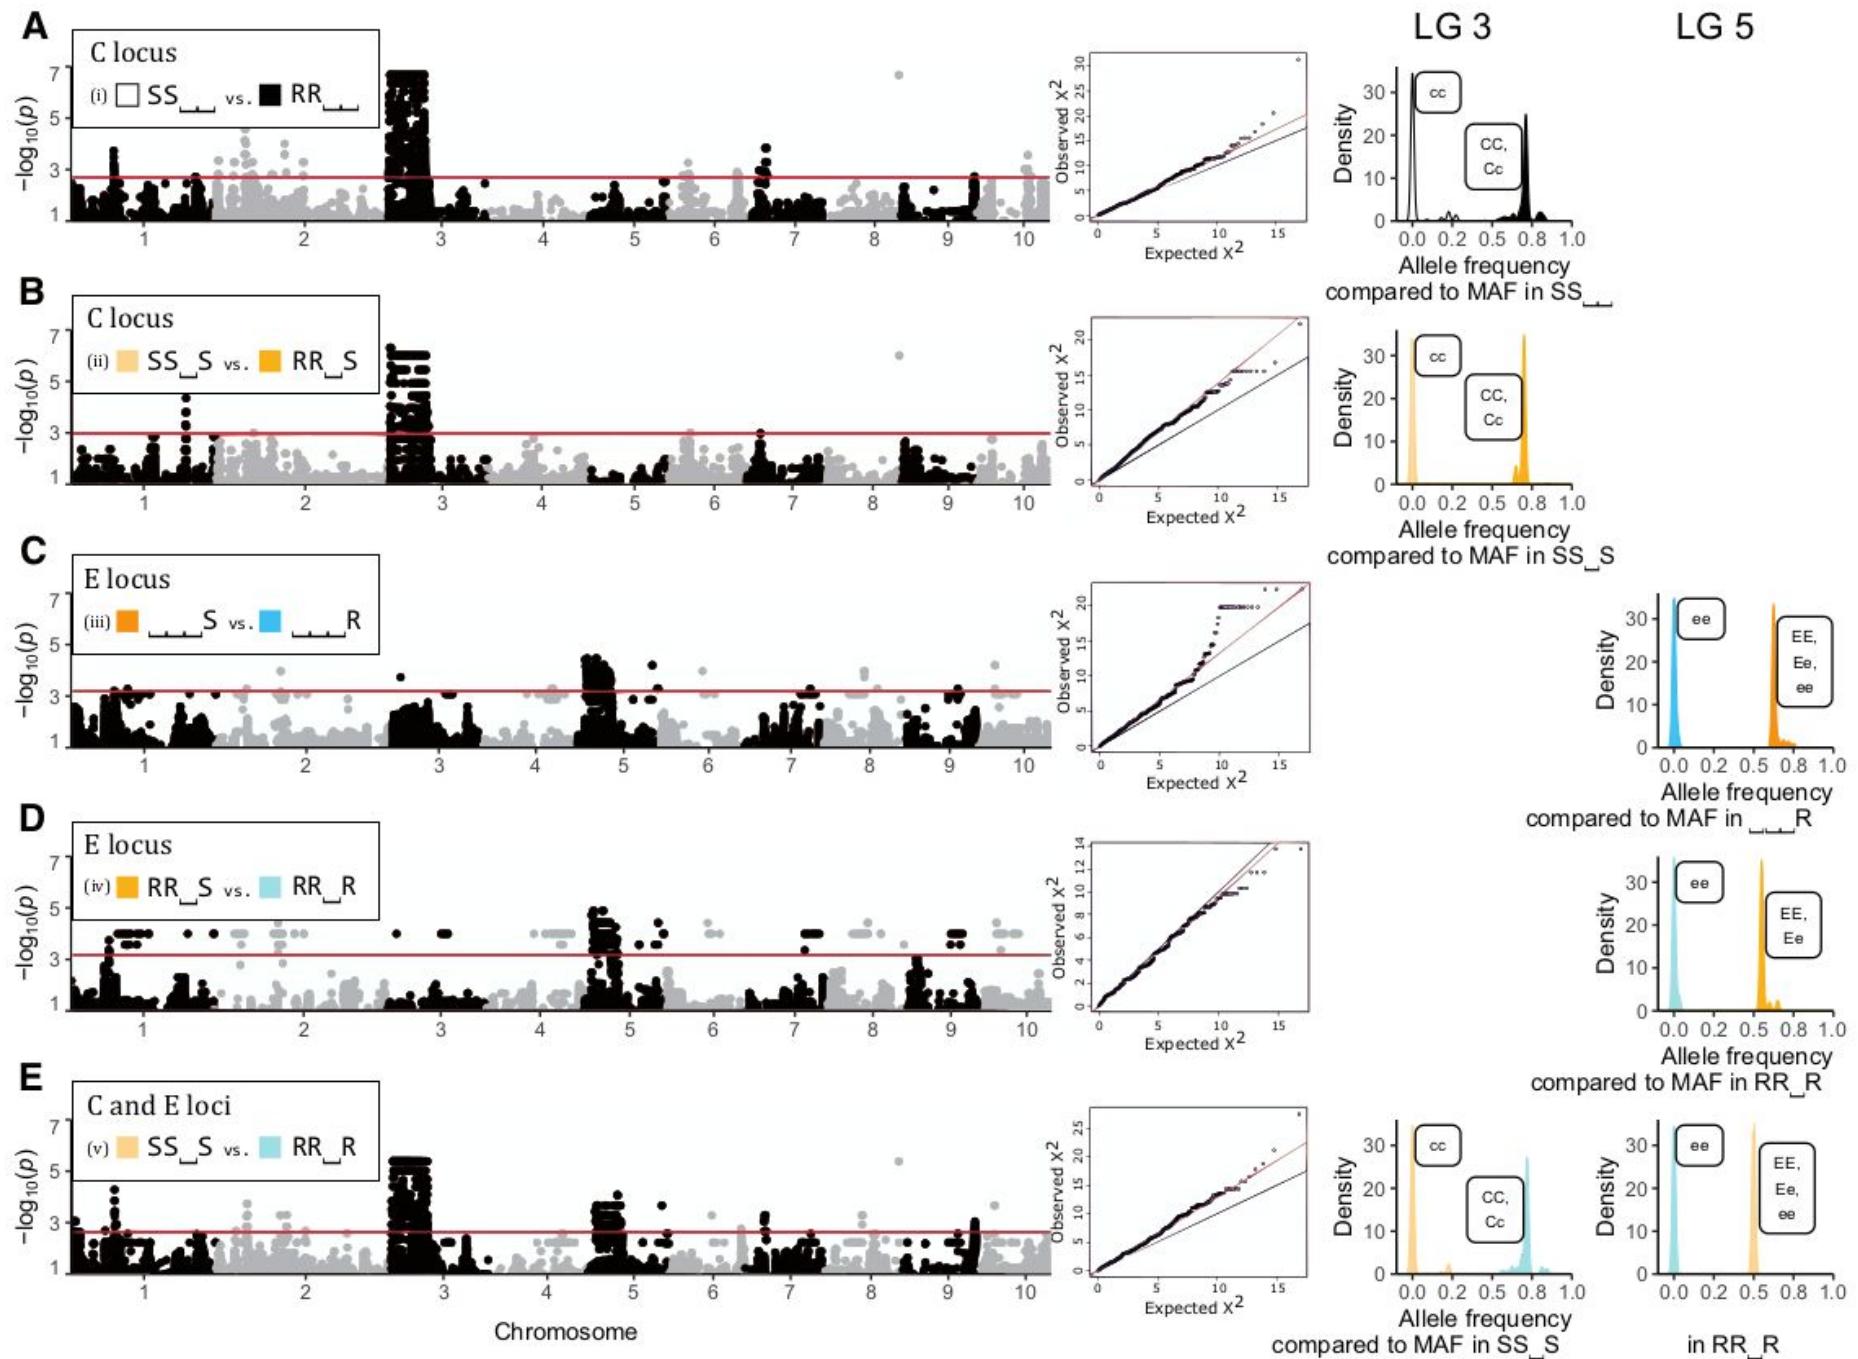
S_i : Testing marker Q : Population structure K : Kinship

S : Pseudo QTNs Q : Population structure Arrow: adjustment

Abstract

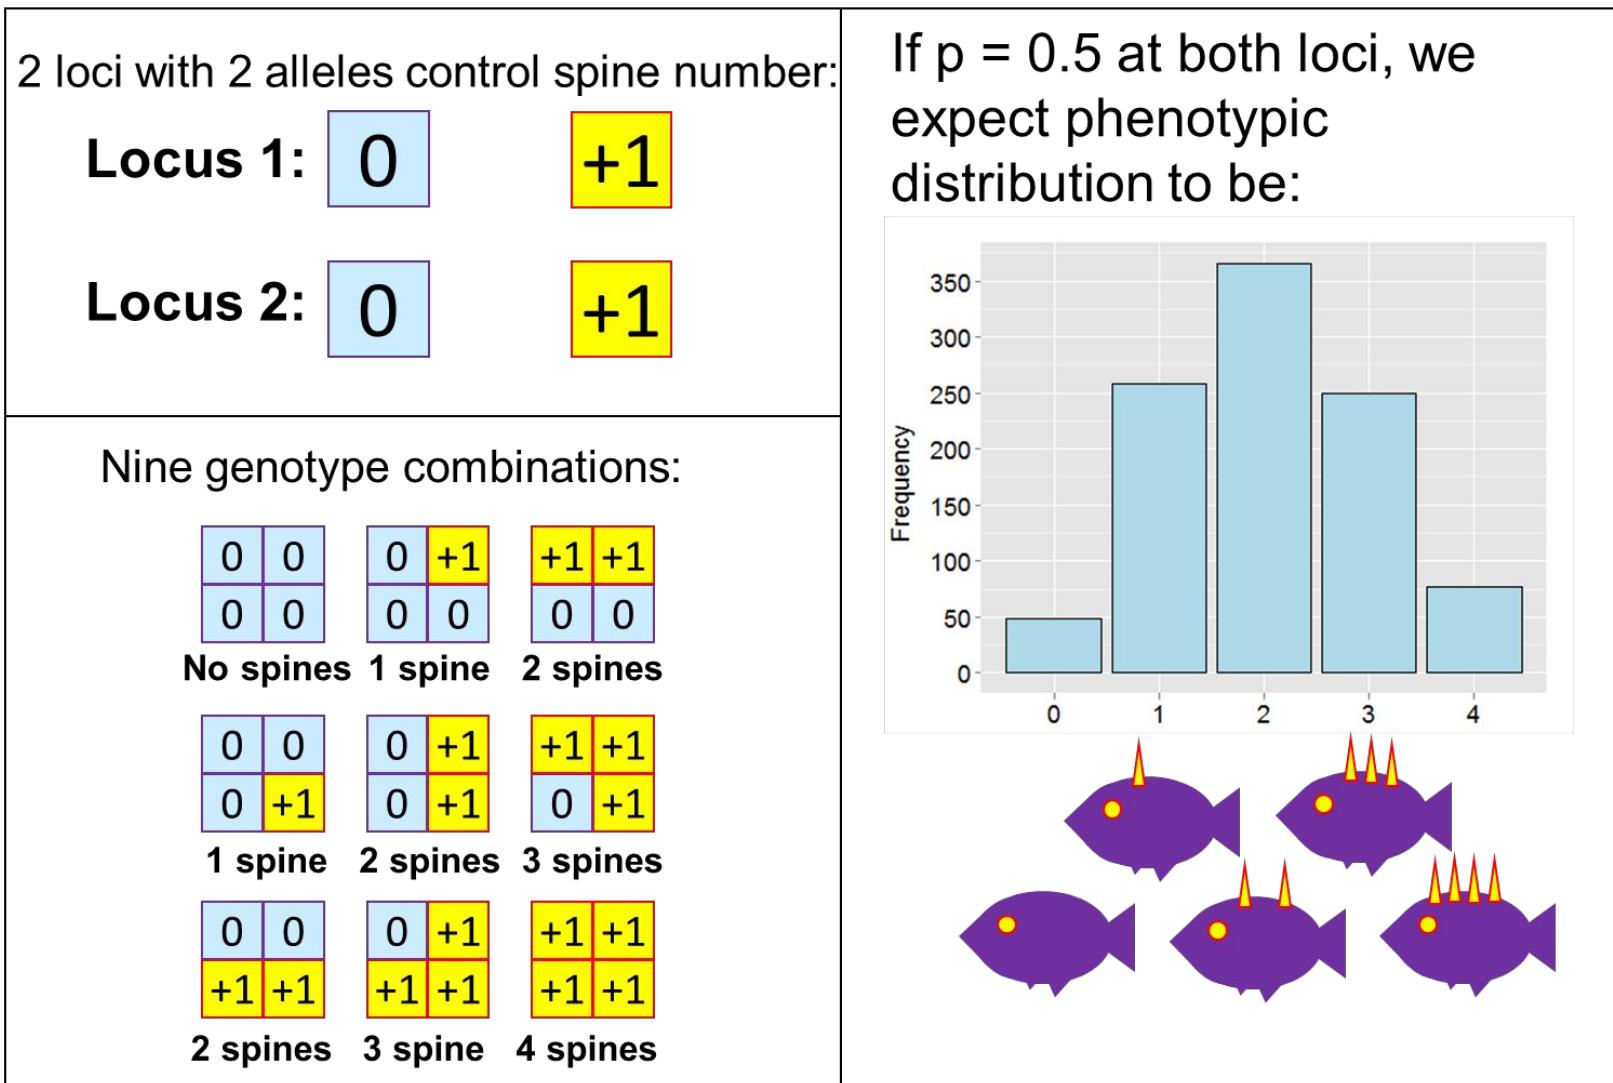
False positives in a Genome-Wide Association Study (GWAS) can be effectively controlled by a fixed effect and random effect Mixed Linear Model (MLM) that incorporates population structure and kinship among individuals to adjust association tests on markers; however, the adjustment also compromises true positives. The modified MLM method, Multiple Loci Linear Mixed Model (MLMM), incorporates multiple markers simultaneously as covariates in a stepwise MLM to partially remove the confounding between testing markers and kinship. To completely eliminate the confounding, we divided MLMM into two parts: Fixed Effect Model (FEM) and a Random Effect Model (REM) and use them iteratively. FEM contains testing markers, one at a time, and multiple associated markers as covariates to control false positives. To avoid model over-fitting problem in FEM, the associated markers are estimated in REM by using them to define kinship. The P values of testing markers and the associated markers are unified at each iteration. We named the new method as Fixed and random model Circulating Probability Unification (FarmCPU). Both real and simulated data analyses demonstrated that FarmCPU improves statistical power compared to current methods. Additional benefits include an efficient computing time that is linear to both number of individuals and number of markers. Now, a dataset with half million individuals and half million markers can be analyzed within three days.

The second strategy changes the definition of kinship among individuals. Only the associated genetic markers are used as pseudo Quantitative Trait Nucleotides (QTNs) to derive kinship instead of all, or a random sample of genetic markers. Pseudo QTNs are expected to closely track some of the causative QTNs, and are selectively used to derive kinship for a specific testing marker. Whenever a pseudo QTN is correlated with the testing marker, it is excluded from those used to derive kinship. In the FaST-LMM-Select method, a pseudo QTN is considered correlated if it is within a 2Mb interval on either side of the testing marker[23]. Instead of using a 2Mb interval, the Settlement of MLM Under Progressively Exclusive Relationship (SUPER) method applies a threshold on Linkage Disequilibrium (LD) between the pseudo QTNs and the testing marker. Selectively including and/or excluding pseudo QTNs to derive kinship for a specific testing marker improves statistical power compared to deriving a overall kinship from all, or a random sample of genetic markers[24].

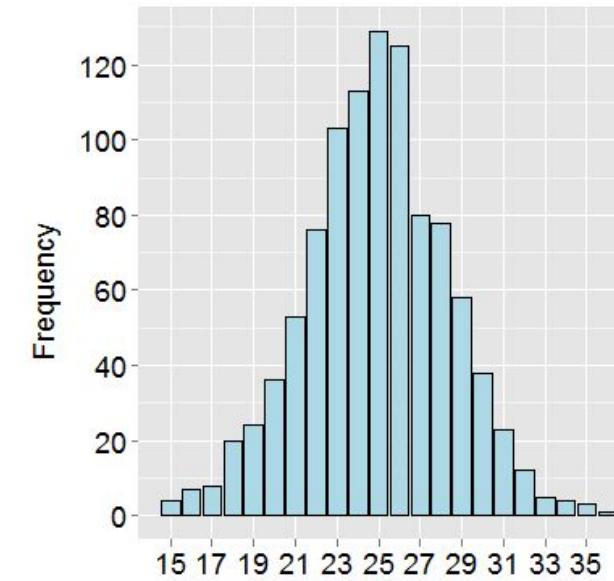
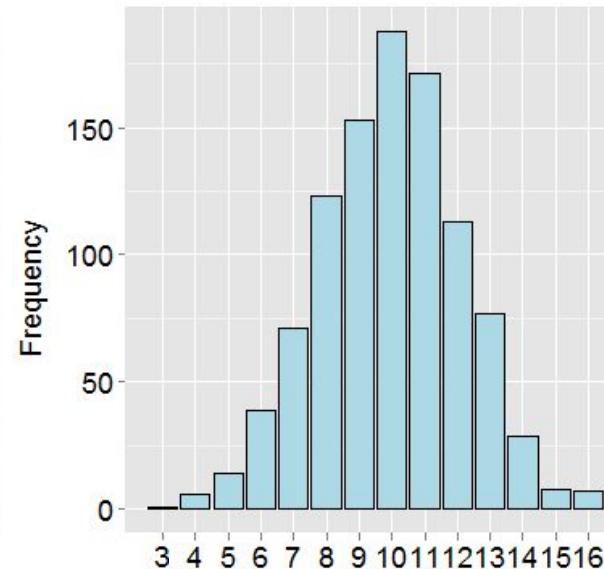
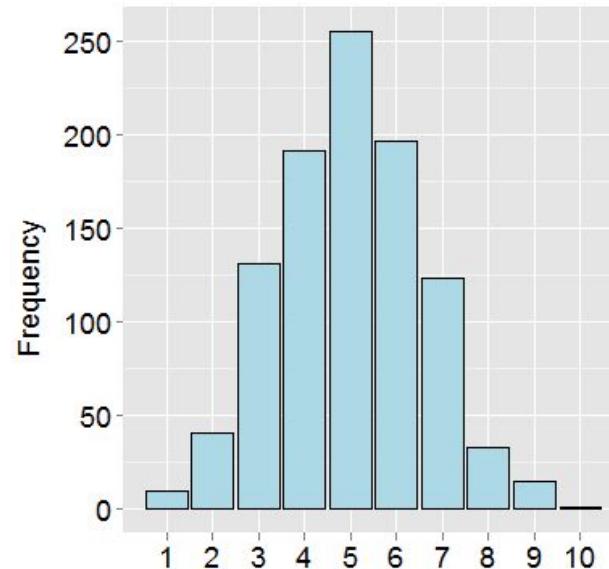


Back to quantitative genetics

- Polygenic inheritance
- Refers to quantitative characteristics controlled by cumulative effects of many genes
- Each locus still follows Mendel's rules
- May be influenced by environmental factors
- Kind of works, and simplifies things...



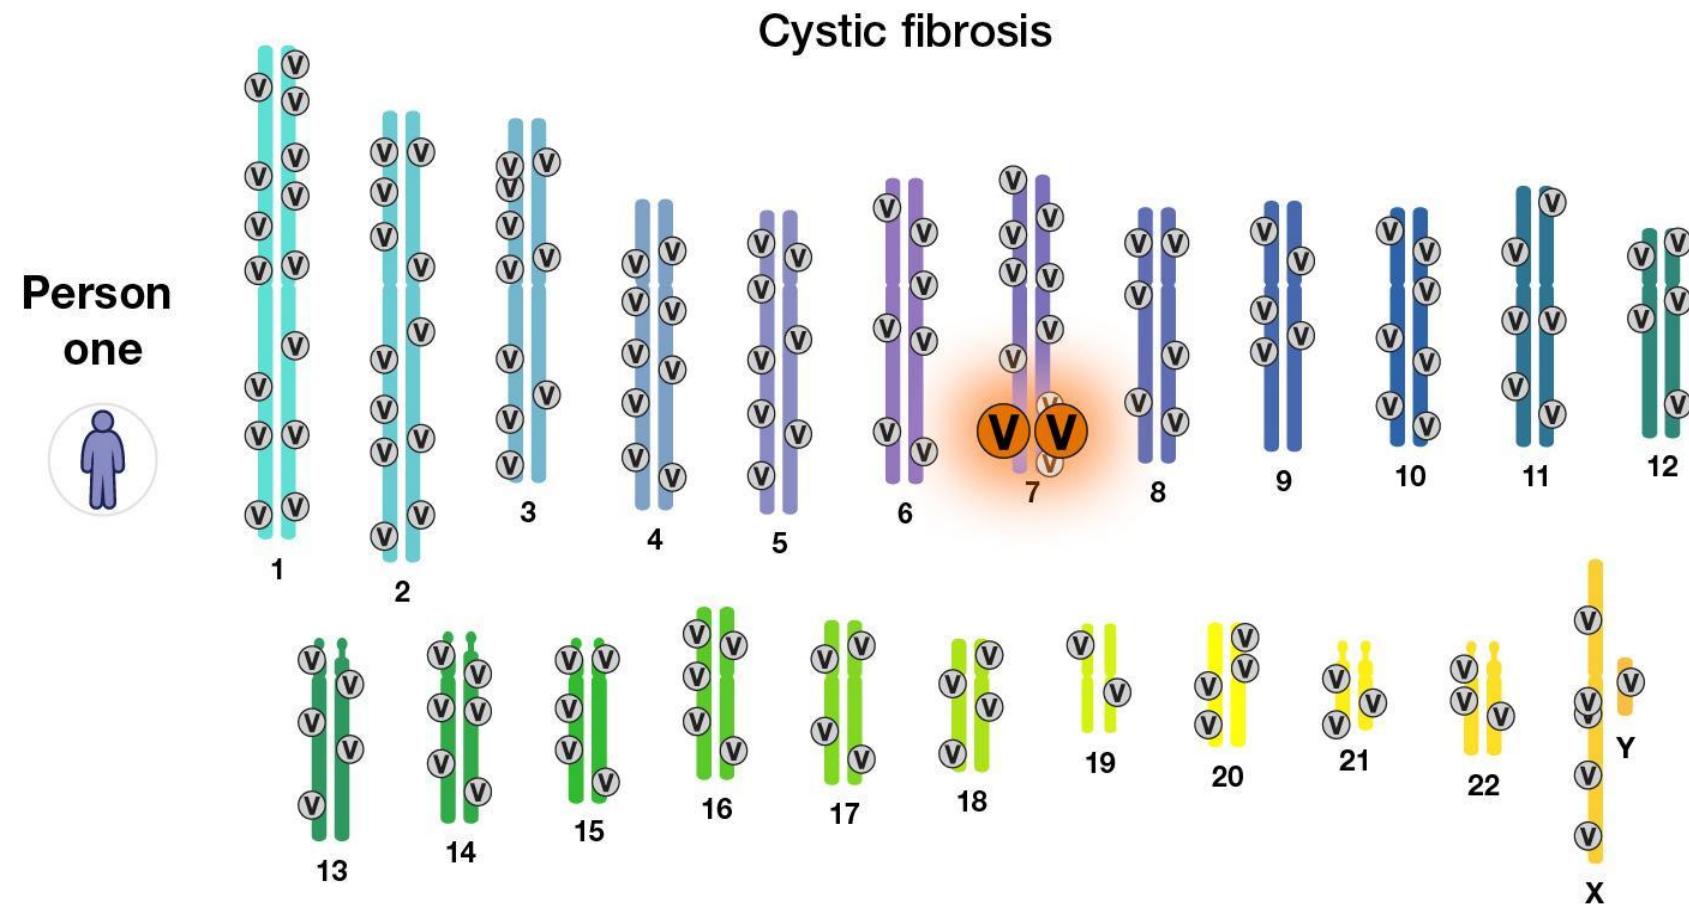
Back to quantitative genetics



- As more loci control the trait, it becomes difficult to describe the trait in terms of individual loci
- Central Limit theorem: the combination of random draws converges to a normal distribution.

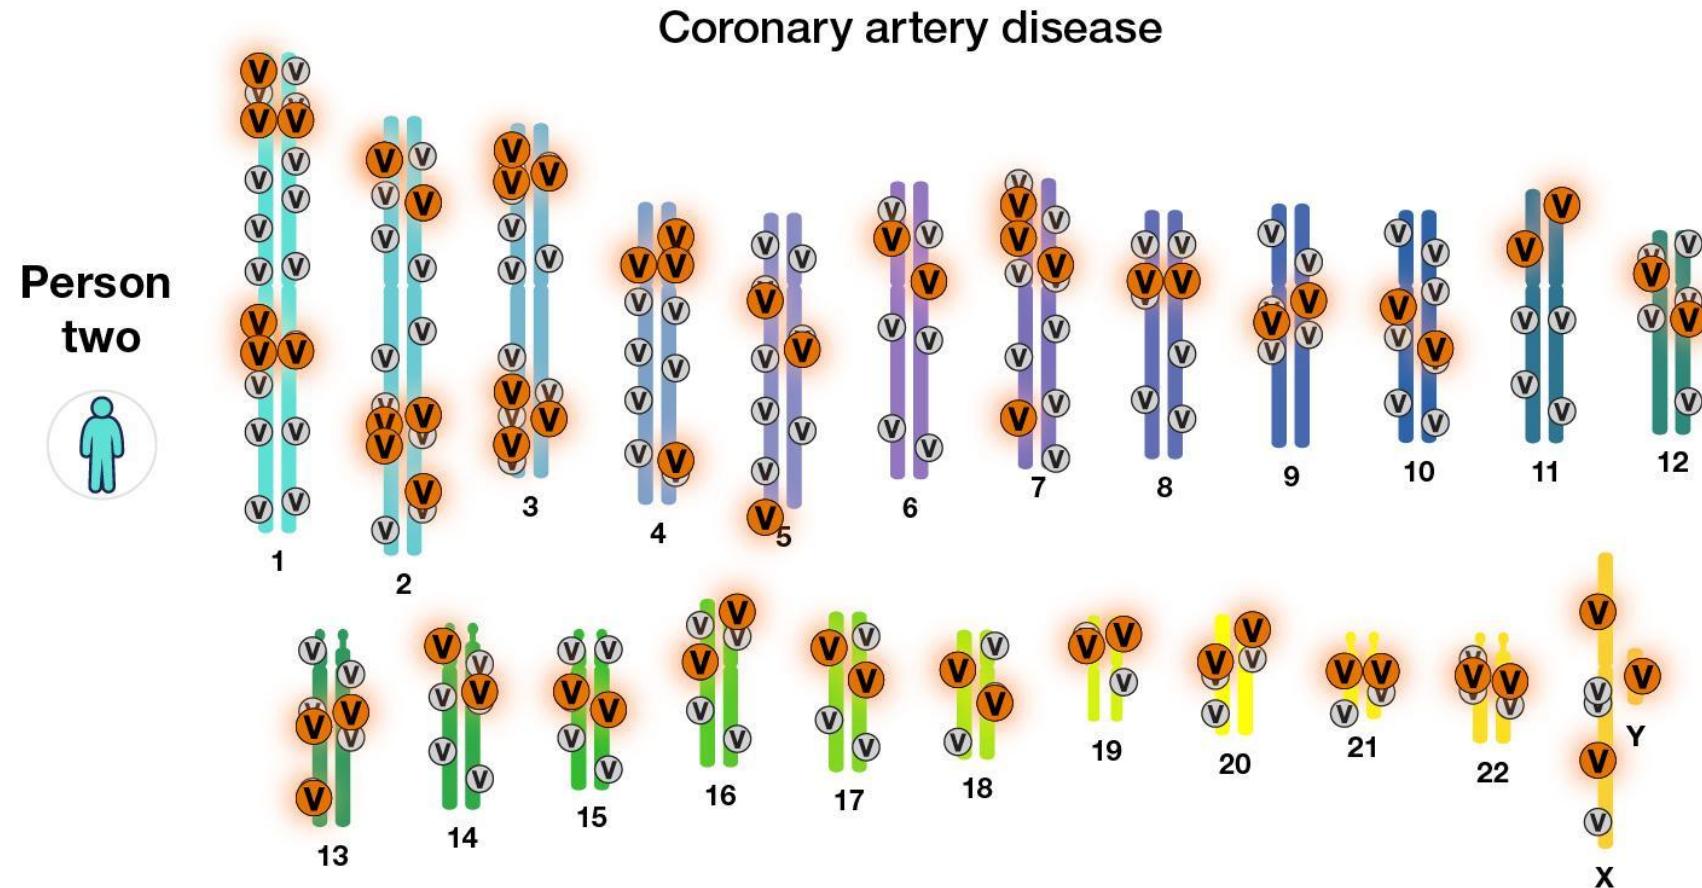
Polygenic risk scores

- We assume additive effects
- The risk is therefore additive.
- Neglects epistasy.
- Using genotypes at associated loci to predict an outcome
- Here with a single locus



Polygenic risk scores

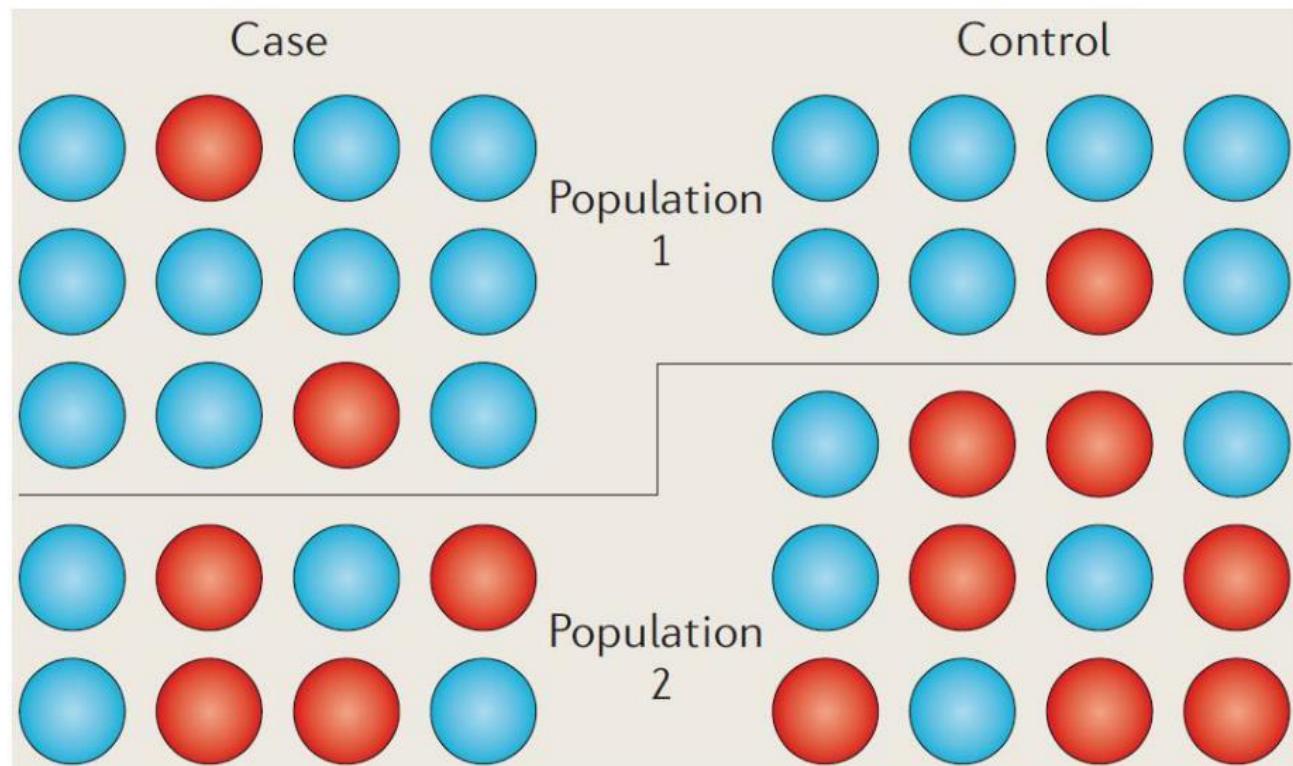
- Using genotypes at associated loci to predict an outcome
- Here with many loci



Genome Wide Scans of Association (GWAS): limitations and pitfalls

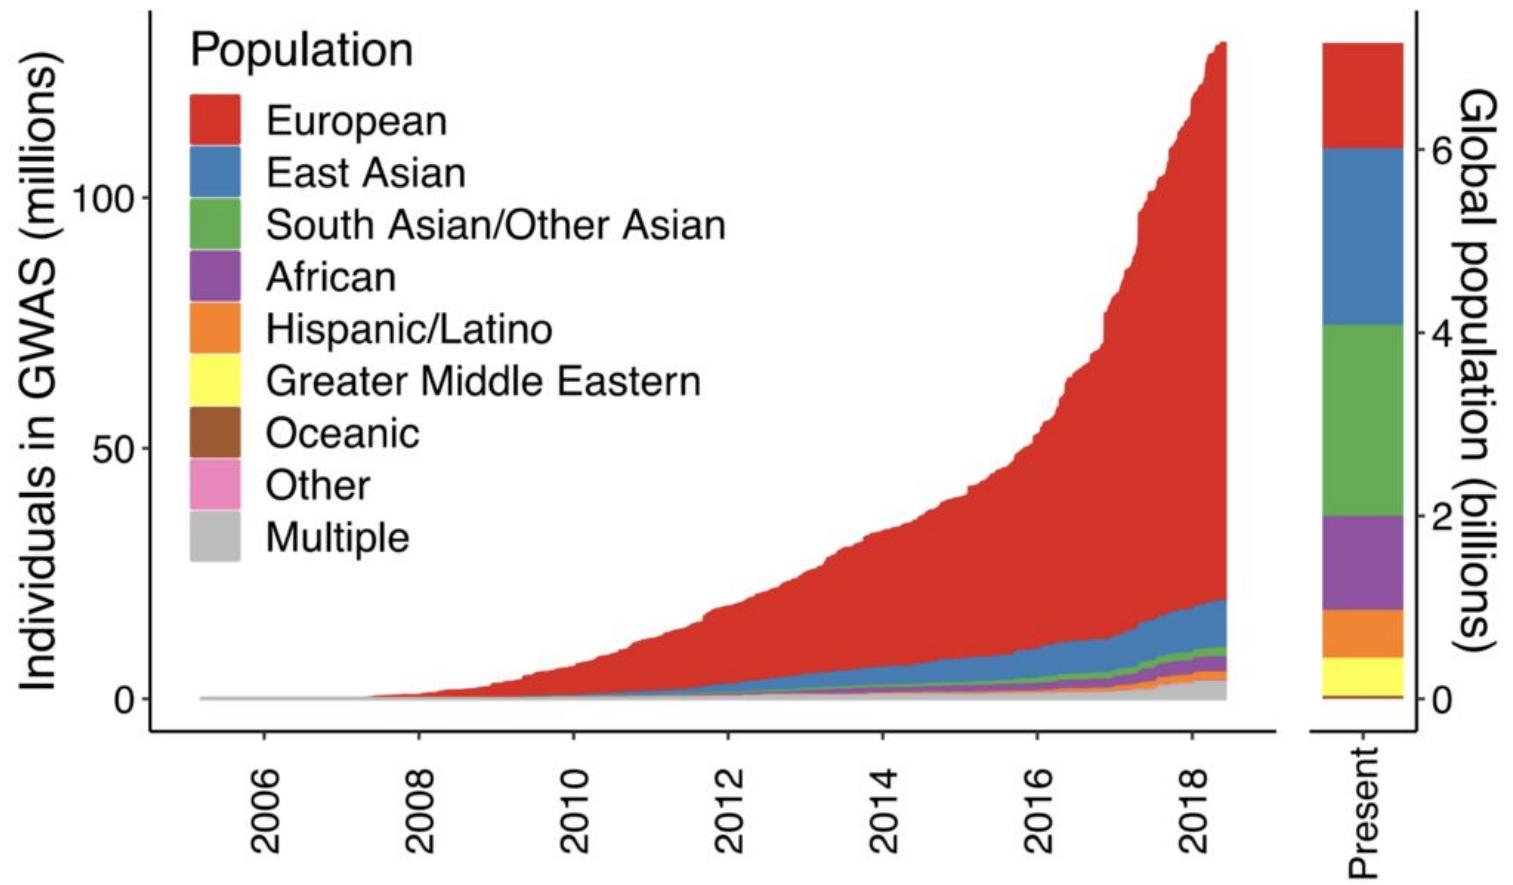
- Small effect sizes
- Biases due to population structure (stratification)
- Ethnic bias
- Story-telling
- Look at the effect size, not just the p-value

What is population stratification?

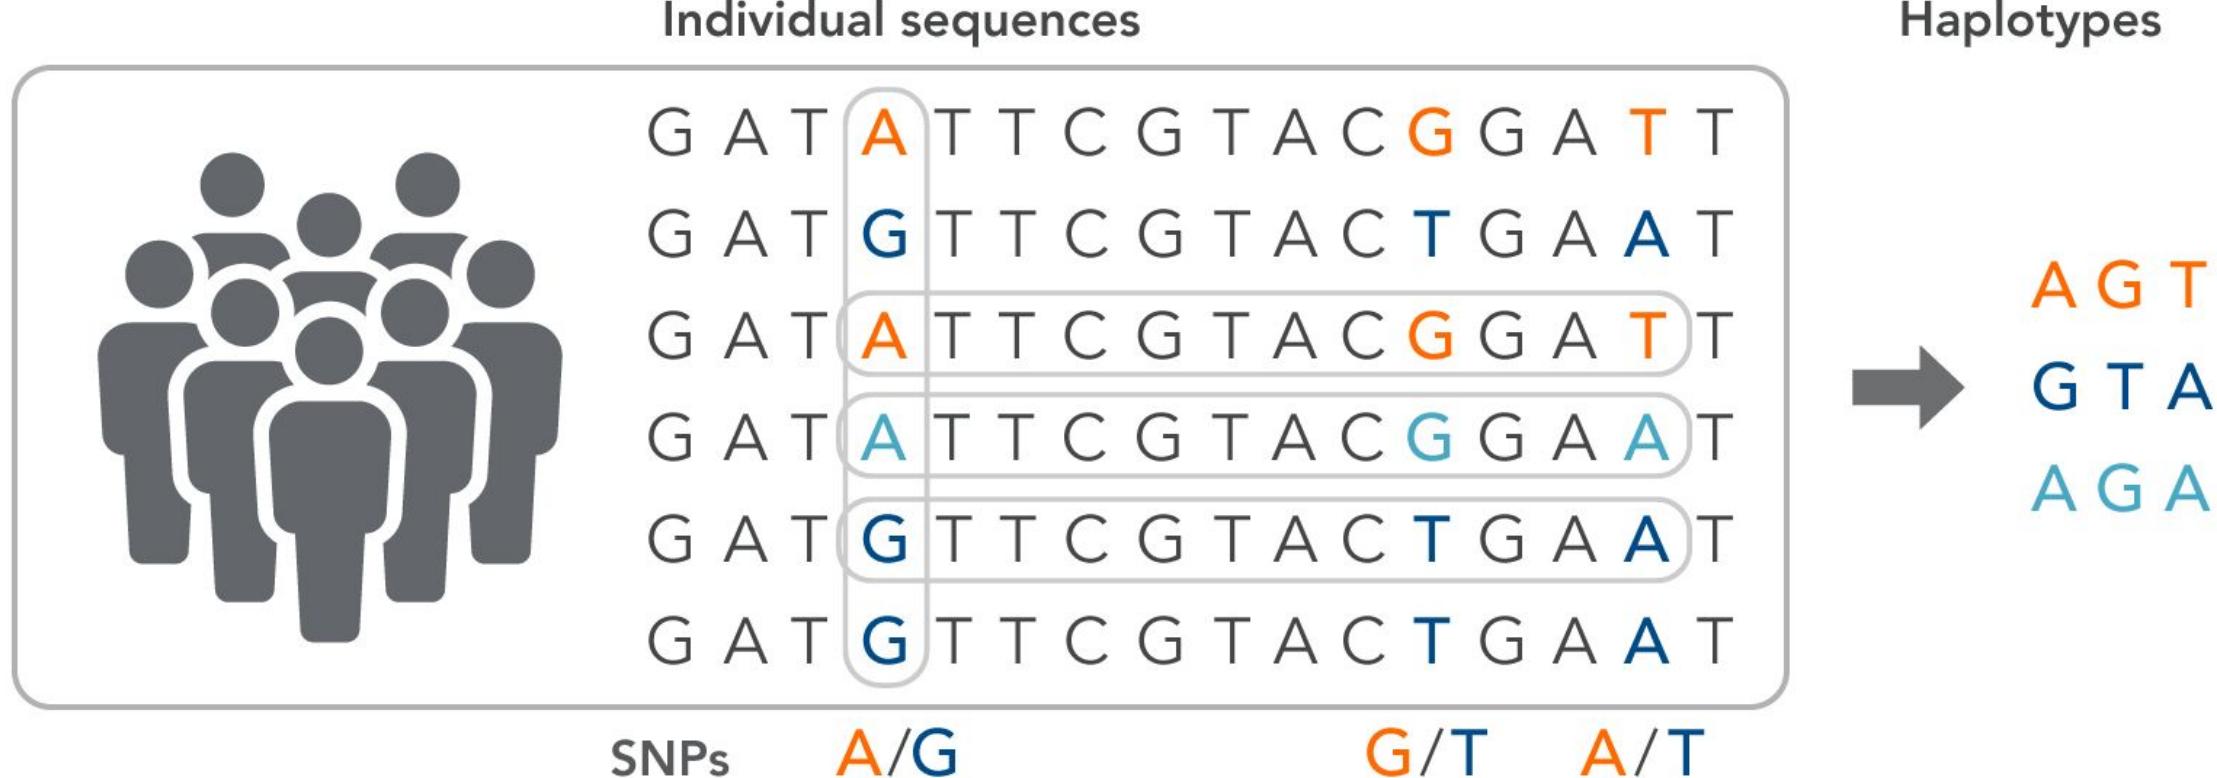


Genome Wide Scans of Association (GWAS): limitations and pitfalls

- Small effect sizes
- Biases due to population structure (stratification)
- Ethnic bias
- Story-telling
- Look at the effect size (β), not just the p-value.

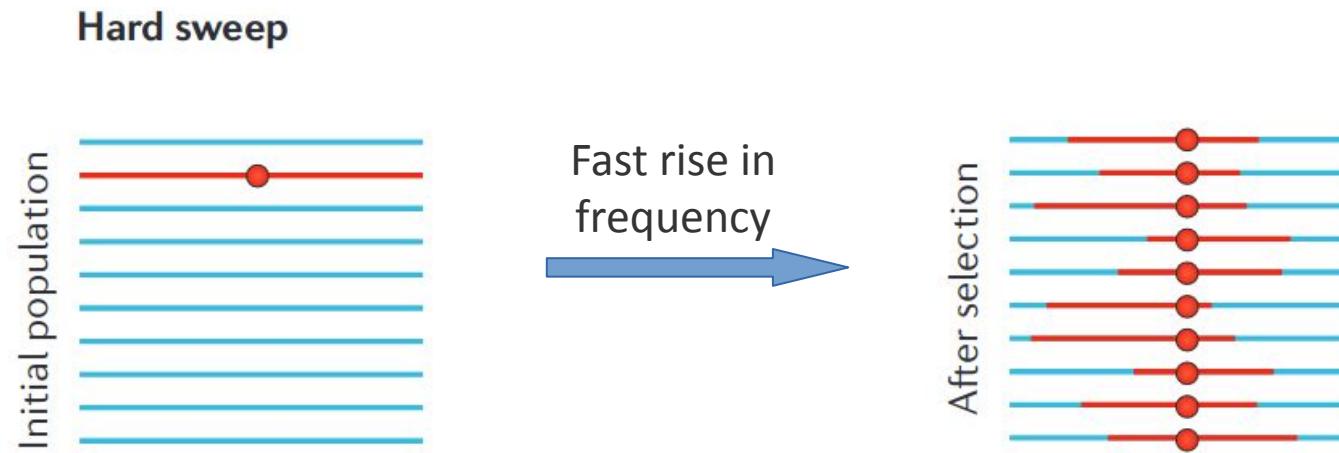


Scans for selection. Quick reminder

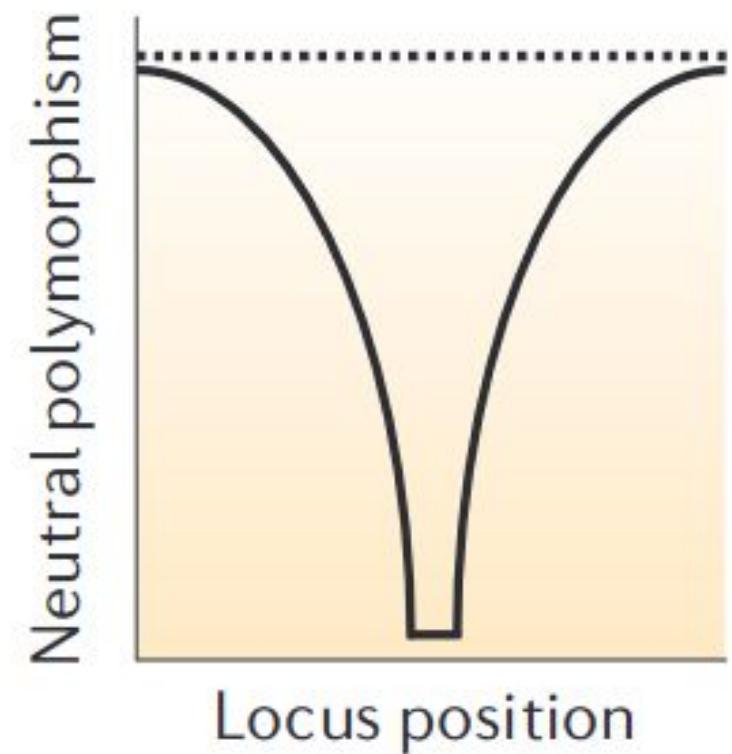
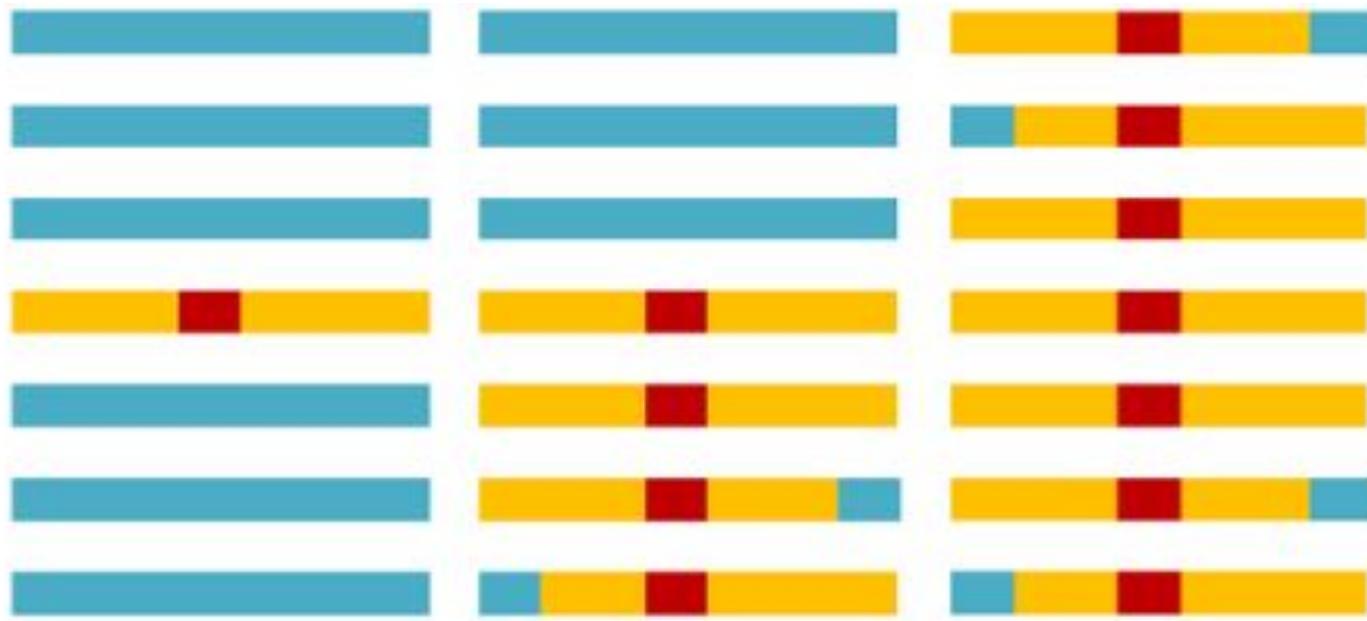


A (hard) selective sweep

- The association between alleles across individuals usually breaks down because of recombination
- Recent selection leads to a fast increase in frequency of the selected allele and those at the vicinity (selective sweep)
- One consequence: drop in ‘how different’ alleles are near the selected locus



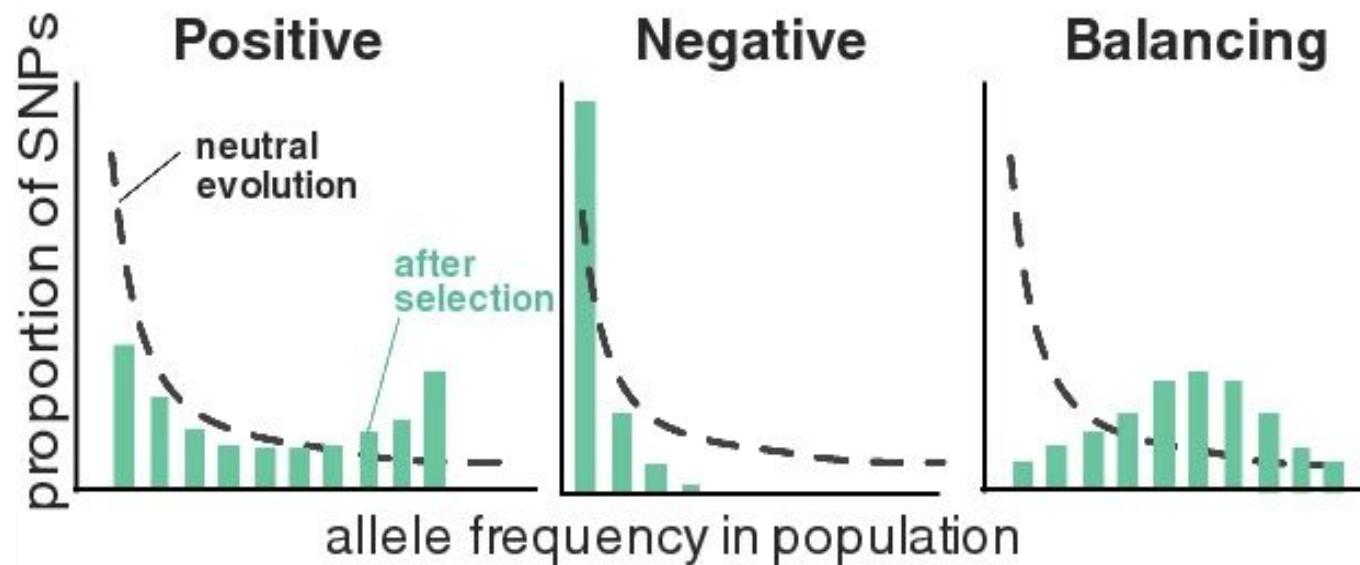
Effect on diversity statistics



Effect on the Allele Frequency Spectrum

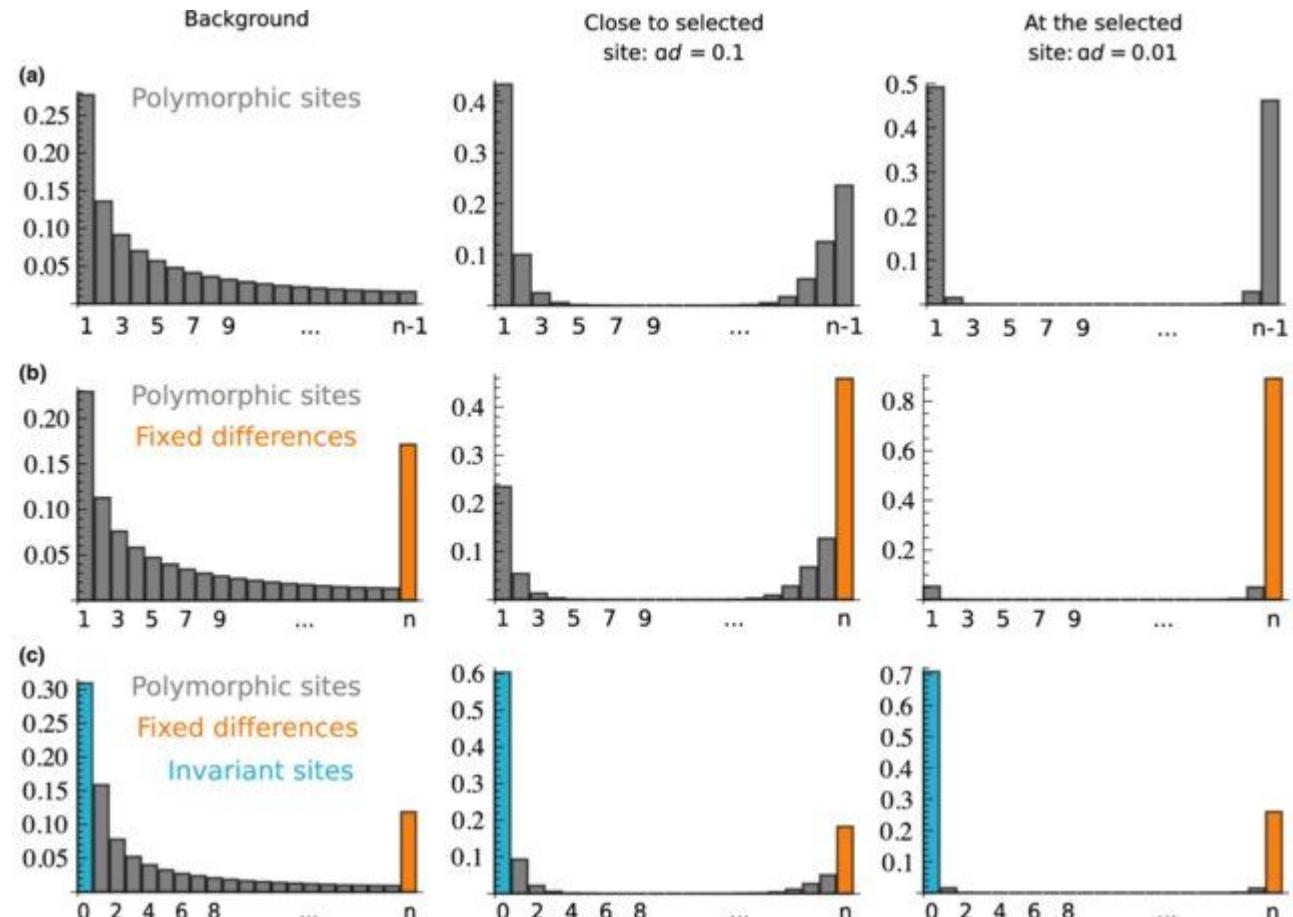
- A selective sweep leads to a change in the frequency of derived alleles
- This signature can be detected by contrasting local and global (or simulated) ‘neutral’ spectra
- Can be described by diversity statistics, such as Tajima’s D

allele frequency distribution



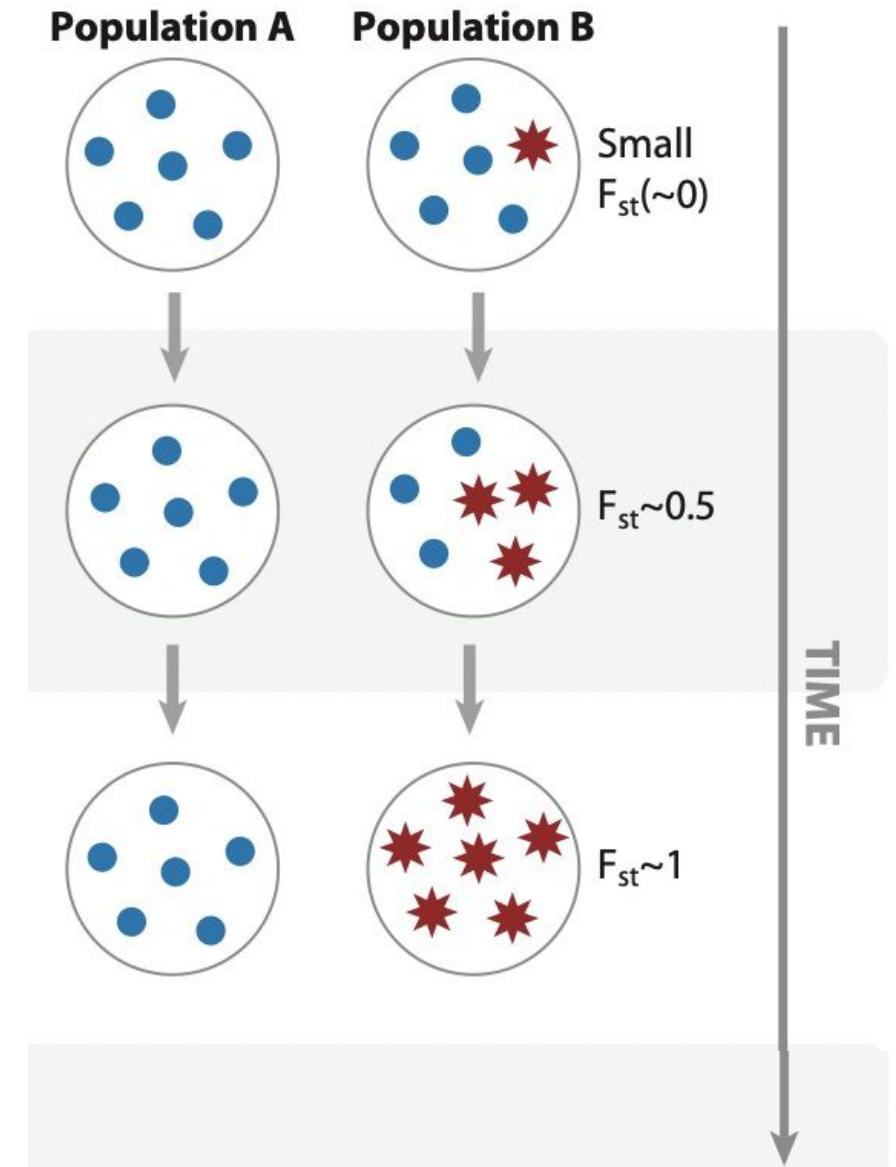
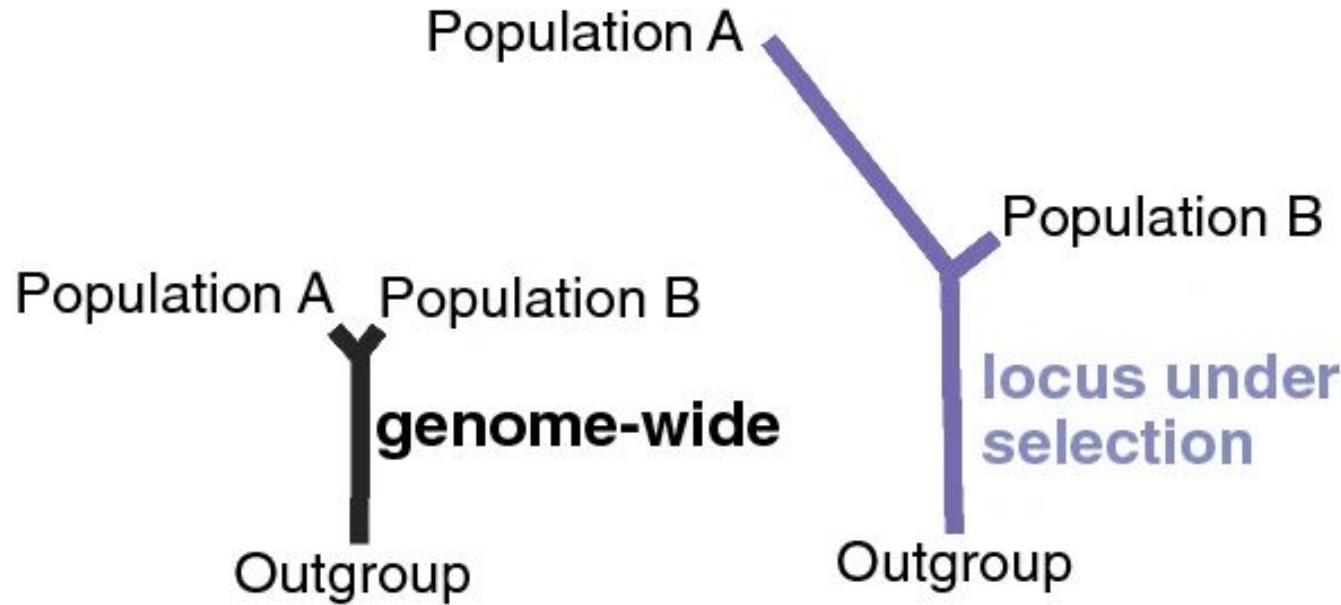
Effect on the Allele Frequency Spectrum

- A selective sweep leads to a change in the frequency of derived alleles
- This signature can be detected by contrasting local and global (or simulated) ‘neutral’ spectra
- Can be described by diversity statistics, such as Tajima’s D
- The entire spectrum can be used more efficiently: Composite Likelihood Ratios (CLR) estimated by SweepFinder/SweeD.

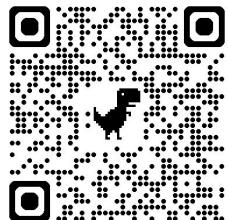
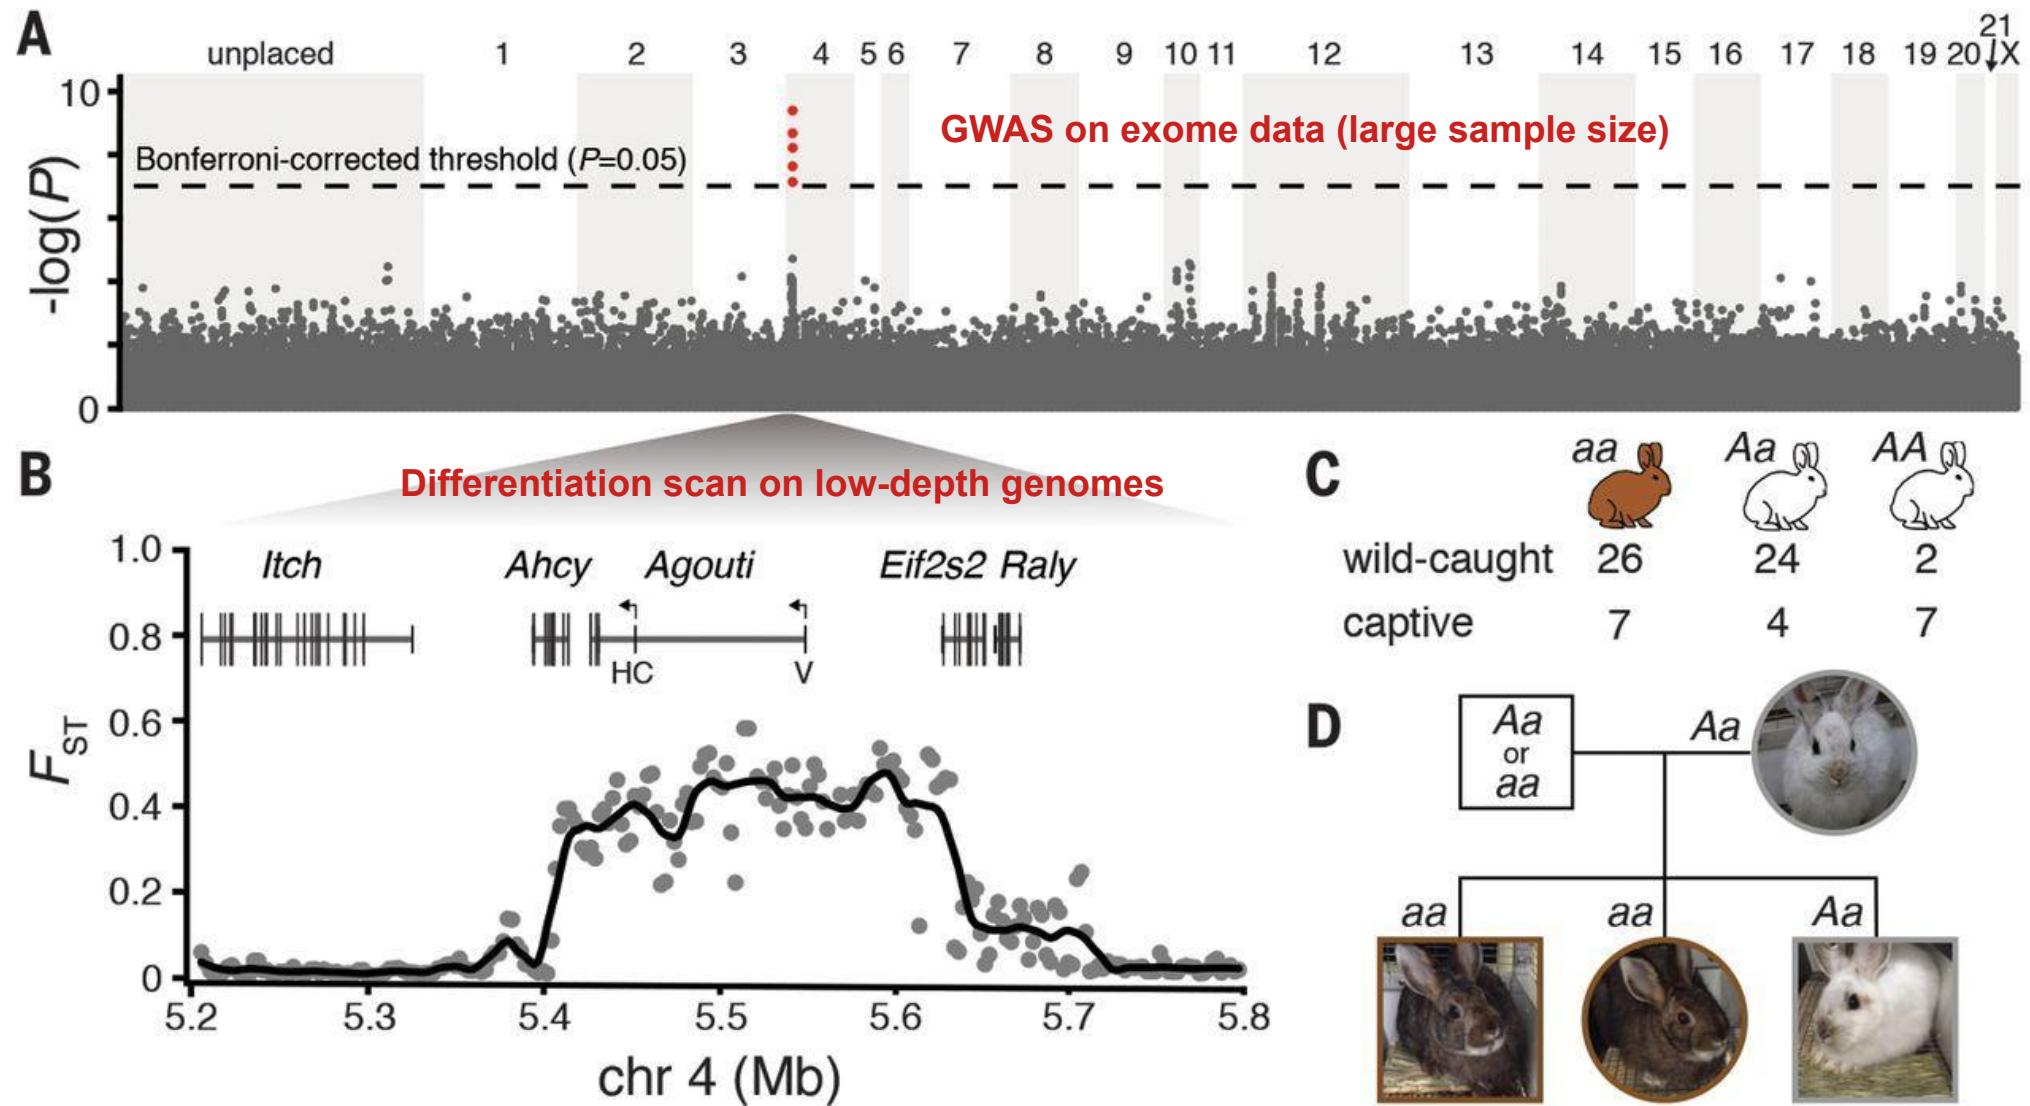


Effect on differentiation statistics

- Increase in F_{ST}
- Increase in Population Branch Tests
- More sophisticated methods exist (e.g. BAYPASS). See Friday's lecture on GEA.

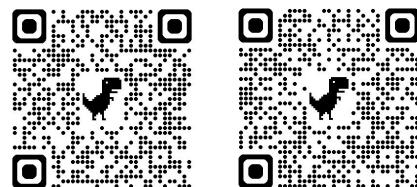
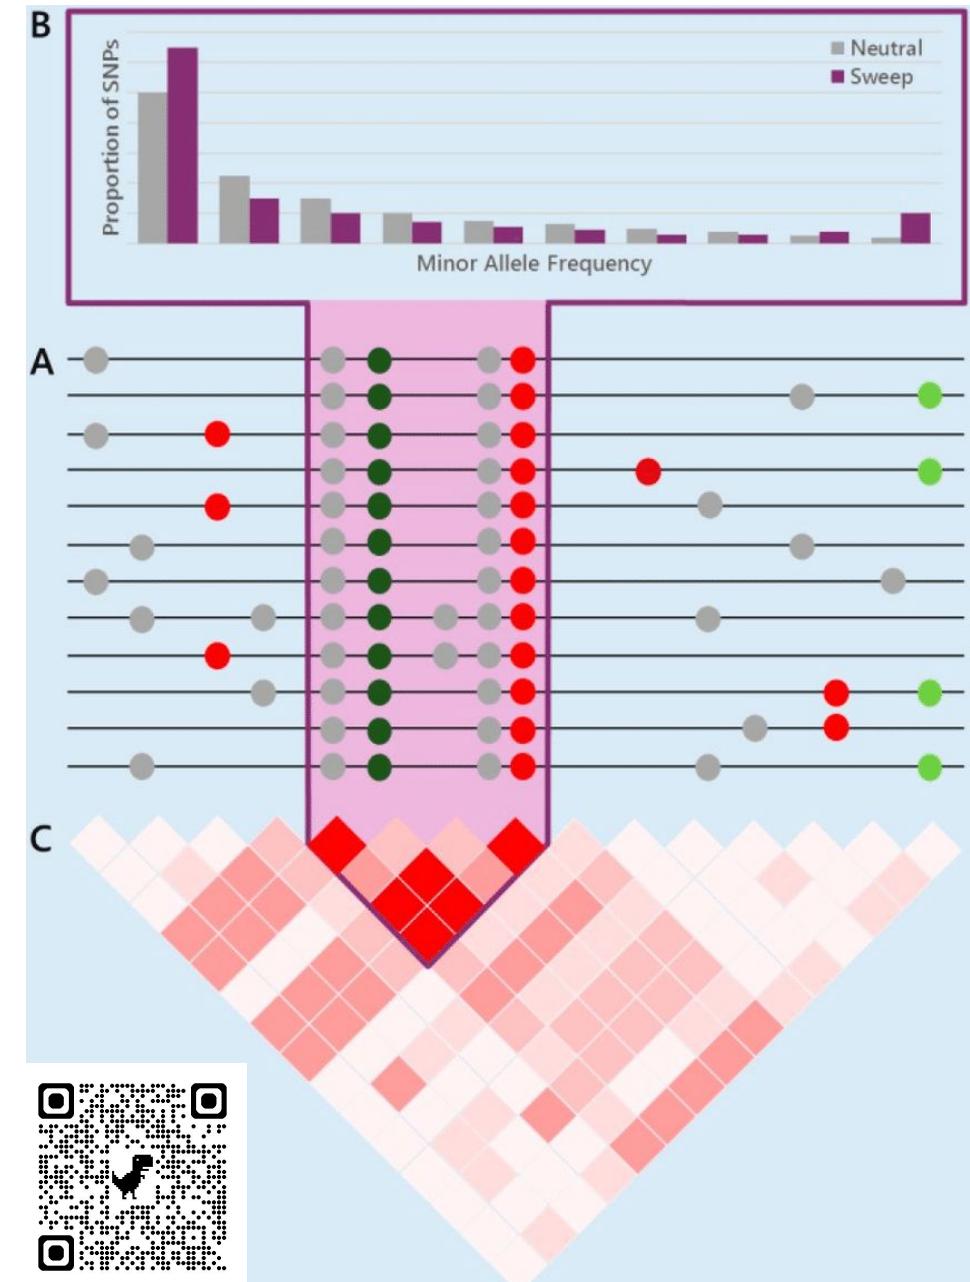


Scans of Differentiation



Effect on Linkage

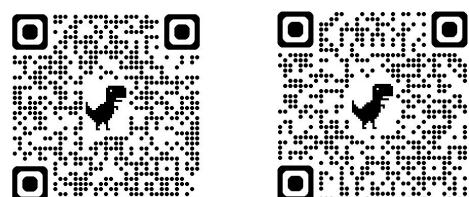
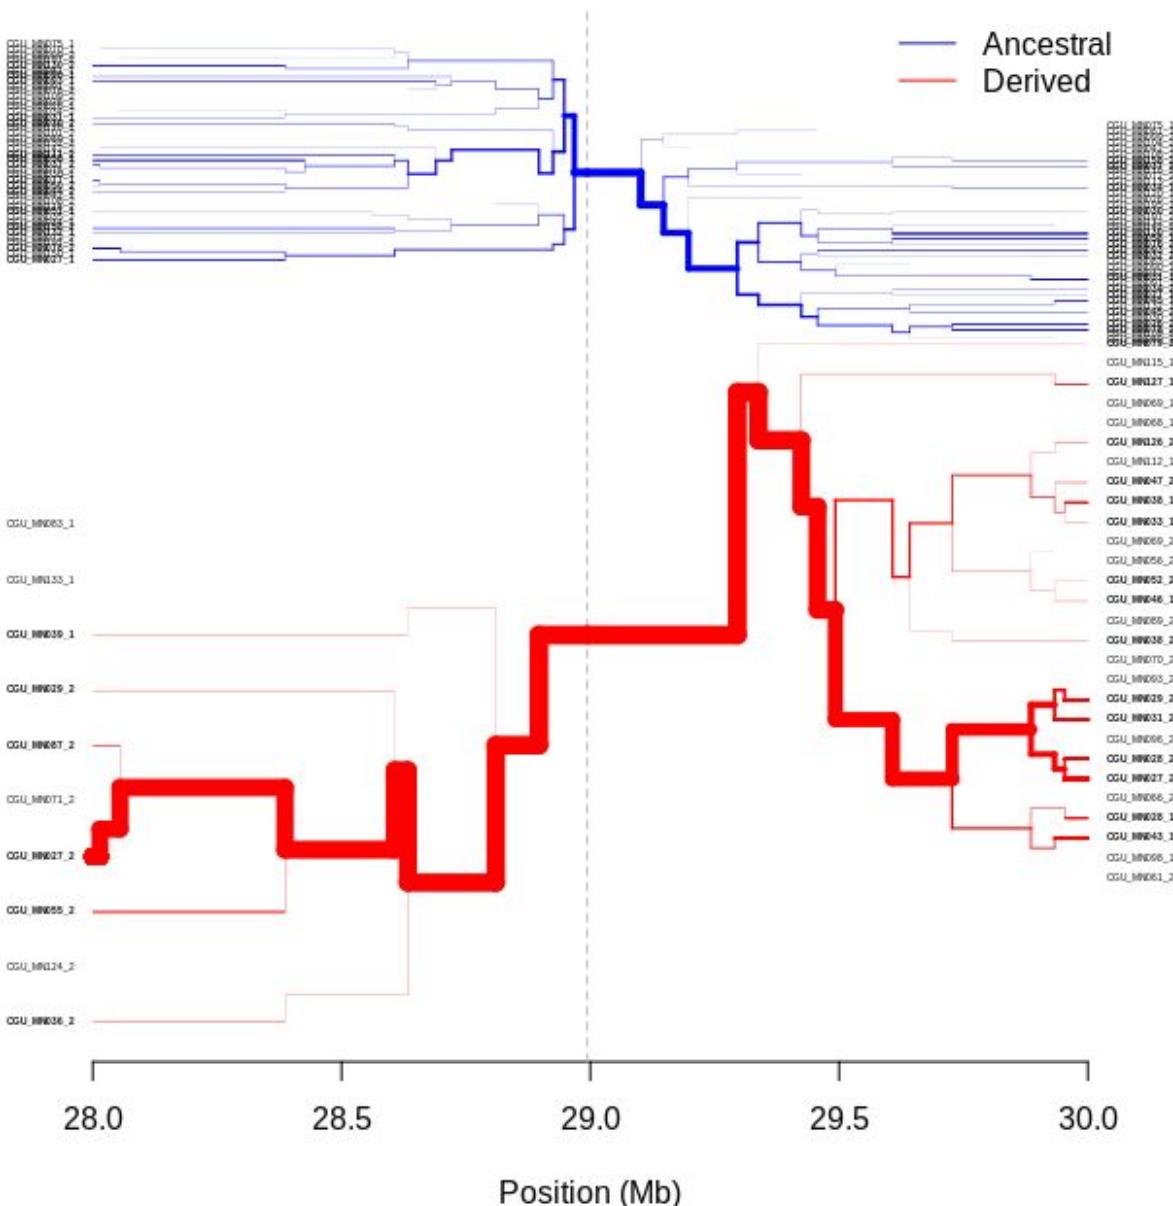
- The association between alleles across individuals breaks down because of recombination
- Recent selection leads to a fast increase in frequency of the selected allele and those at the vicinity (selective sweep)
- Selected haplotypes are frequent but have had less time to recombine with non-selected ones
- Can be detected by methods based on haplotype extension (Sabeti et al. 2006, see methods such as rehh, in the workshop). OmegaPlus is also based on LD.
- Requires phasing to be efficient.



Effect on Linkage

- The association between alleles across individuals breaks down because of recombination
- Recent selection leads to a fast increase in frequency of the selected allele and those at the vicinity (selective sweep)
- Can be detected by methods based on haplotype extension (e.g. iHS score, in rehh, see workshop)
- Possible to compare extension across populations instead of within (XP-EHH, Rsb).

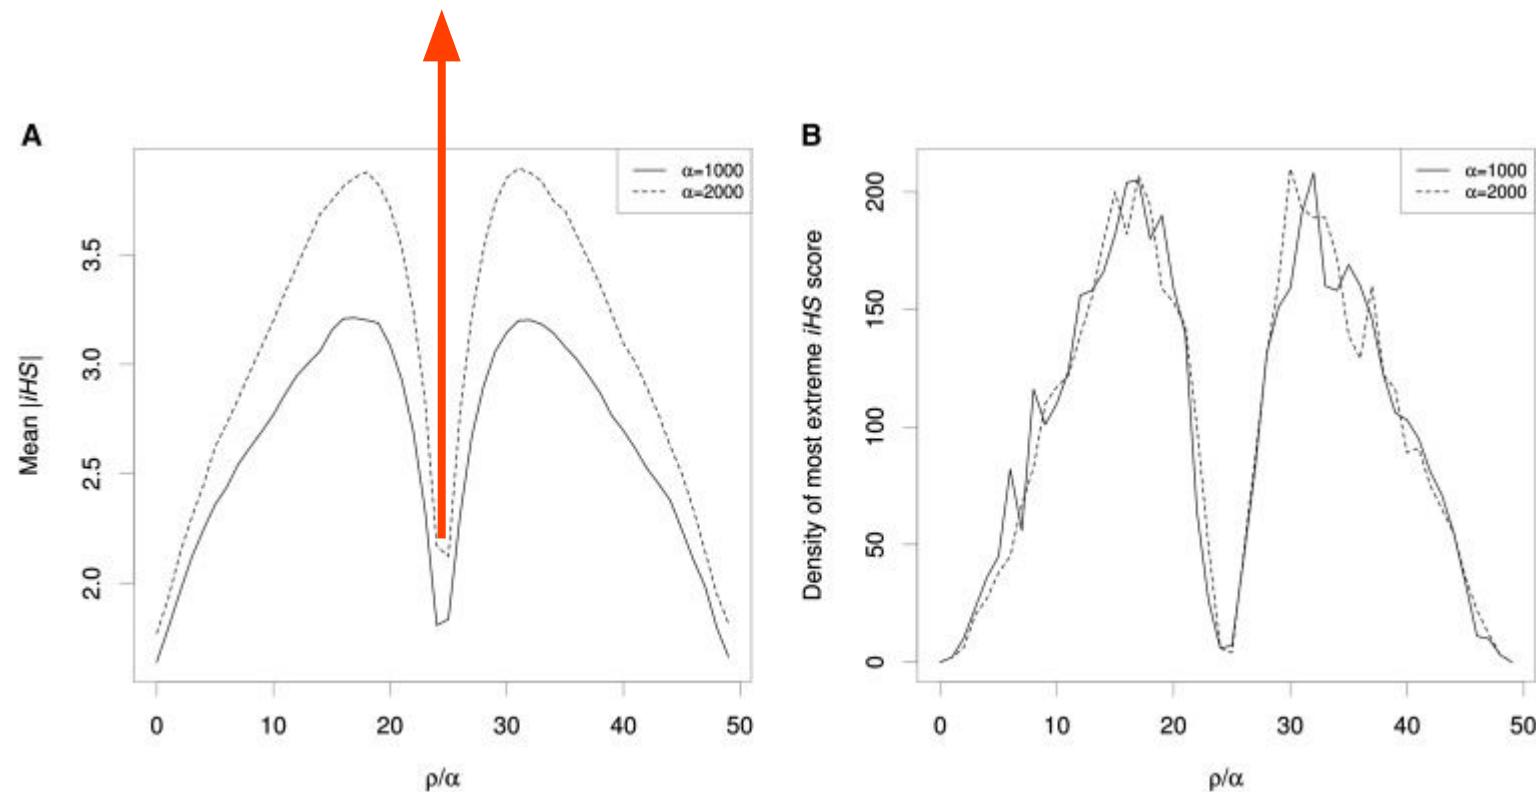
Haplotype furcations around 'F1205400'



Effect on Linkage

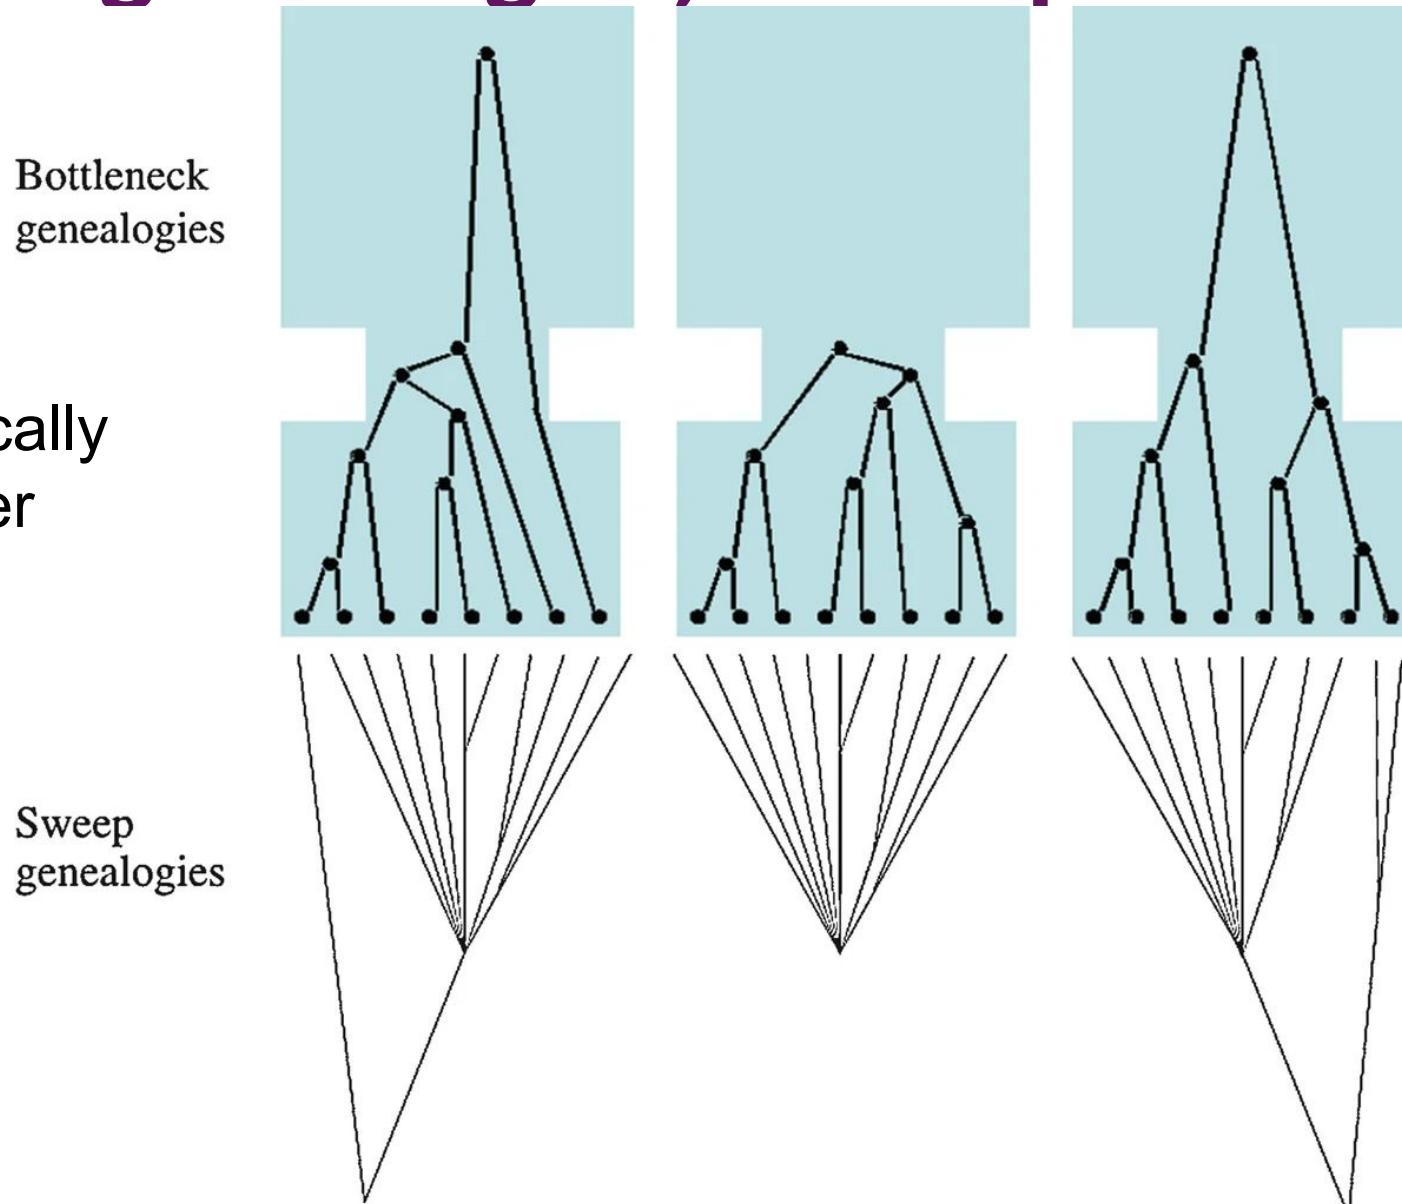
- The association between alleles across individuals breaks down because of recombination
- Recent selection leads to a fast increase in frequency of the selected allele and those at the vicinity (selective sweep)
- Can be detected by methods based on haplotype extension (e.g. rehh, see workshop)

If the sweep is fixed, there is only the derived allele at the centre, and the iHS score can not be computed

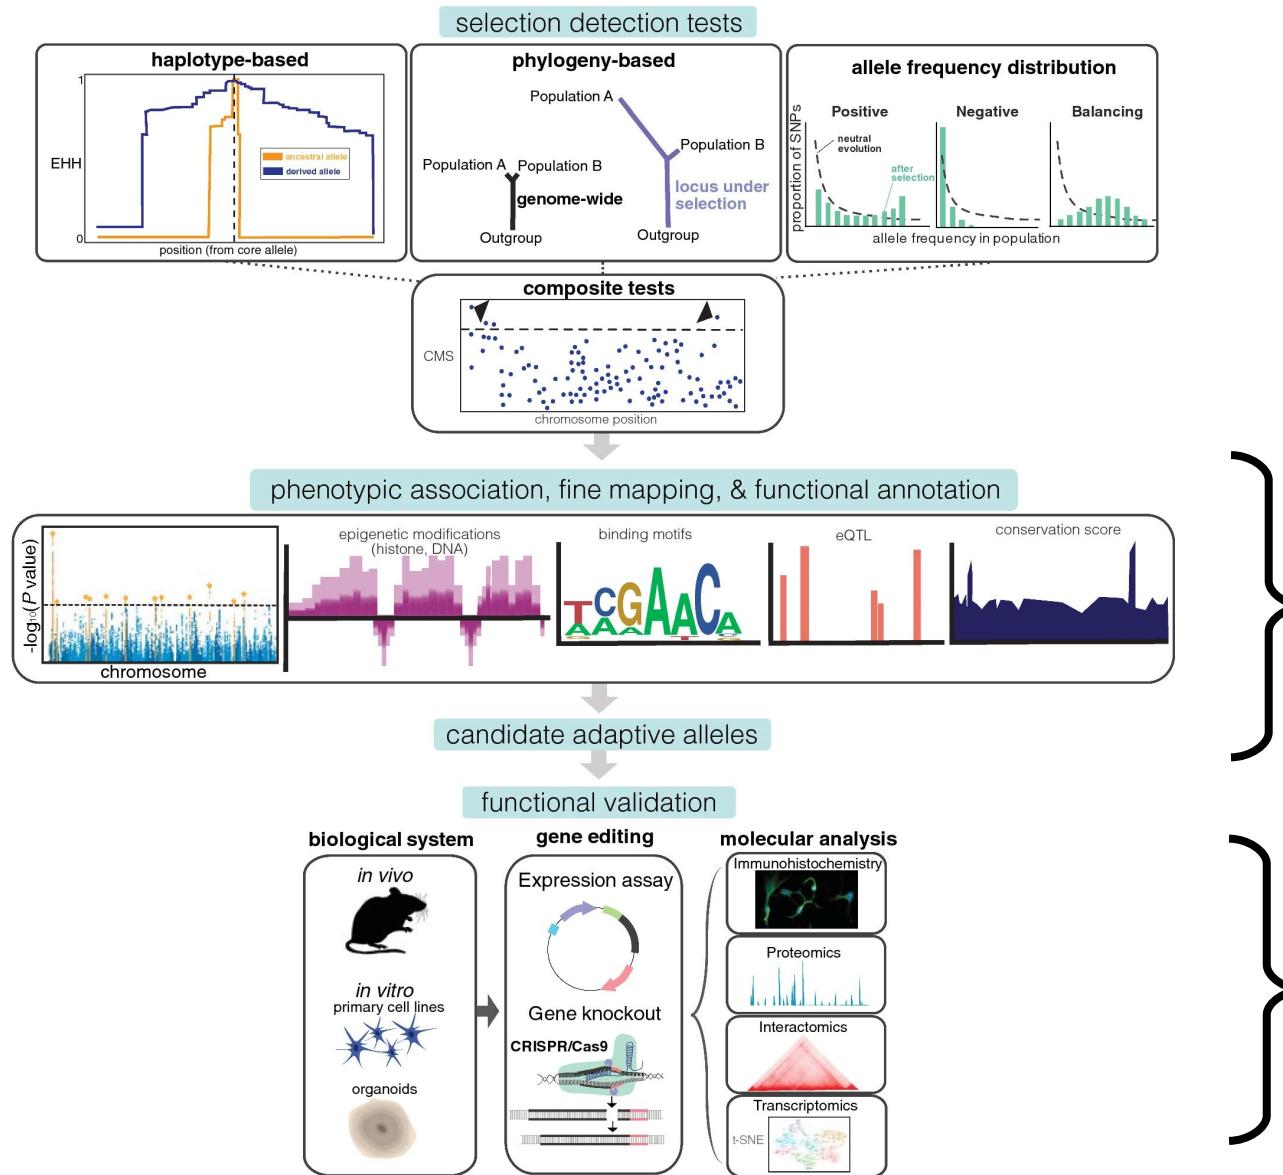


Effect on local genealogies, and a problem...

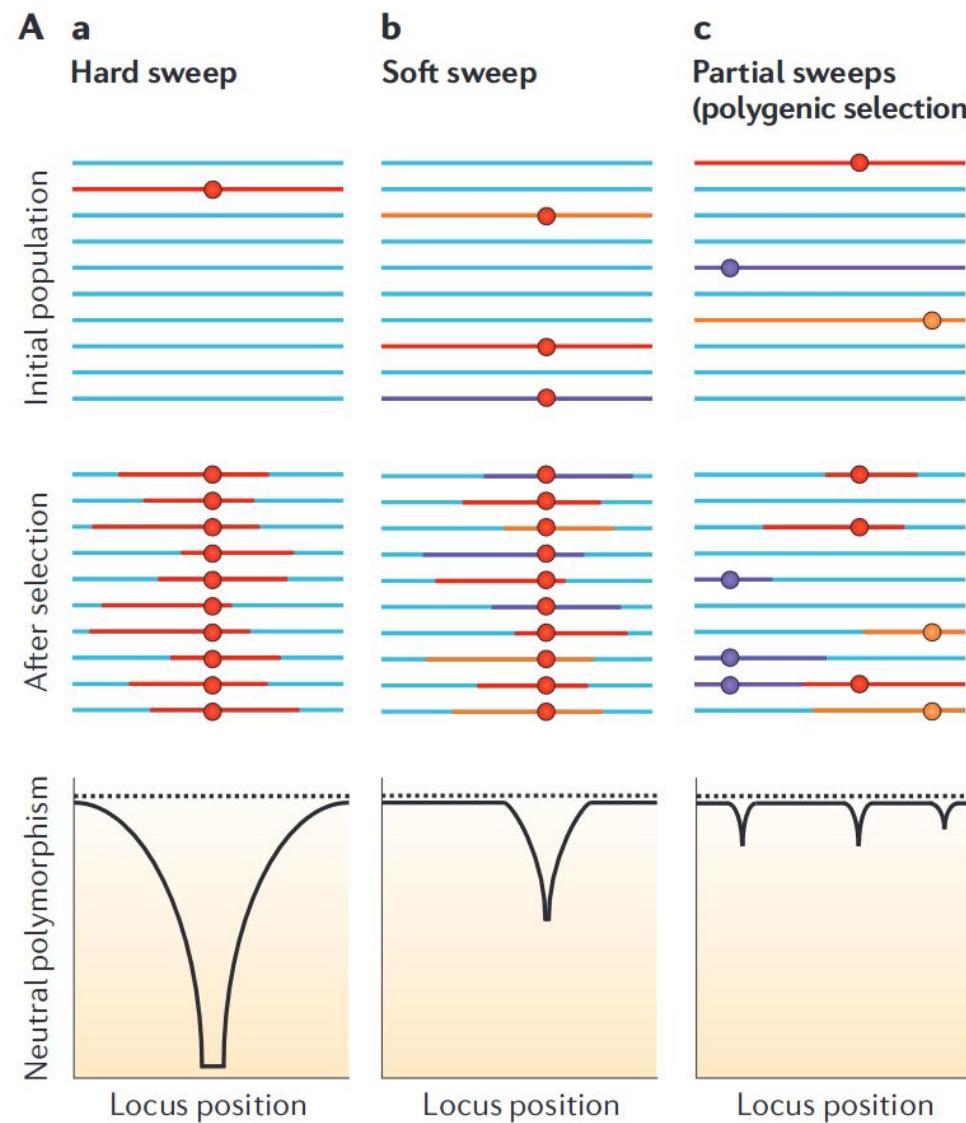
- But selection acts locally and demography over the whole genome.
- Right??
- Issue of variance...



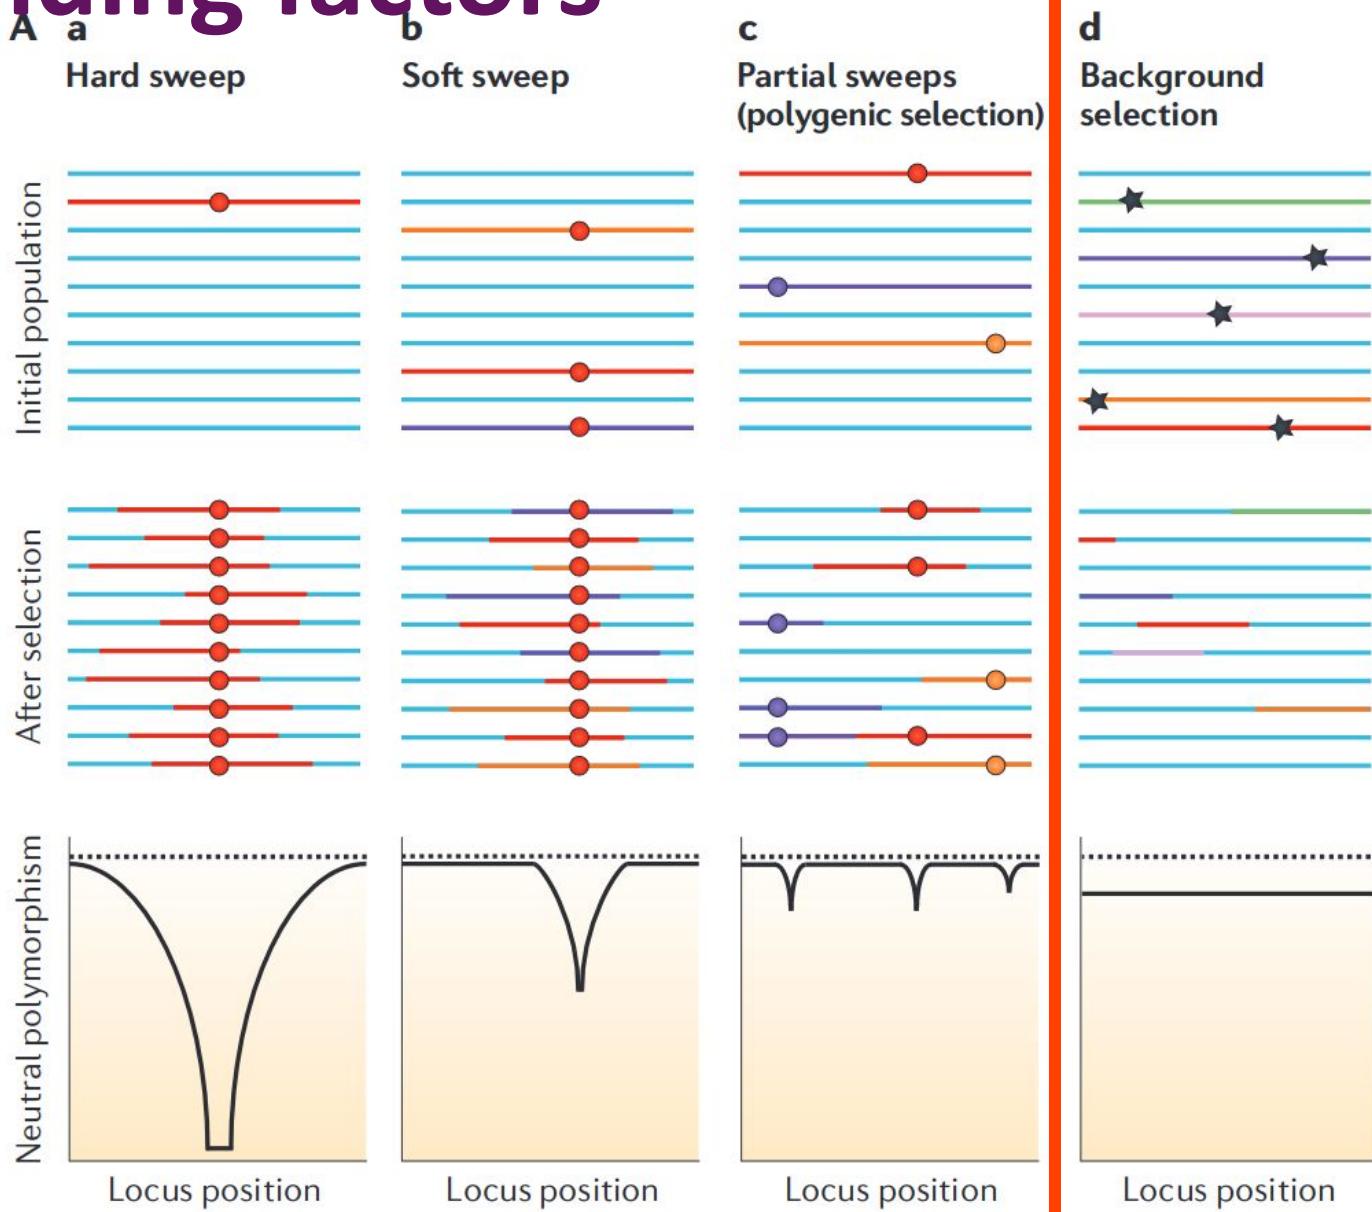
The ideal situation



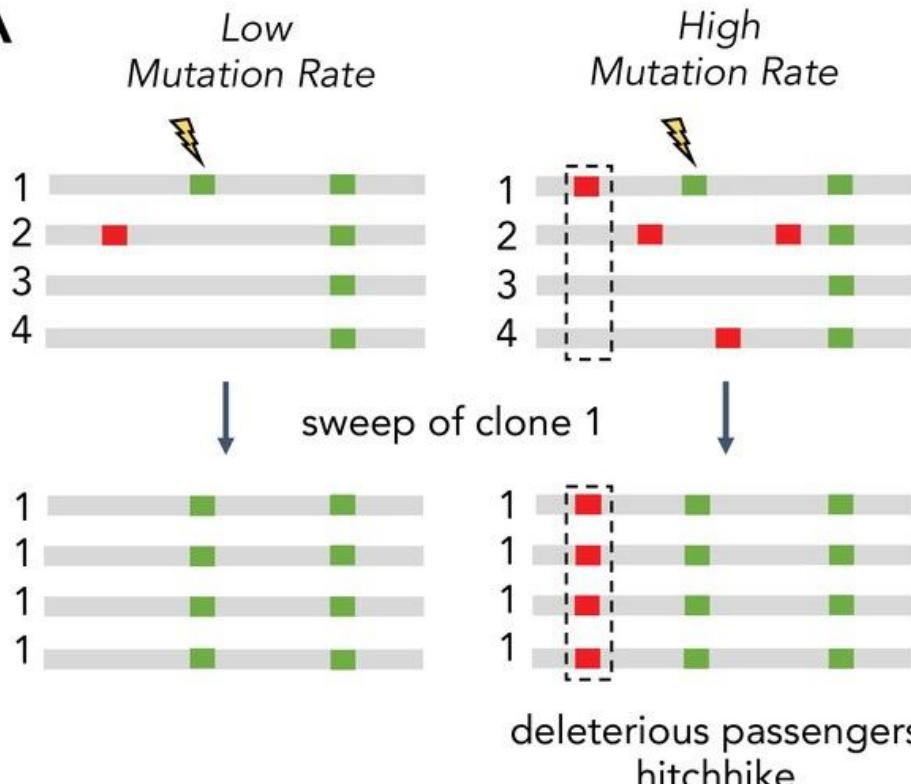
More complicated. Soft and partial sweeps



Confounding factors



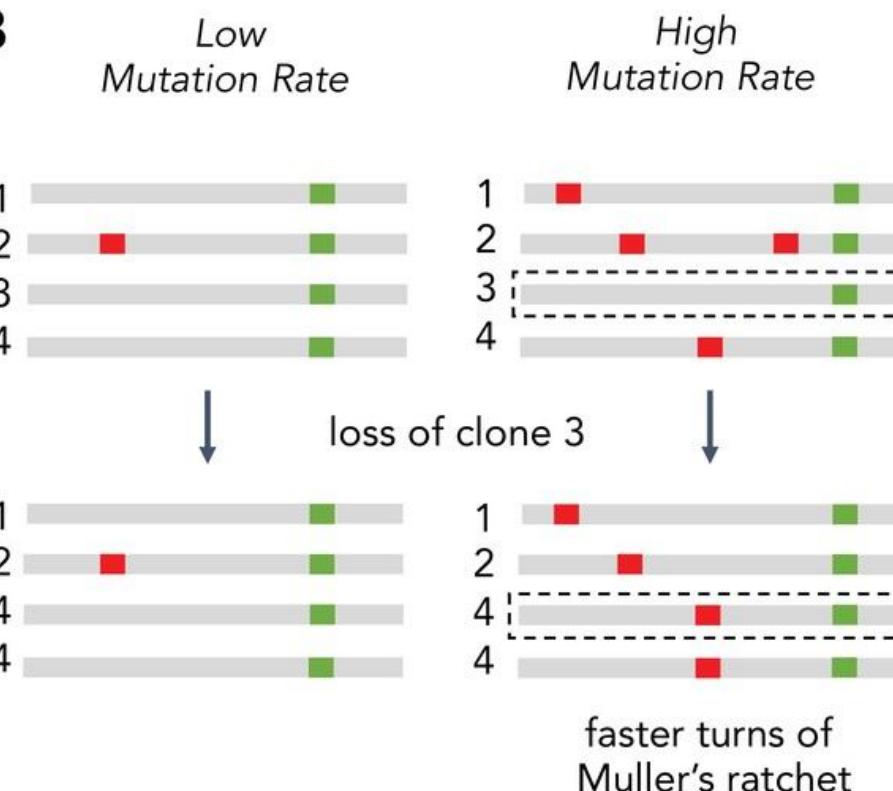
'Linked' selection

A

Genetic Hitchhiking

fixation of drivers causes fixation of deleterious passengers

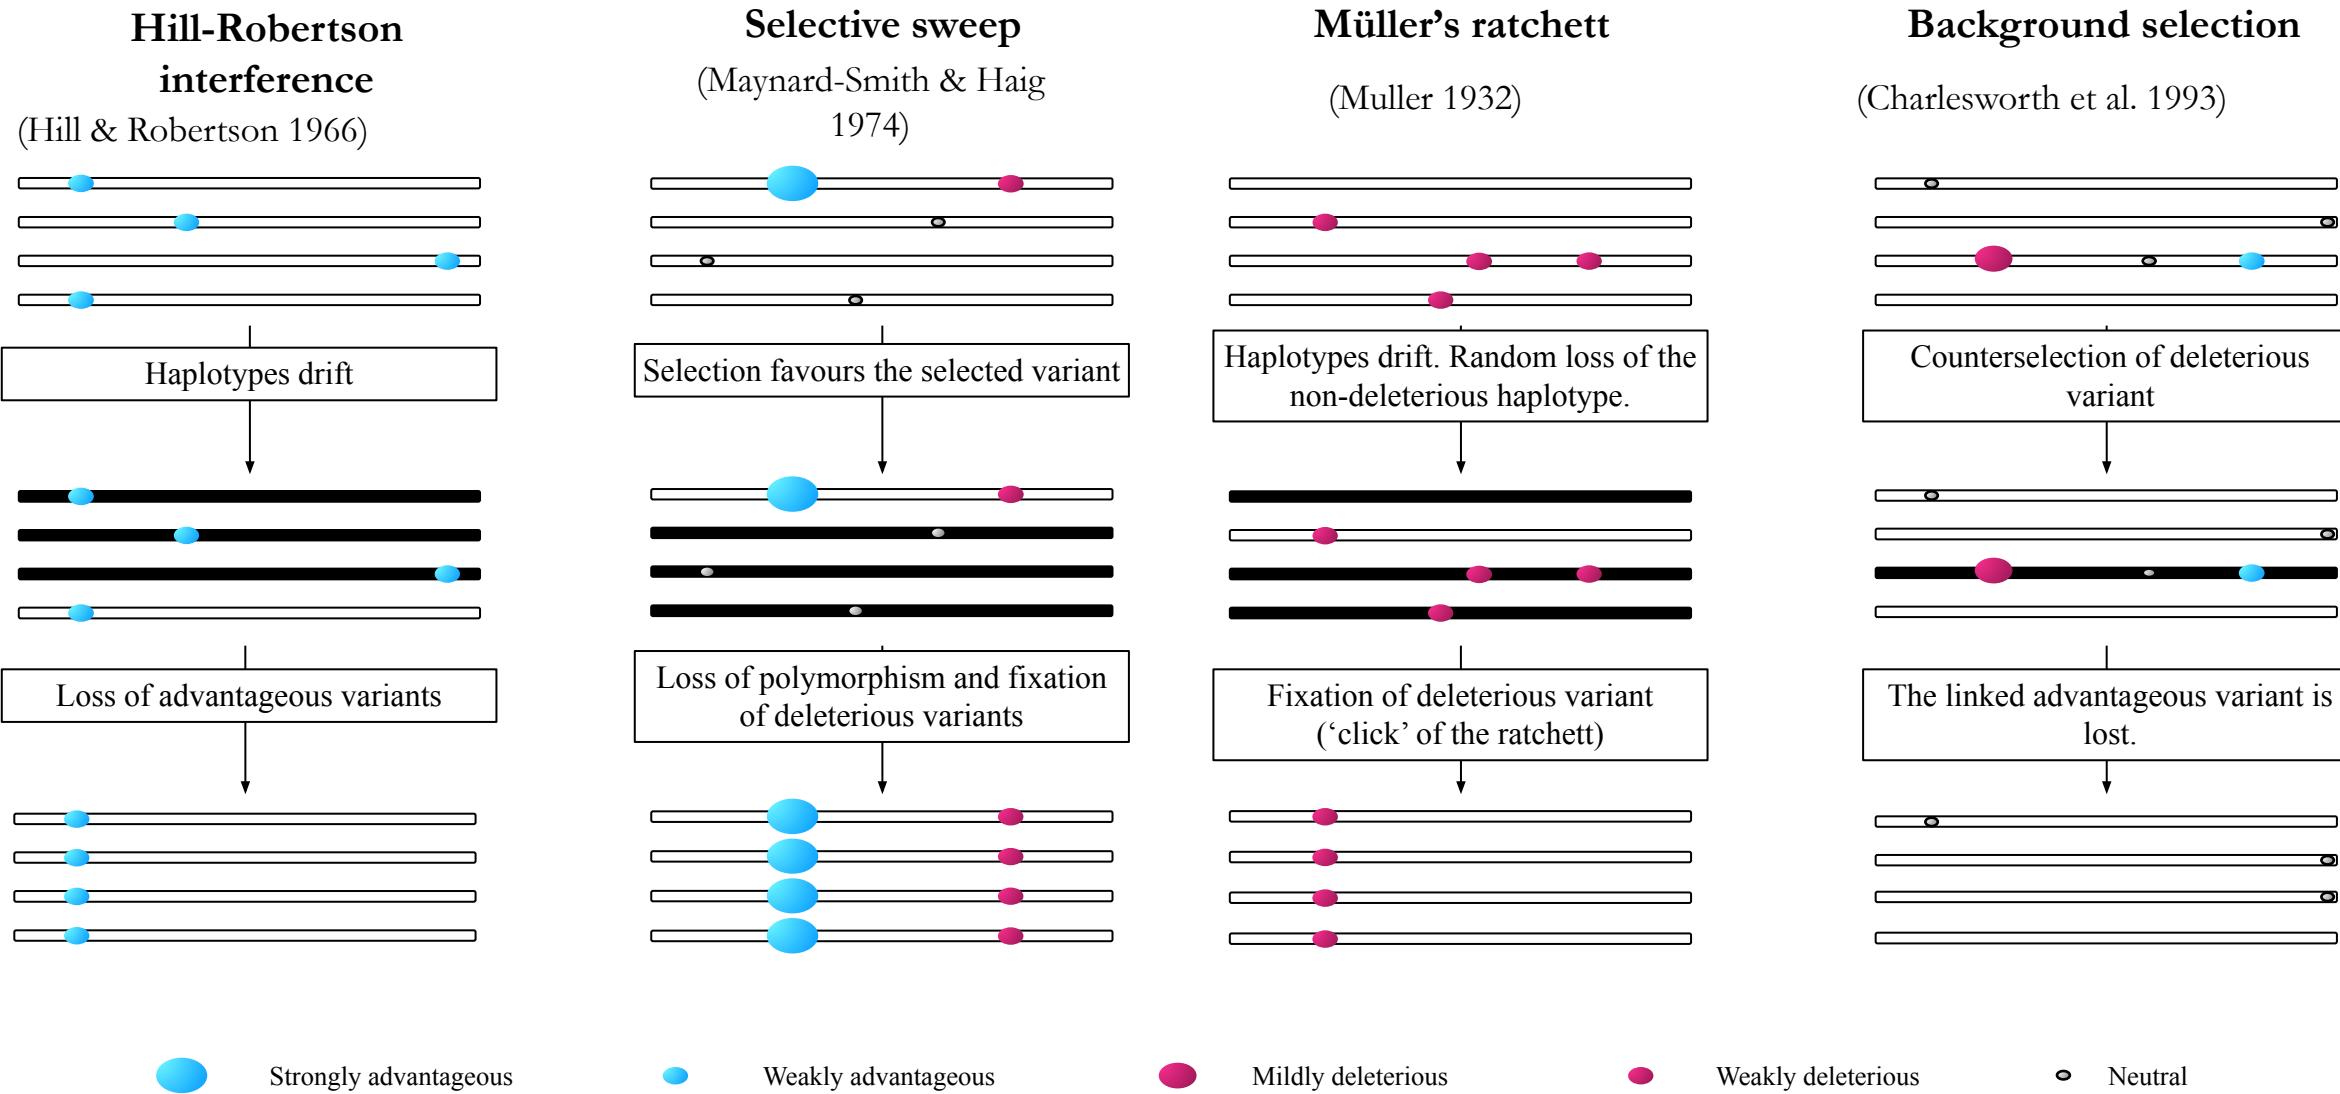
Note that we consider populations that sexually reproduce, but if recombination is locally low relative to mutation, we converge to this clonal case

B

Muller's Ratchet

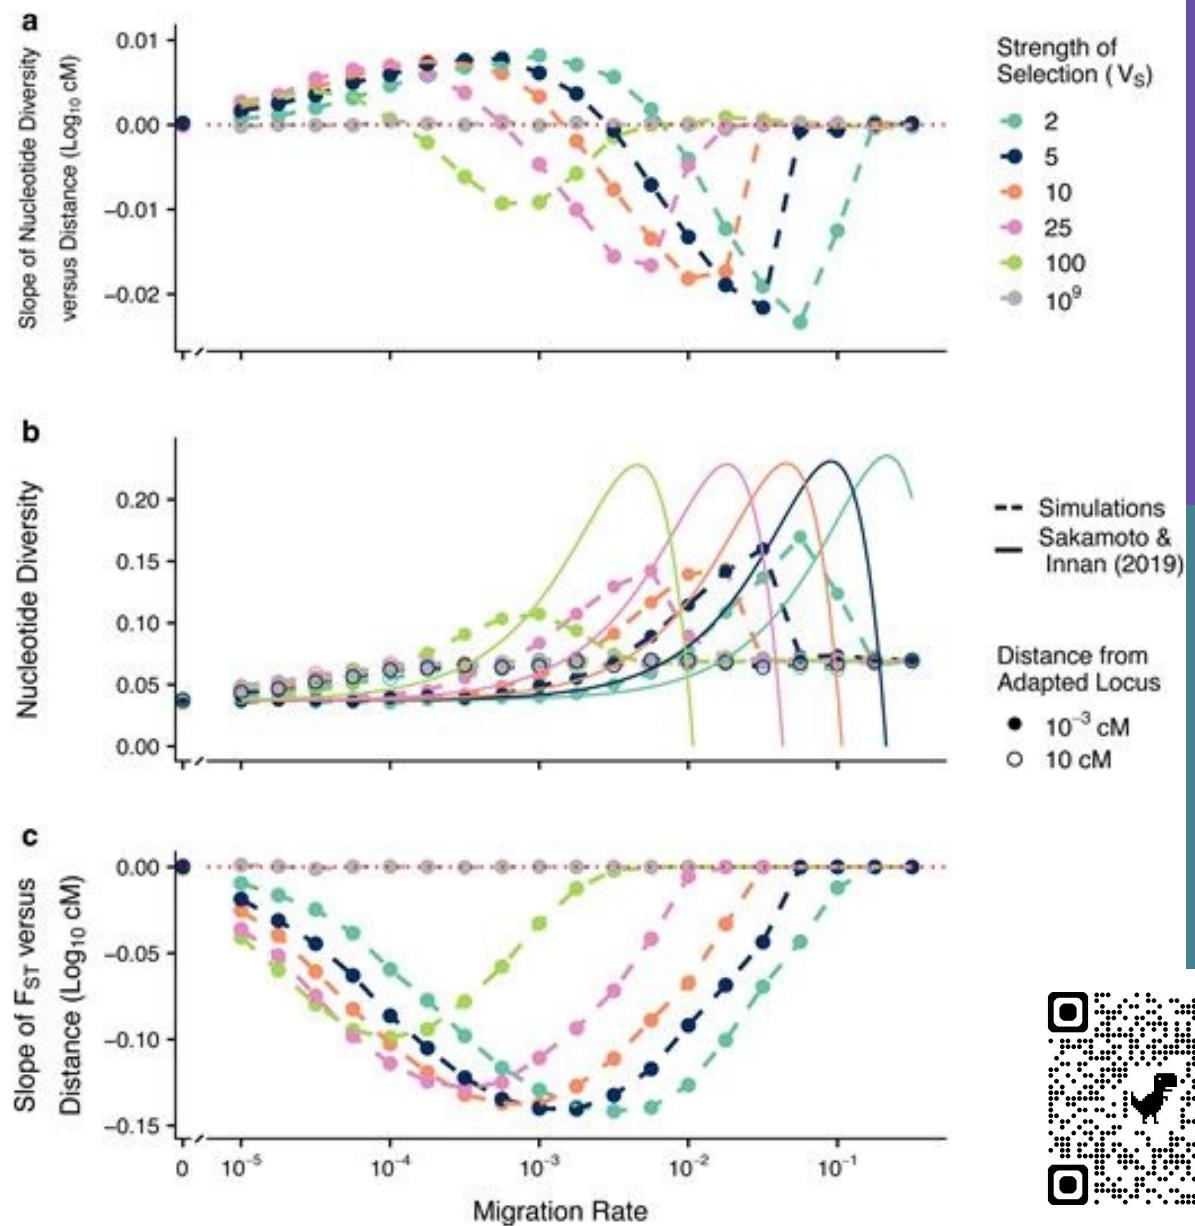
even fittest clone accumulates deleterious passengers

■ beneficial driver
■ deleterious passenger



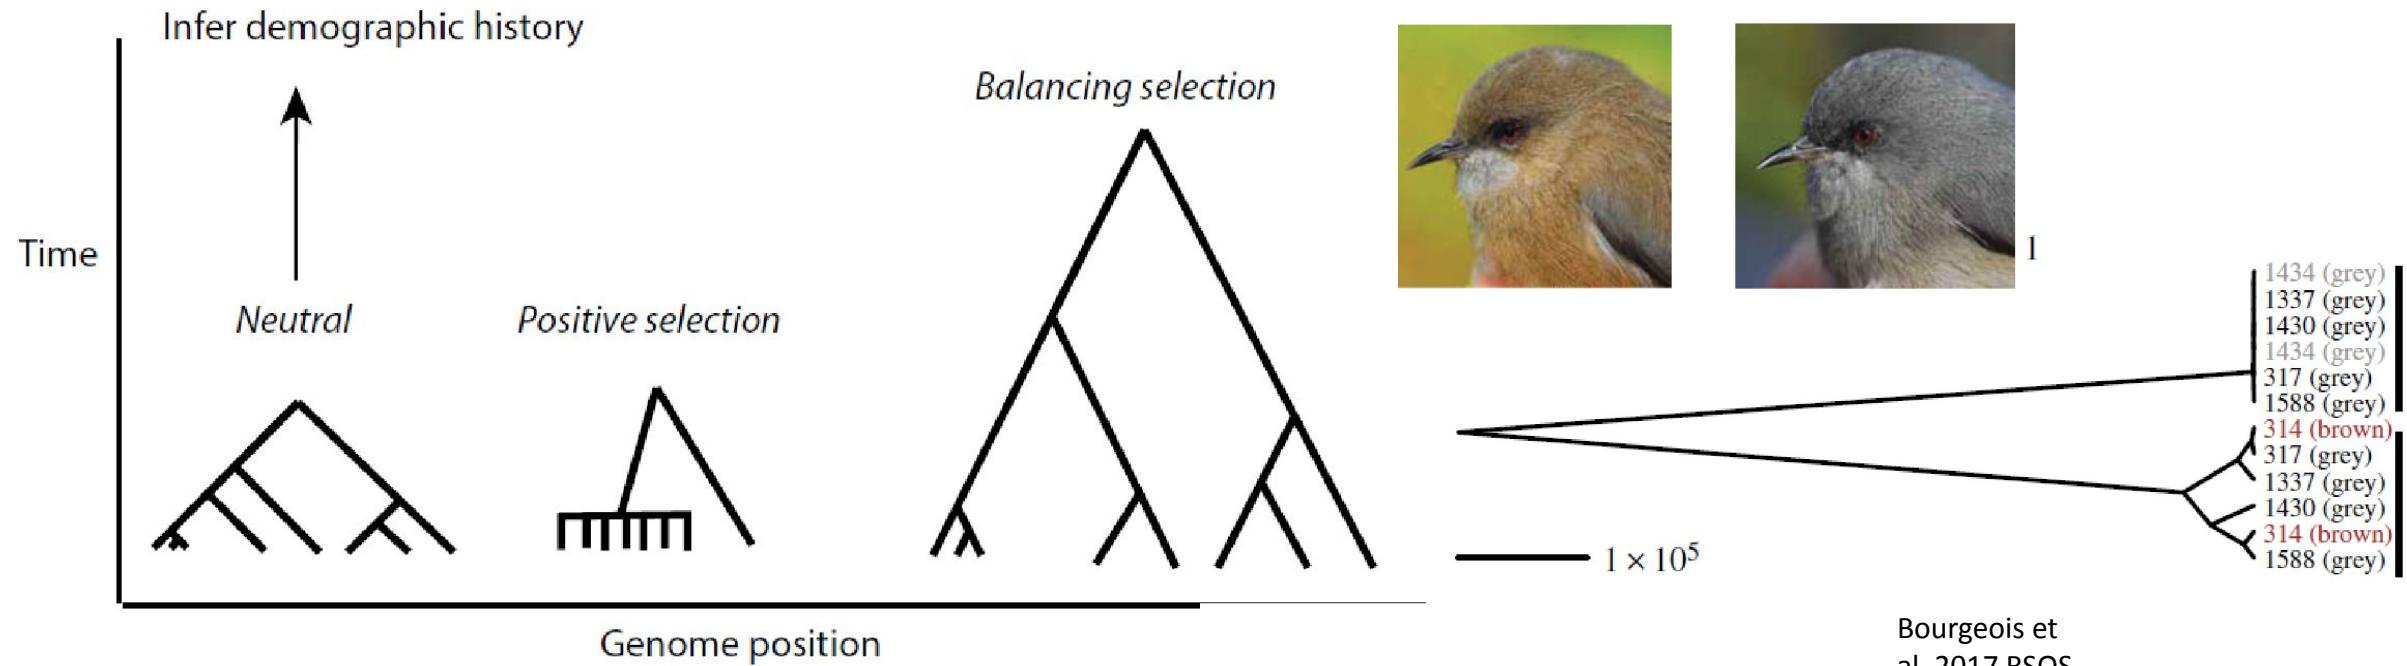
Confounding factors

- Local adaptation might show more complicated patterns
- This is due to the interplay between selection, recombination, and migration
- Recent simulation studies suggest that both peaks and troughs in diversity are possible
- Think of the scale at which the sweep may have happened.
- Whenever possible, run simulations to get a sense of what you may expect/detect.



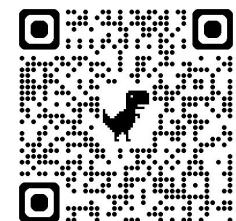
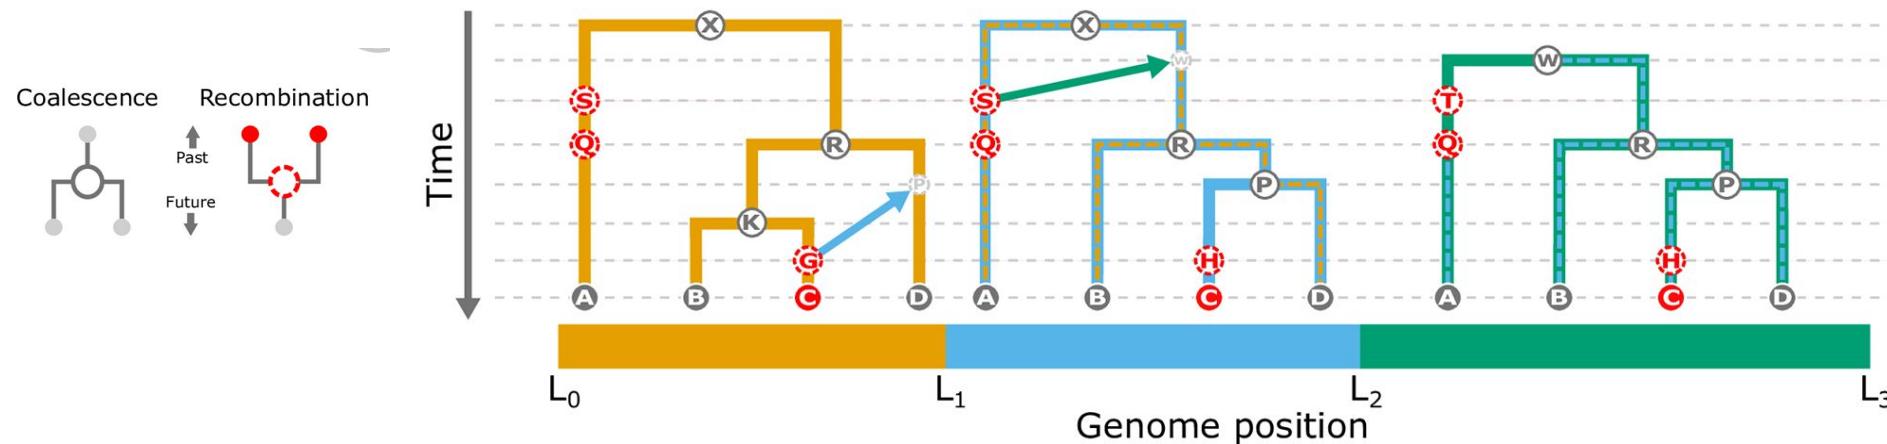
Reconstruct past genealogies and population structure

- Coalescence rate and topology will change if there is selection at a gene
- Genealogies at distant loci are less similar than ones at closely linked loci (because of recombination breaking associations between alleles within a chromosome)



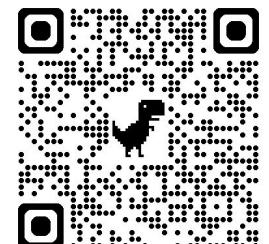
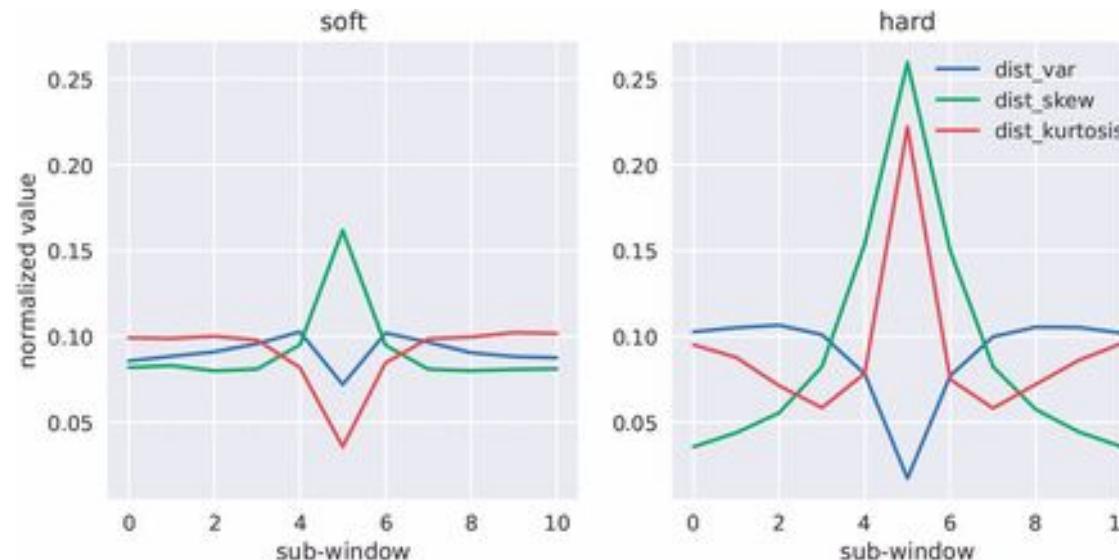
Reconstruct past genealogies and population structure

- One can reconstruct genealogies using polymorphisms along the genome
- The rate at which new branches merge from present to past is inversely correlated with population size (small populations -> high rates of **coalescence**)
- Local deviations in genealogies and recombination can be fully described (ARGWeaver, Relate)
- Can be used to reconstruct allele frequency over time and infer s (CLUES2)



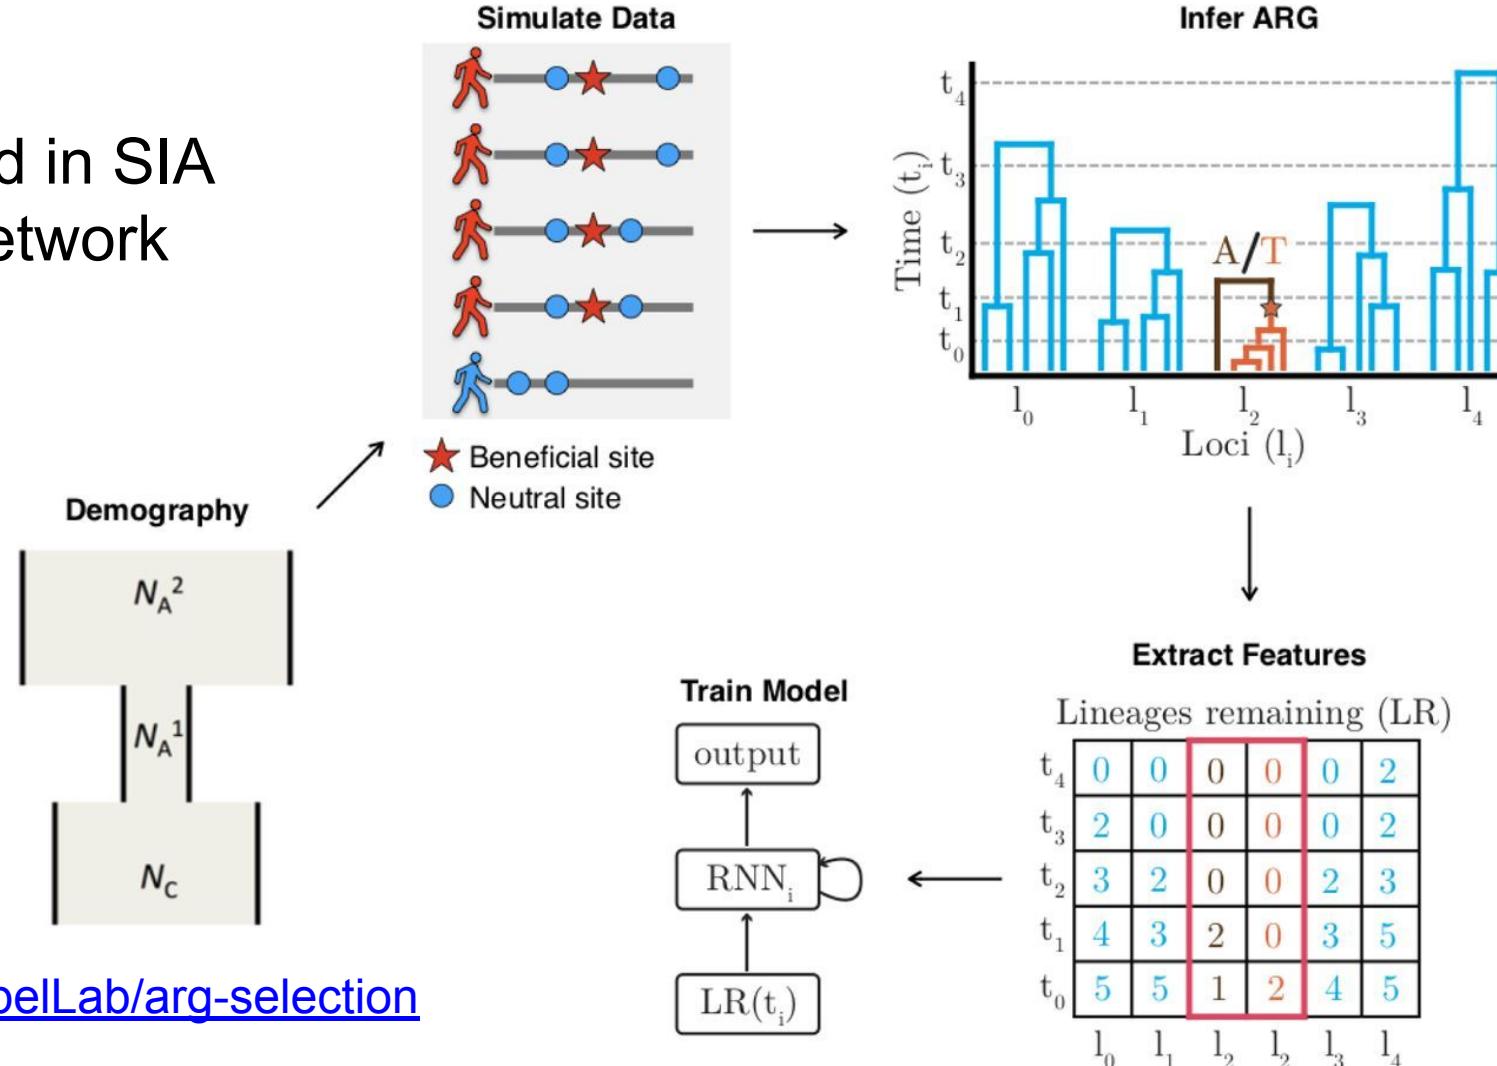
Integration of signals: machine learning (summary statistics)

- Machine and deep learning share a principle with ABC (and they overlap)
- We simulate neutral and selected datasets, obtain summary statistics, and compare them to observations through a trained algorithm.
- Possible advantage compared to ABC: fewer issues with the choice of summary statistics.
- diploS/HIC and S/HIC



Integration of signals: machine learning (Ancestral Recombination Graphs)

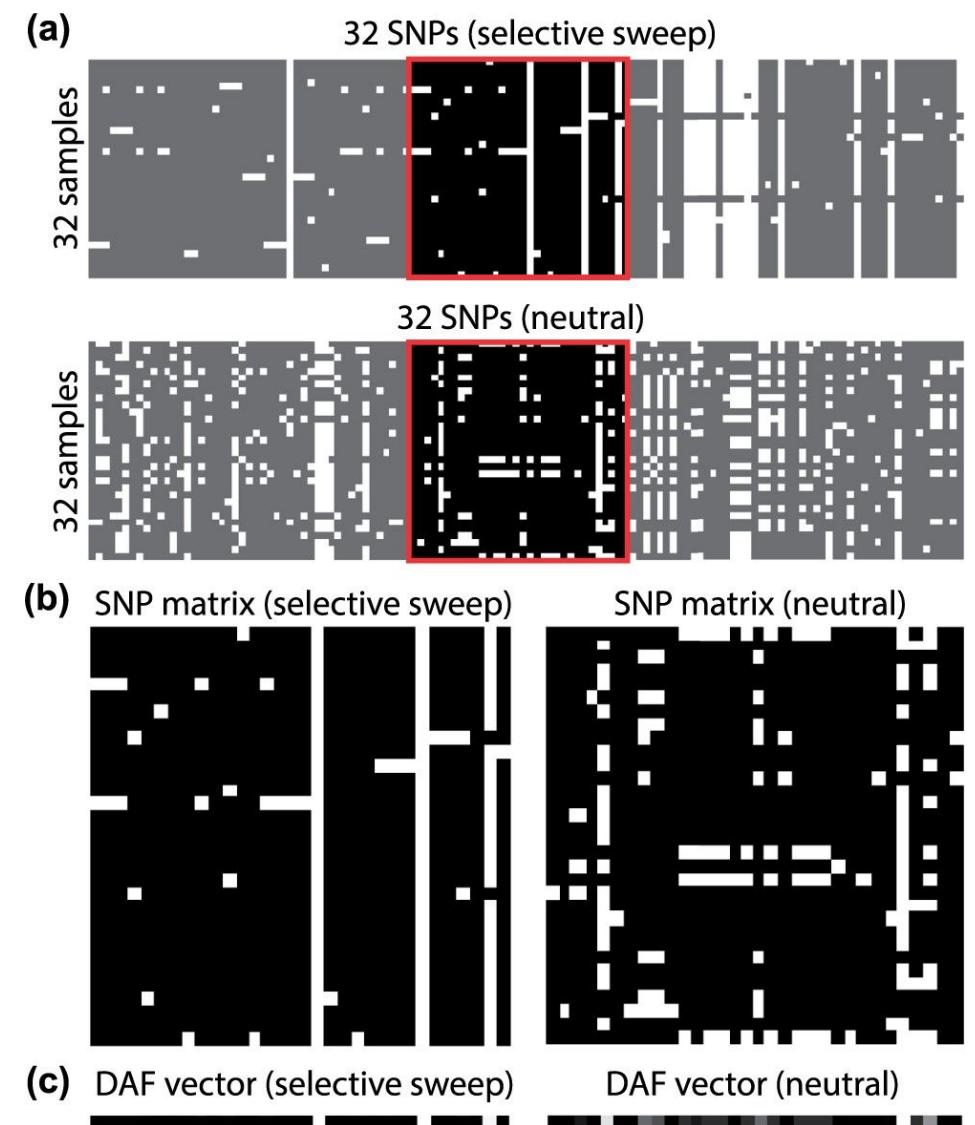
Method implemented in SIA
Recurrent Neural Network



<https://github.com/CshlSiepelLab/arg-selection>

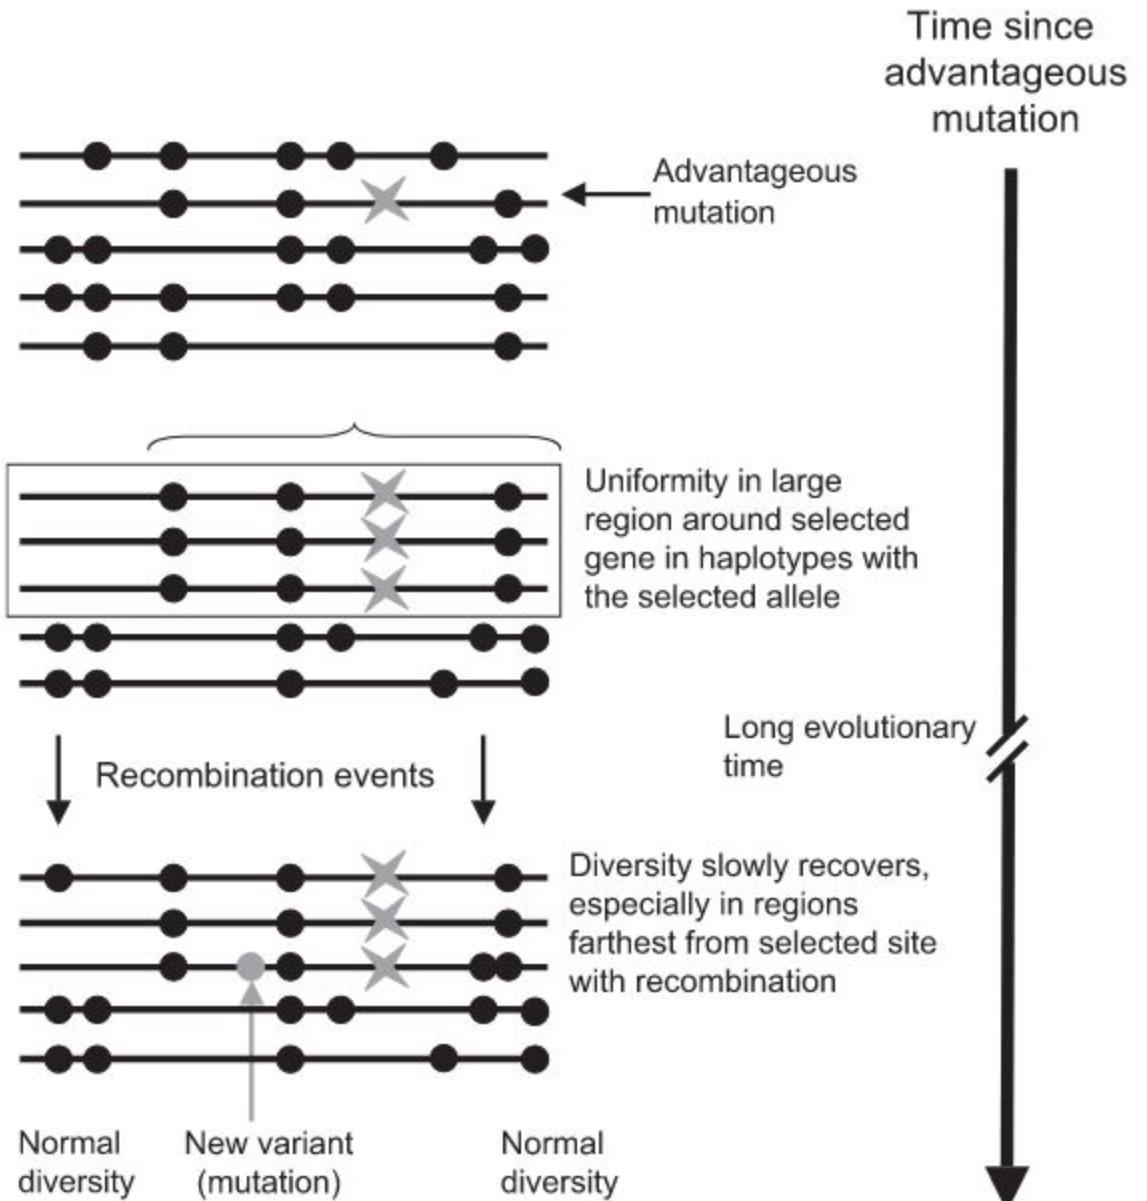
Integration of signals: machine learning (image recognition)

- Direct image recognition is also seducing in principle to get rid of summary statistics
- Take the alignment and train the algorithm to recognize images that fit with a sweep
- Fast- NN (takes the vector of frequencies as an image).
- SweepNet (takes the alignment as an image).



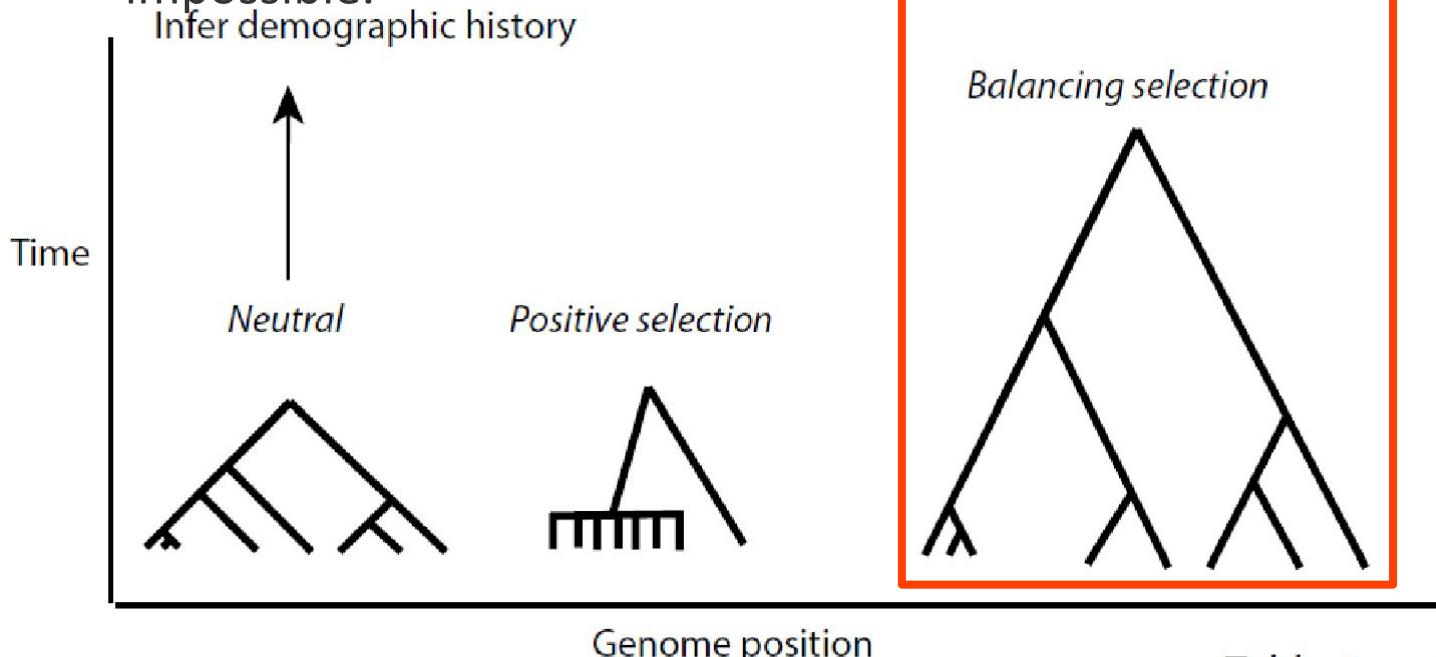
Balancing selection.

- This is an umbrella term, covering processes that maintain genetic diversity over time at a locus. Especially interesting to people working on diversifying processes, such as host-parasite interaction, or self-incompatibility.
- Think of a partial sweep that never gets fixed.
- The main challenge here is the loss of linkage over evolutionary times, which may make the signal very narrow.



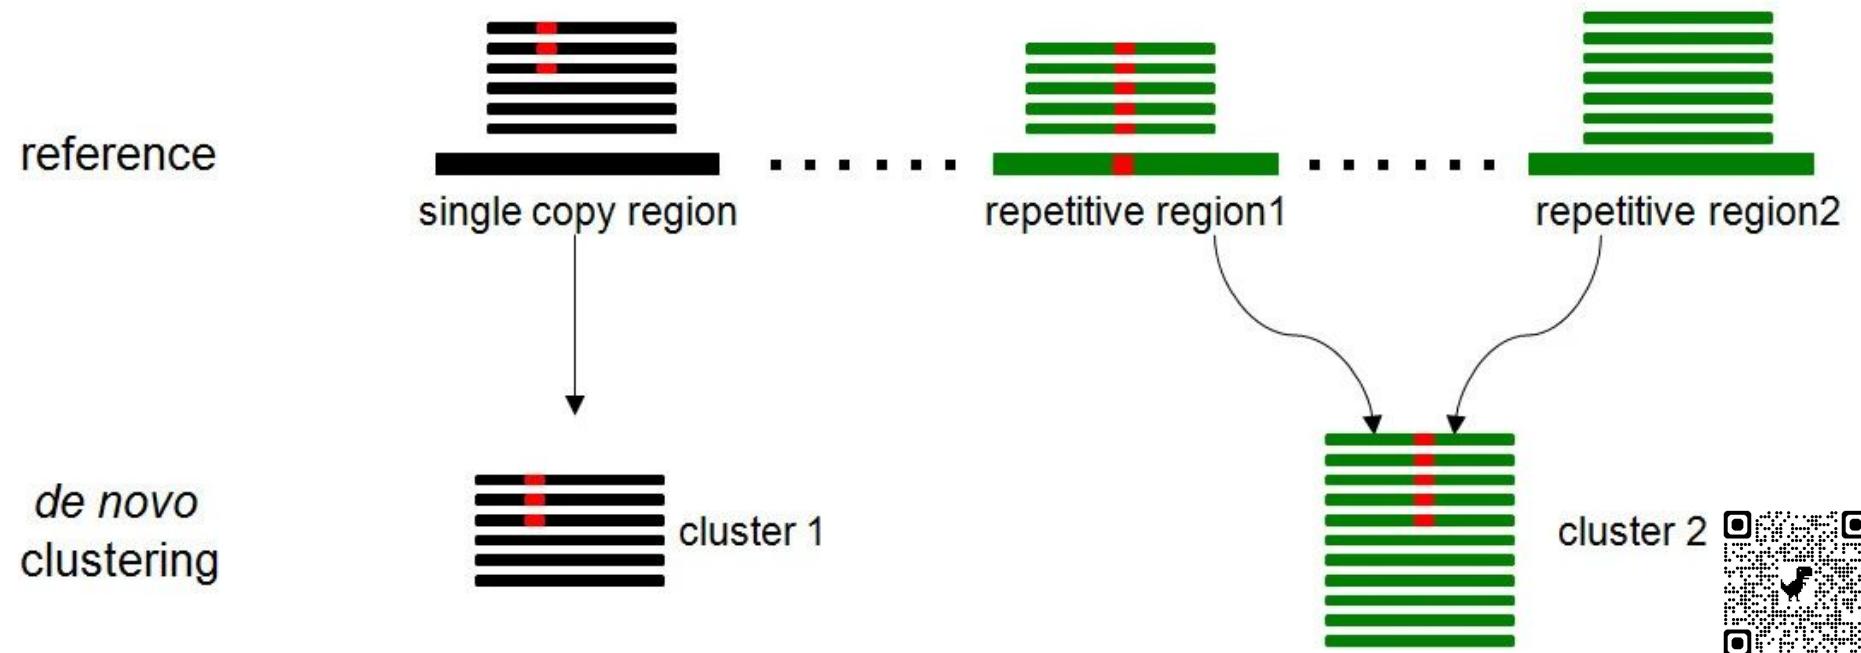
Long-term balancing selection.

- The genealogies near a locus under long-term balancing selection should be older. ARGs can be useful here.
- However, this is also dangerous: misalignment in your (short read) data along undetected paralog loci may lead to false signals
- Very old balancing selection can also involve very divergent haplotypes, making alignment impossible.



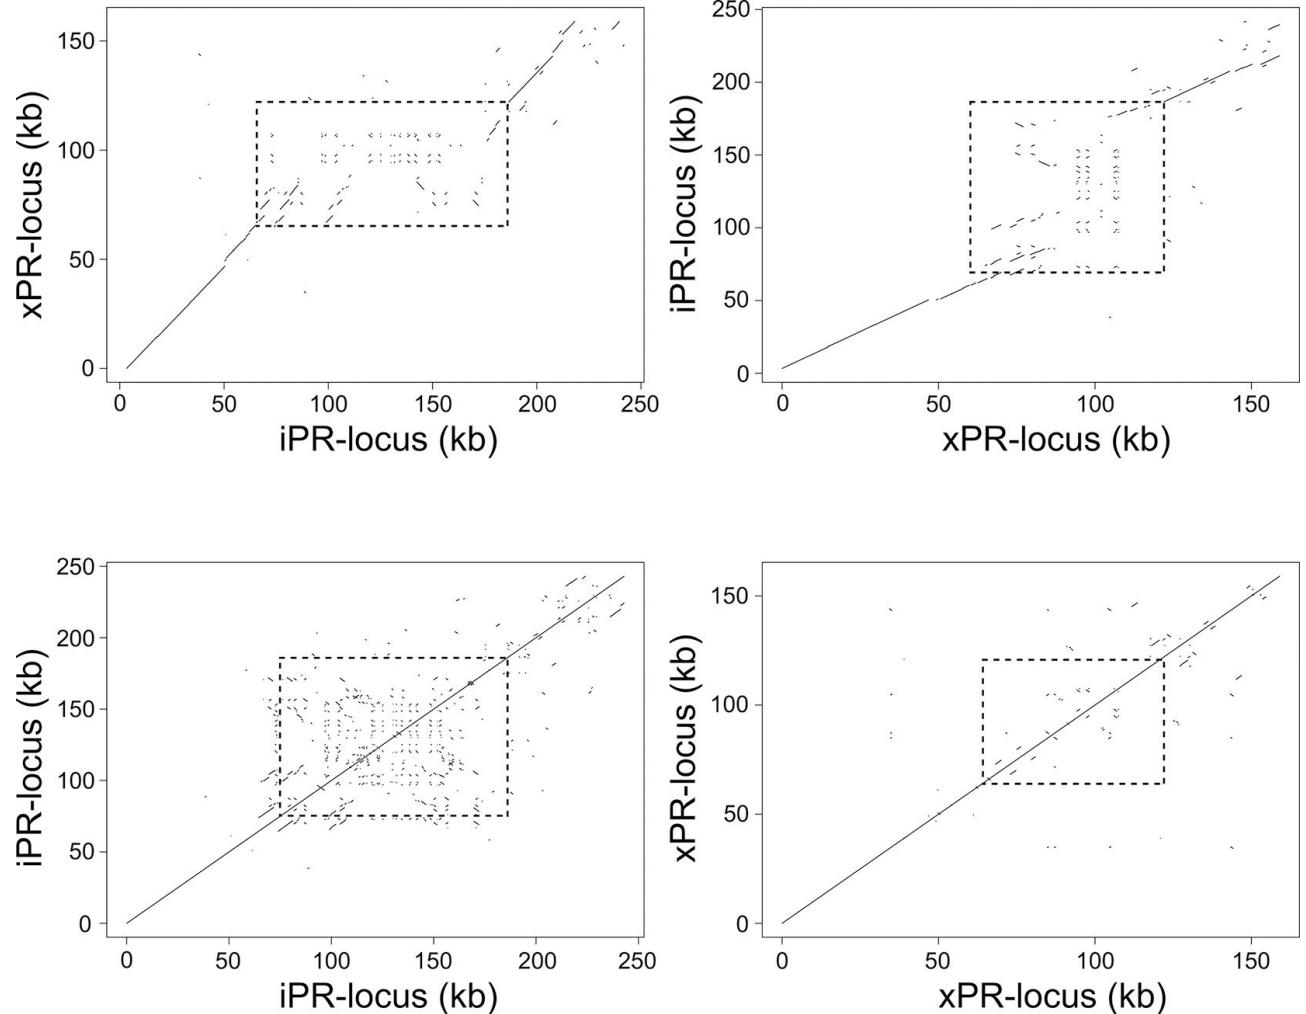
Long-term balancing selection.

- The genealogies near a locus under long-term balancing selection should be older. ARGs can be useful here.
- However, this is also dangerous: misalignment in your (short read) data along undetected paralog loci may lead to false signals
- Very old balancing selection can also involve very divergent haplotypes, making alignment impossible



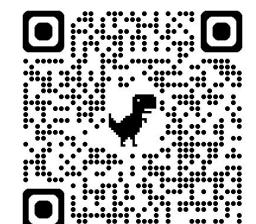
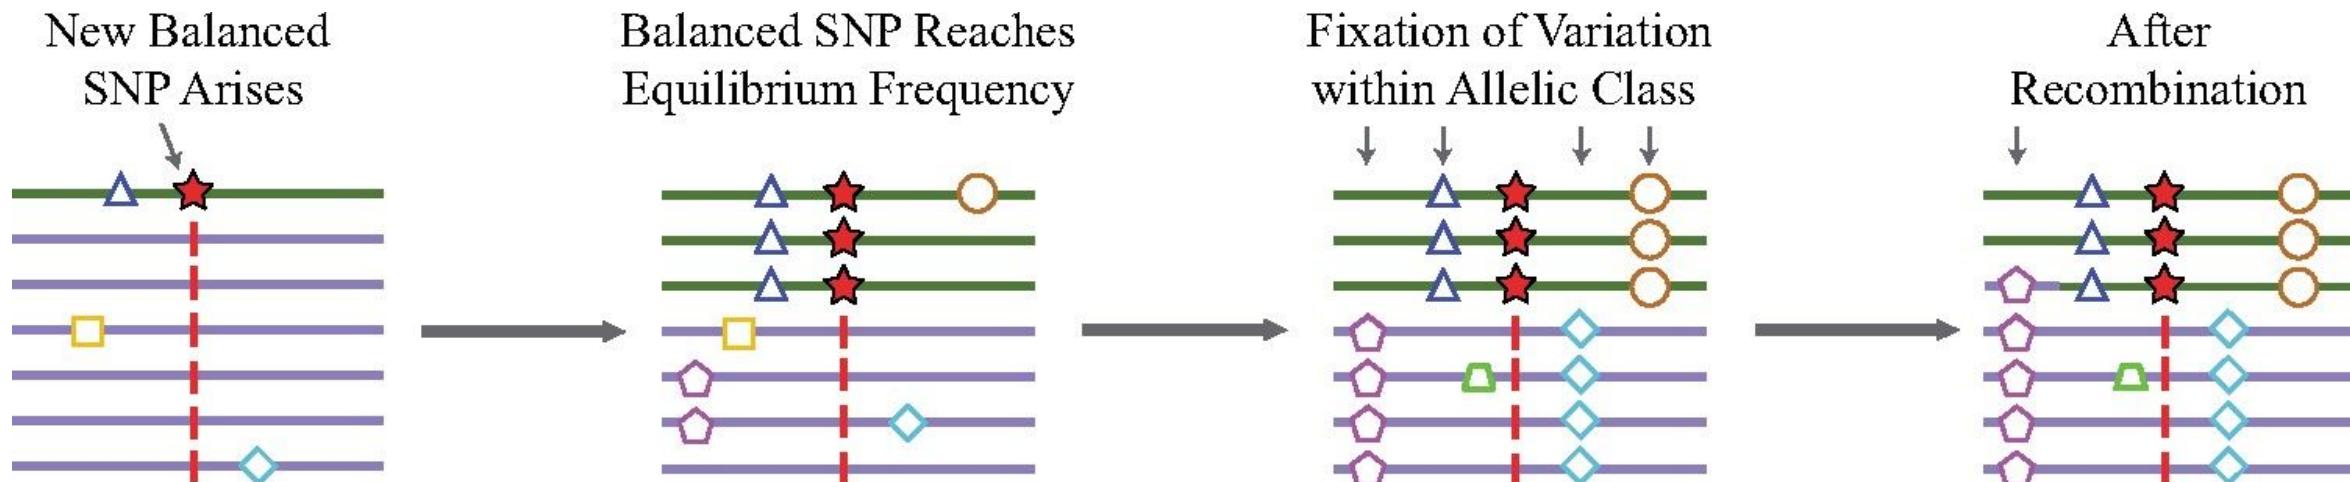
Long-term balancing selection.

- The genealogies near a locus under long-term balancing selection should be older. ARGs can be useful here.
- However, this is also dangerous: misalignment in your (short read) data along undetected paralog loci may lead to false signals
- Very old balancing selection or gene conversion can also lead to very divergent haplotypes, making alignment impossible.



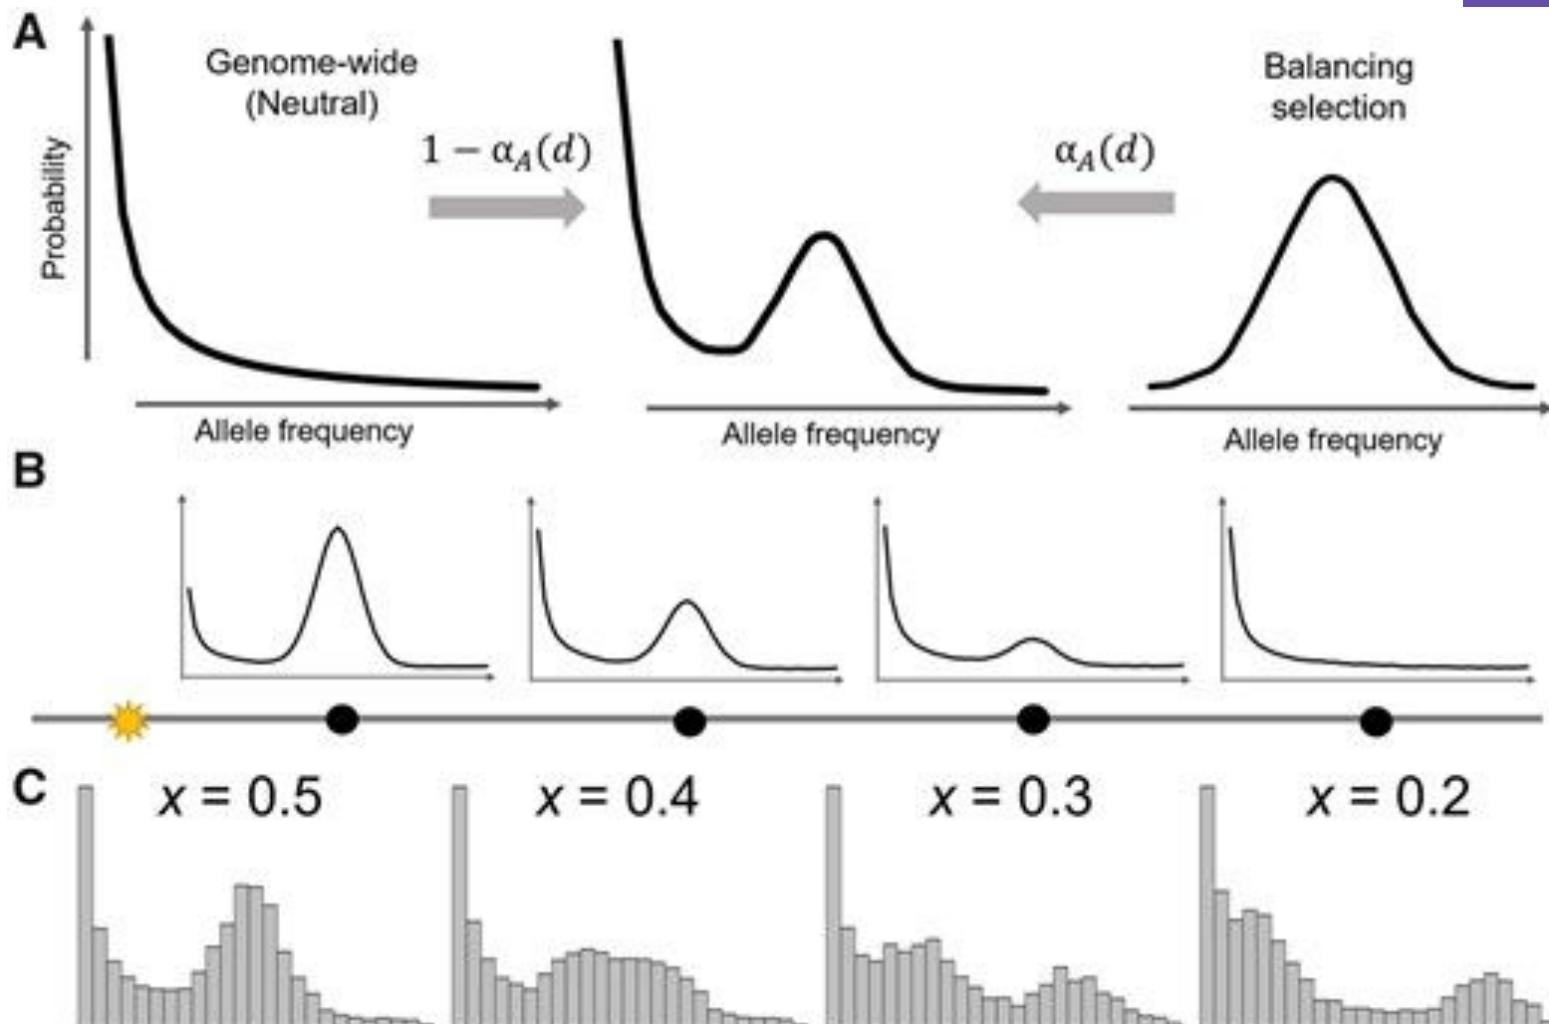
Detecting Balancing selection: LD.

- The β statistic method can work without specifying ancestral and derived alleles (useful in non-model species)
- It is fast.
- Same issues of mappability of short reads to take into account.



Detecting Balancing selection: Spectrum.

- BalLeRMix+. Method can work without specifying ancestral and derived alleles (useful in non-model species)
- It is fast. Robust to the choice of window's size.
- Benefits from an outgroup to include divergence with polymorphism.



Detecting Balancing selection: Spectrum.

- Those are two relatively user-friendly examples
- For a full list of recent methods (2023), see:
- Bárbara D Bitarello, Débora Y C Brandt, Diogo Meyer, Aida M Andrés, Inferring Balancing Selection From Genome-Scale Data, *Genome Biology and Evolution*, Volume 15, Issue 3, March 2023, evad032,
<https://doi.org/10.1093/gbe/evad032>

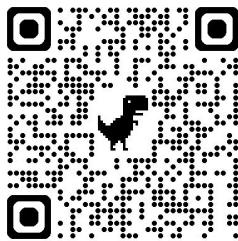
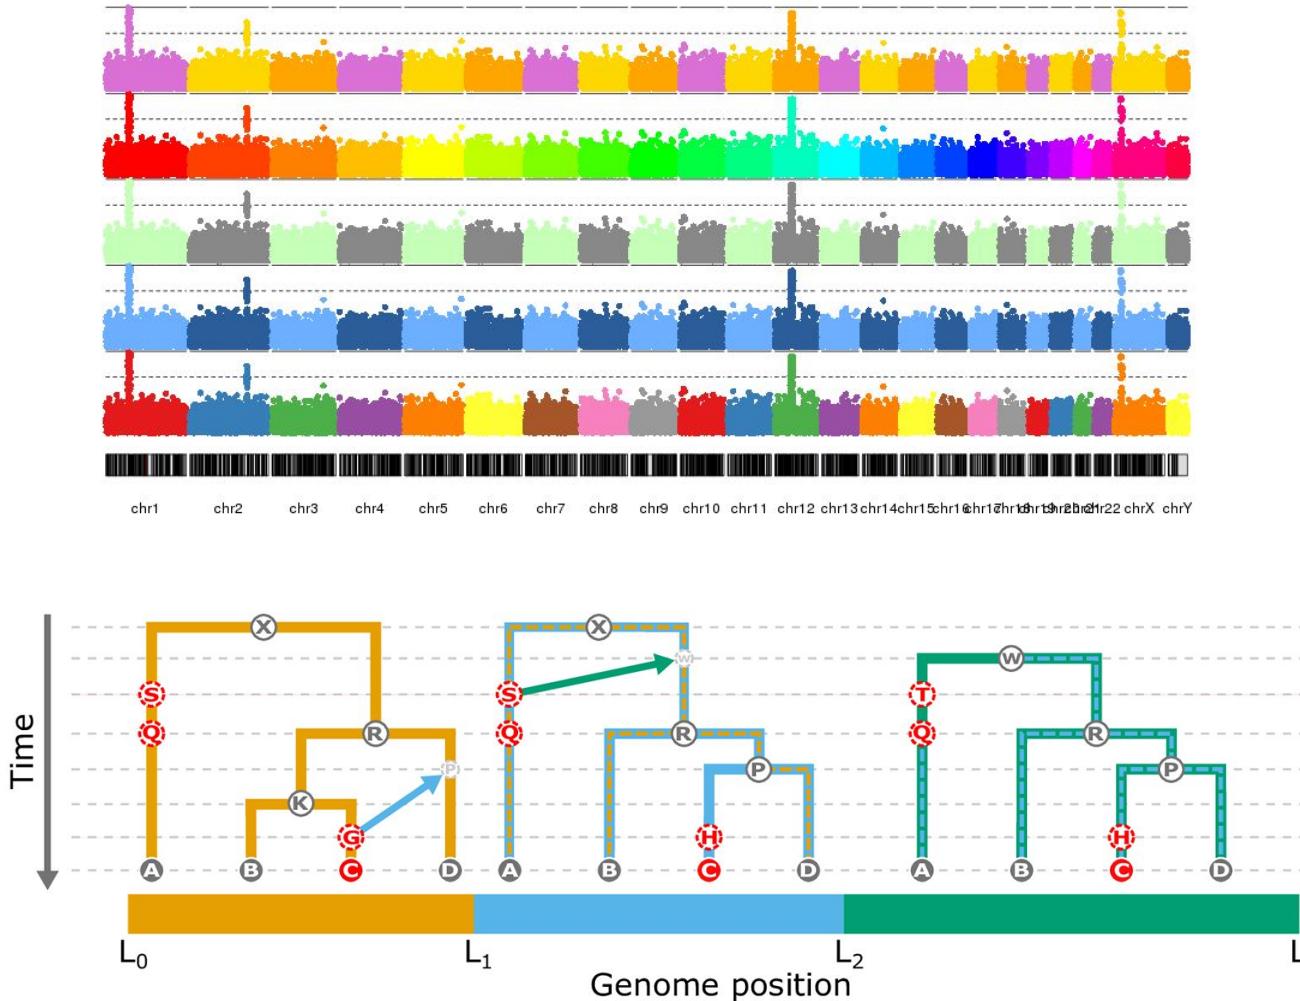


Table 4

Available Software that Implements Tests for BLS

Software	Methods or Tests Implemented	URL
Betascan (python)	$\beta^{(1)*}; \beta^{(1)}; \beta_{std}^{(1)*}; \beta^{(2)}; \beta_{std}^{(2)}$	https://github.com/ksiewert/BetaScan
BALLET (C)	$T_1; T_2$	http://degiorgiogroup.fau.edu/ballet.html
NCD-Statistics (R)	$NCD1; NCD2$	https://github.com/bbitarello/NCD-Statistics
balselr (R)	$NCD1; NCD2$	https://github.com/bbitarello/balselr
MuteBaSS (python)	$NCD2_{trans}; NCD2_{mid,trans}; NCD2_{opt,trans}; HKA_{trans}$	https://github.com/bioXiaoheng/MuteBass
MULLET (C)	$T_{1,trans}; T_{2,trans}$	http://degiorgiogroup.fau.edu/mullet.html
BalLeRMix (python)	$B_0; B_{0,MAF}; B_1; B_2; B_{2,MAF}$	https://github.com/bioXiaoheng/BalLeRMix
BalLeRMix + (python)	$B_0; B_{0,MAF}; B_1; B_2; B_{2,MAF}$	https://github.com/bioXiaoheng/BallerMixPlus
BaSe (python)	Artificial and convoluted neural networks.	https://github.com/ulasisik/balancing-selection
ANGSD	Tajima's D ; Fu & Li's F ; Fu & Li's D ; Fay's H ; Zheng's H	https://github.com/ANGSD/angsd
tskit (python)	Tajima's D	https://github.com/tskit-dev/tskit
Tsel (R)	T_{sel}	https://blogs.cornell.edu/clarklabblog/clark-lab/software
C code	HKA	https://github.com/alanrogers/hka
C code	HKA	https://github.com/andrewkern/hka
C code	HKA	https://bio.cst.temple.edu/~tuf29449/hka_manual

Highly polygenic selection.

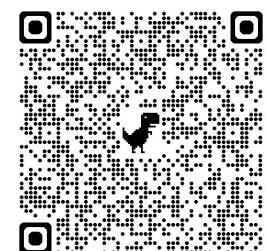


Implemented in
PALM

Obtain effect sizes for SNPs
associated with a trait of interest

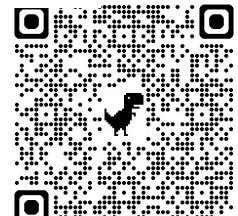
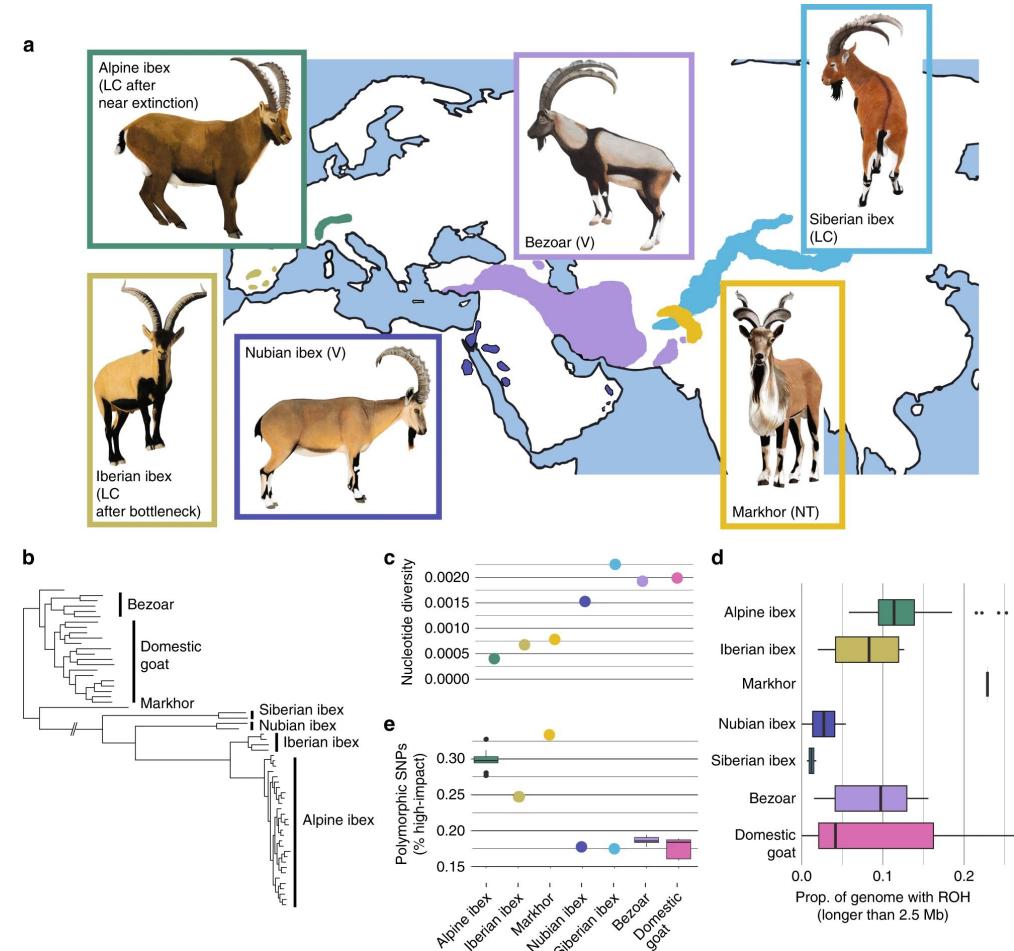
Reconstruct allele trajectories through time
and infer selective coefficients

Obtain genome-wide
genealogies



A final note: the concept of mutation load

- Not everything in selection is positive...
- The genomic burden of deleterious variants is being more and more studied in conservation.
- Demographic processes are important.
- May be extremely important in adaptability.



A final note: the concept of mutation load

- Not everything in selection is positive...
- The genomic burden of deleterious variants is being more and more studied in conservation.
- Demographic processes are important.
- May be extremely important in adaptability.

